

# Robust Fitting

Mathematical Models and Methods for Image Processing

Giacomo Boracchi

<https://boracchi.faculty.polimi.it/>

May 24<sup>th</sup> 2022

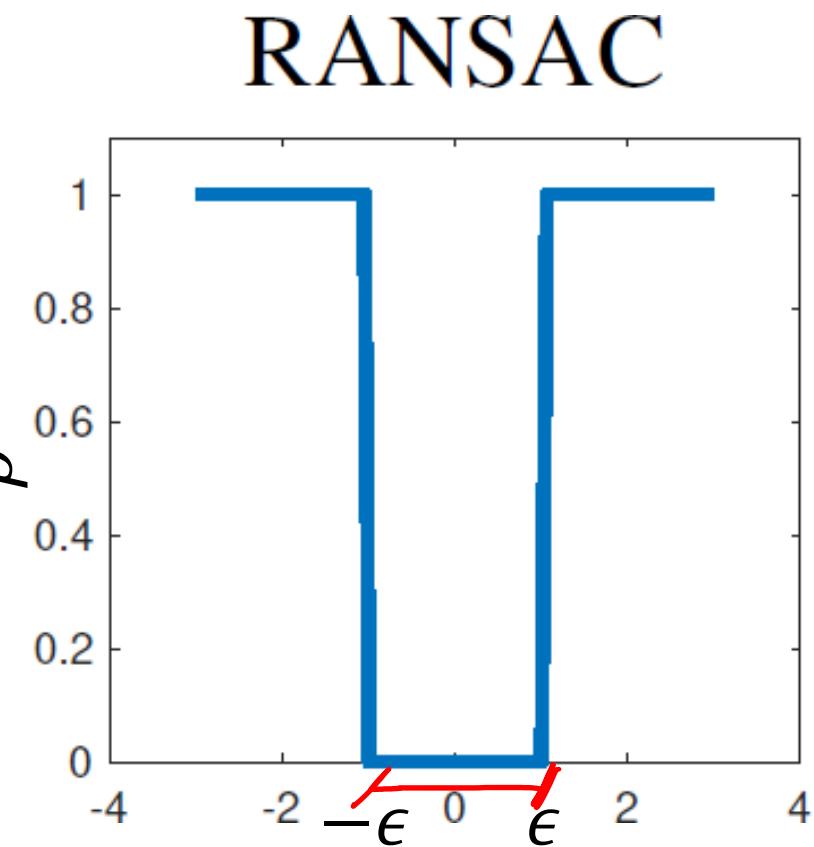
# Ransac as M-estimator

(Steward 1999) RanSaC can be seen as a particular M-estimator since the **loss it minimizes** is the number of points having residual above the inlier threshold  $\epsilon$

$$f(r_i) = \begin{cases} 1, & r_i > \epsilon \\ 0, & r_i \leq \epsilon \end{cases}$$

Of course selecting inlier threshold  $\epsilon$  is very critical

Ransac achieves a theoretical breakdown of 50% of outliers, but in practice, provided a good selection of  $\epsilon$ , this can be even higher

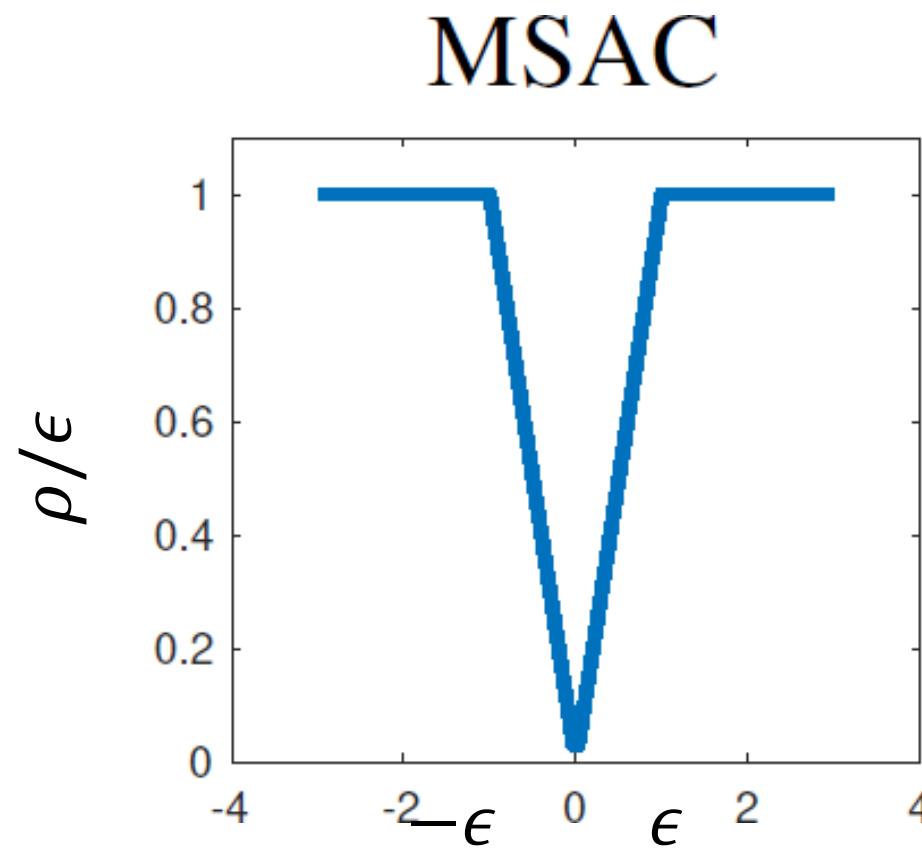


# MSAC

(Torr and Zisserman 2000) a different loss function to be minimized within the RanSaC framework

$$f(r_i) = \begin{cases} \epsilon, & |r_i| > \epsilon \\ |r_i|, & |r_i| \leq \epsilon \end{cases}$$

This turns to be more effective and should be preferred to RanSaC



# Ransac vs MSaC

**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in S} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = +\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Estimate inlier set  $I = \{x \in X : r(x, \theta)^2 < \epsilon^2\}$ ;

Evaluate  $J(\theta) = \sum_{x \in I} r(x, \theta) + (|X| - |I|)\epsilon$ ;

**if**  $J(\theta) < J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

**end**

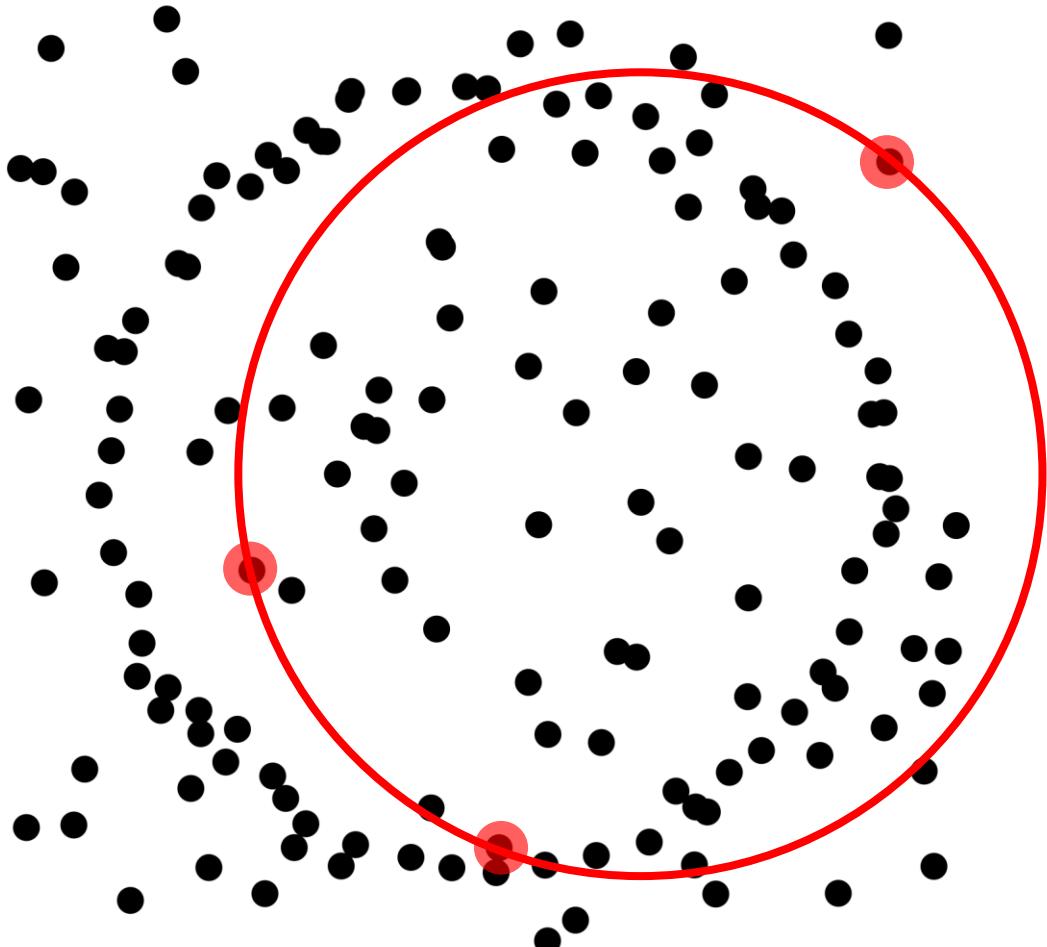
$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Least Median of Squares

# L-meds: Least Median of Squares, Rousseeuw e Leroy (1987)



**Input:**  $X$  data,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = +\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;  
Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \text{median}_{x \in X}(r(x, \theta))$ ;

**if**  $J(\theta) < J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

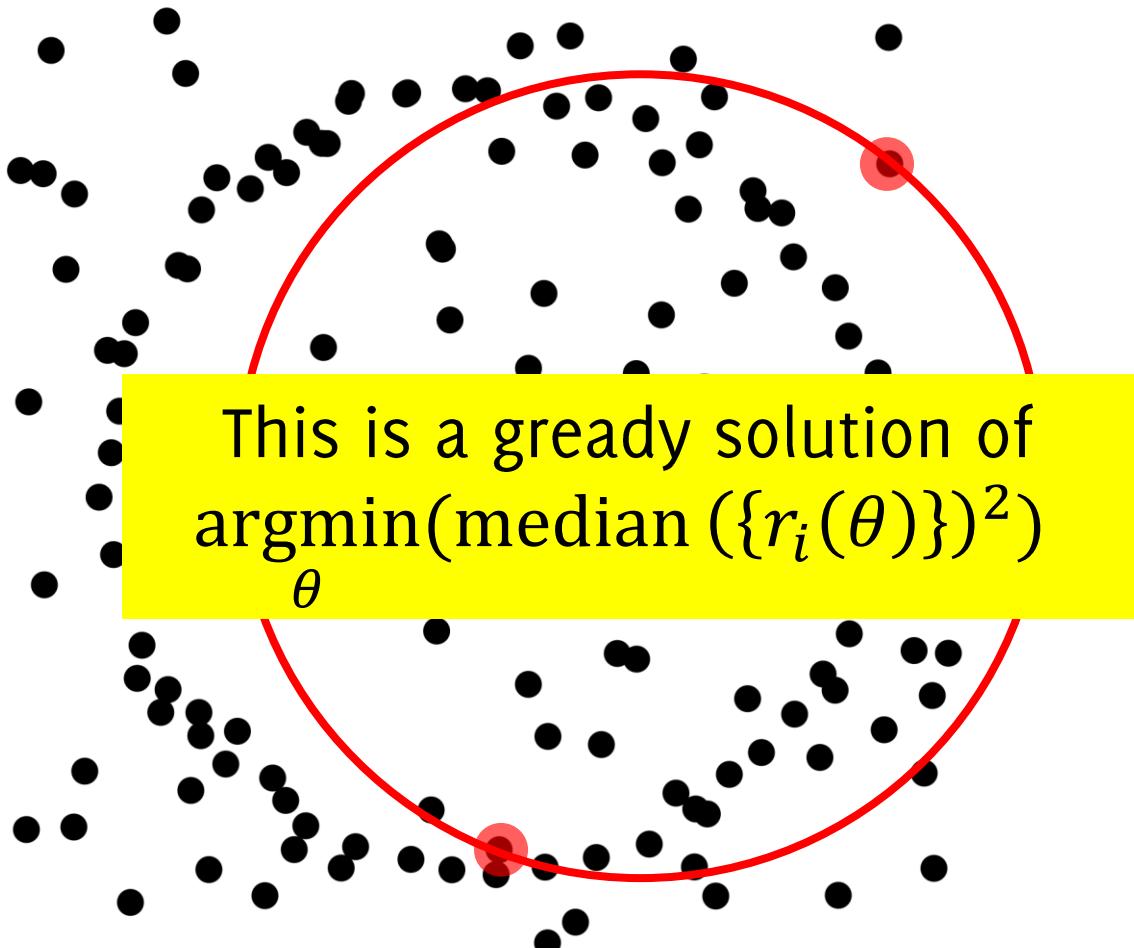
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

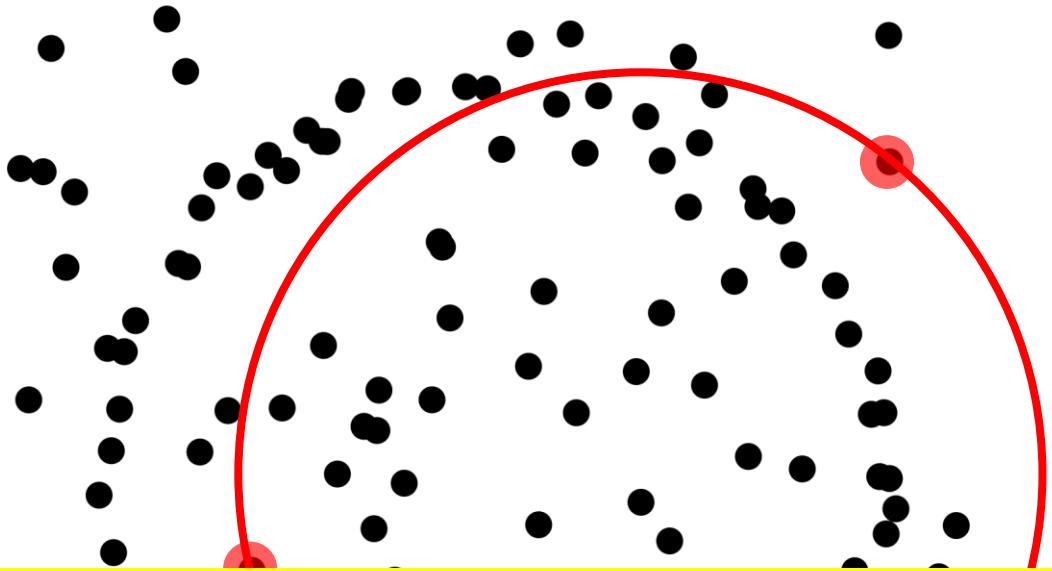
Optimize  $\theta^*$  on its inliers.

# L-meds: Least Median of Squares, Rousseeuw e Leroy (1987)



**Input:**  $X$  data,  $k_{\max}$  max iteration  
**Output:**  $\theta^*$  model estimate  
 $J^* = +\infty, k = 0;$   
**repeat**  
    Select randomly a minimal sample set  $S \subset X$ ;  
    Estimate parameters  $\theta$  on  $S$ ;  
    Evaluate  $J(\theta) = \text{median}_{x \in S} (r(x, \theta))$ ;  
    **if**  $J(\theta) < J^*$  **then**  
         $\theta^* = \theta$ ;  
         $J^* = J(\theta)$ ;  
    **end**  
     $k = k + 1$ ;  
**until**  $k > k_{\max}$ ;  
Optimize  $\theta^*$  on its inliers.

# L-meds: Least Median of Squares, Rousseeuw e Leroy (1987)



Since there is no explicit definition of inliers here, inliers can be identified as points having residuals (w.r.t. to the final model) that are smaller than  $2.5\sigma$

**Input:**  $X$  data,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = +\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;  
Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \text{median}_{x \in X}(r(x, \theta))$ ;

**if**  $J(\theta) < J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

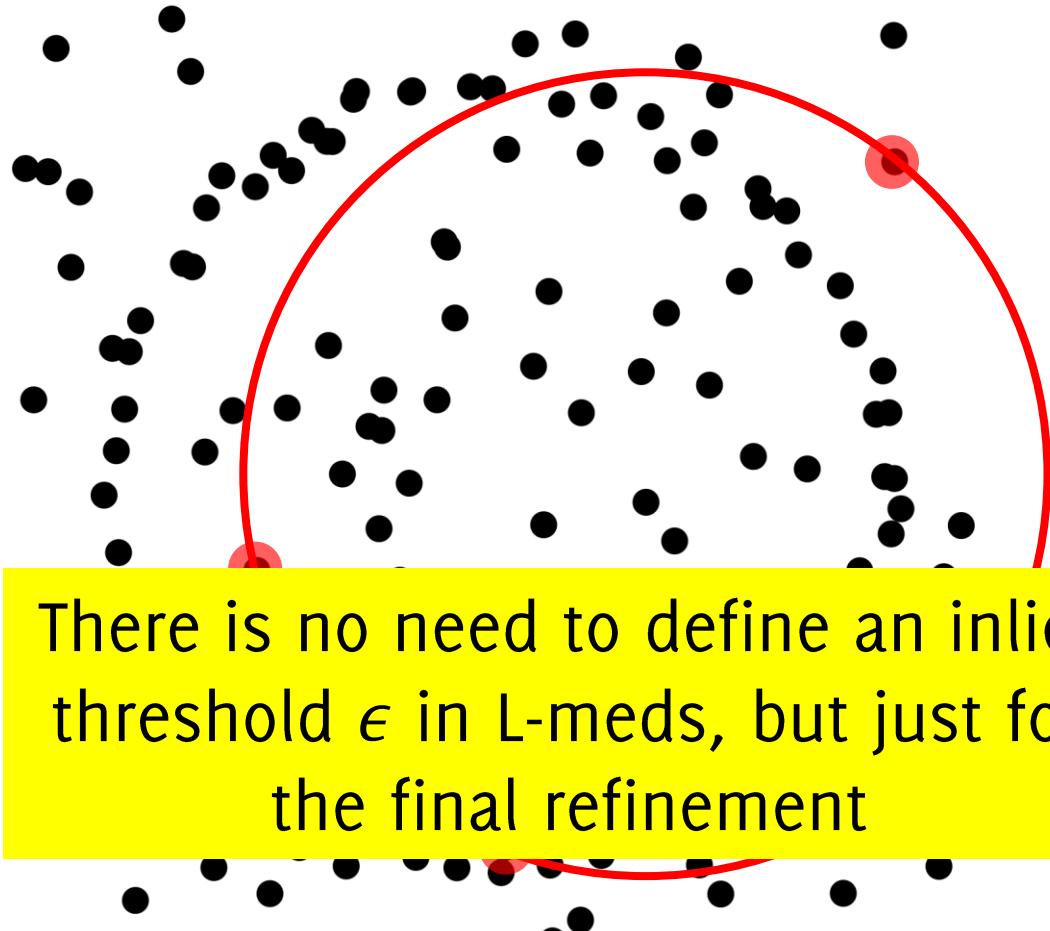
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# L-meds: Least Median of Squares, Rousseeuw e Leroy (1987)



**Input:**  $X$  data,  $k_{\max}$  max iteration  
**Output:**  $\theta^*$  model estimate  
 $J^* = +\infty, k = 0;$   
**repeat**  
    Select randomly a minimal sample set  $S \subset X$ ;  
    Estimate parameters  $\theta$  on  $S$ ;  
    Evaluate  $J(\theta) = \text{median}_{x \in X}(r(x, \theta))$ ;  
    **if**  $J(\theta) < J^*$  **then**  
         $\theta^* = \theta$ ;  
         $J^* = J(\theta)$ ;  
    **end**  
     $k = k + 1$ ;  
**until**  $k > k_{\max}$ ;  
Optimize  $\theta^*$  on its inliers.

The problem of fitting multiple  
geometric primitives  
is ubiquitous in Computer Vision

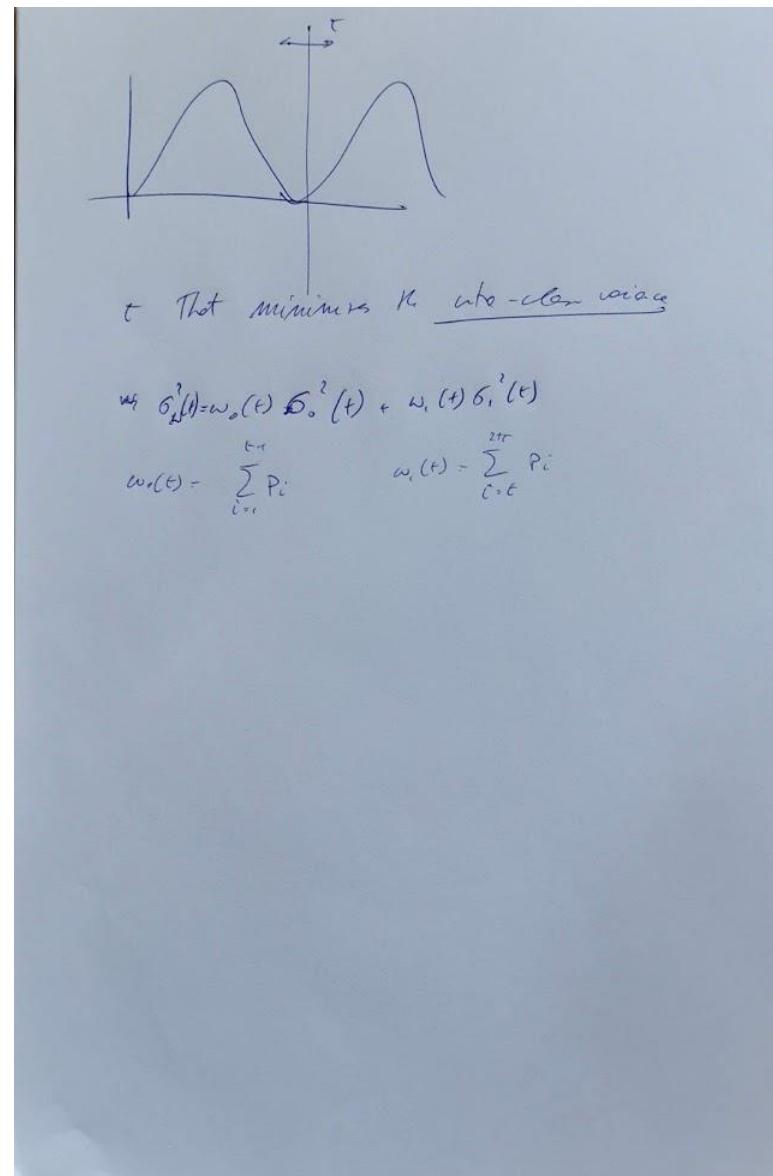
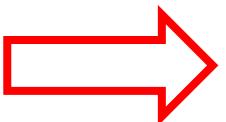
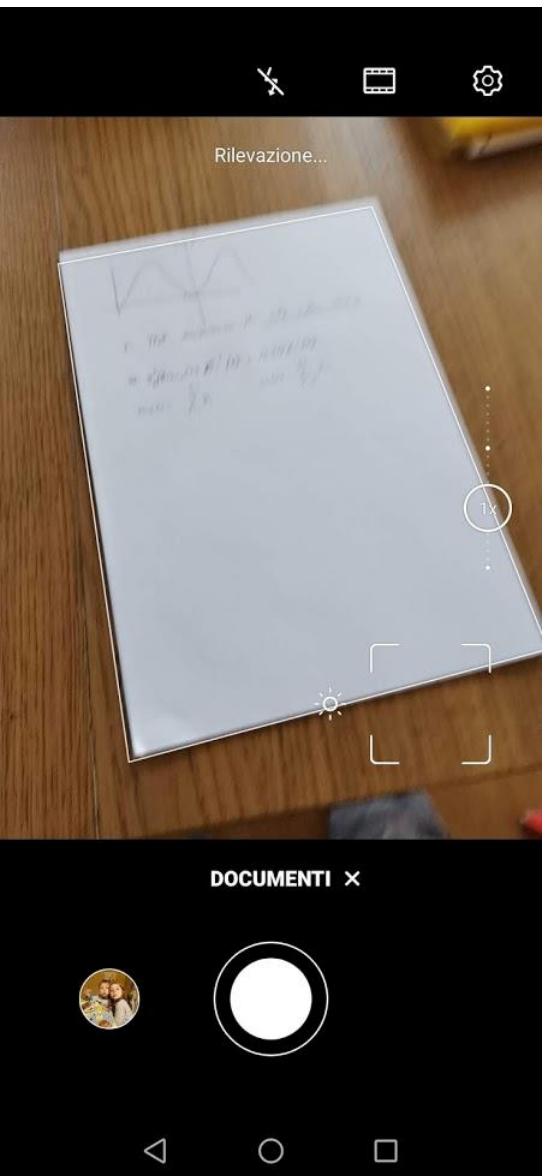
# Example: Line Fitting for Vanishing Point Estimation



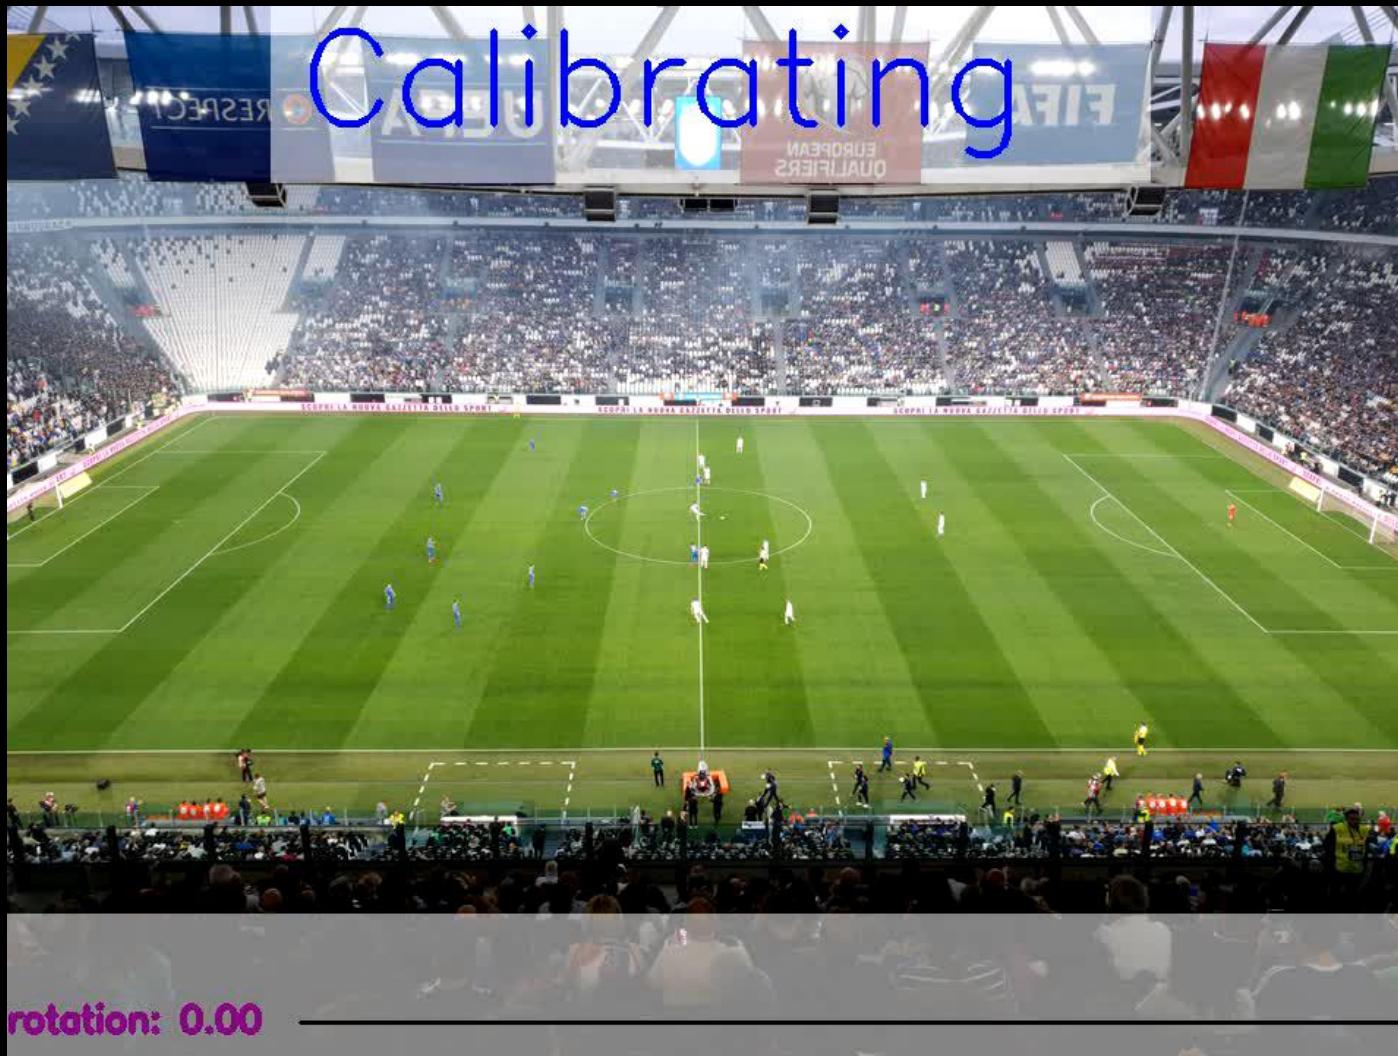
# Example: Conic Fitting



# Line Detection is Important



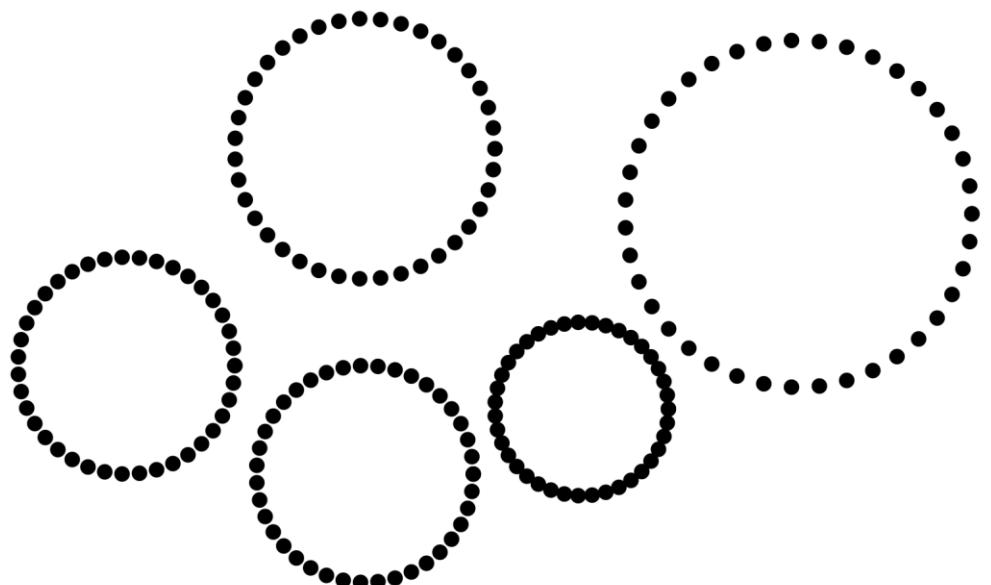
## Calibrating



Powered by

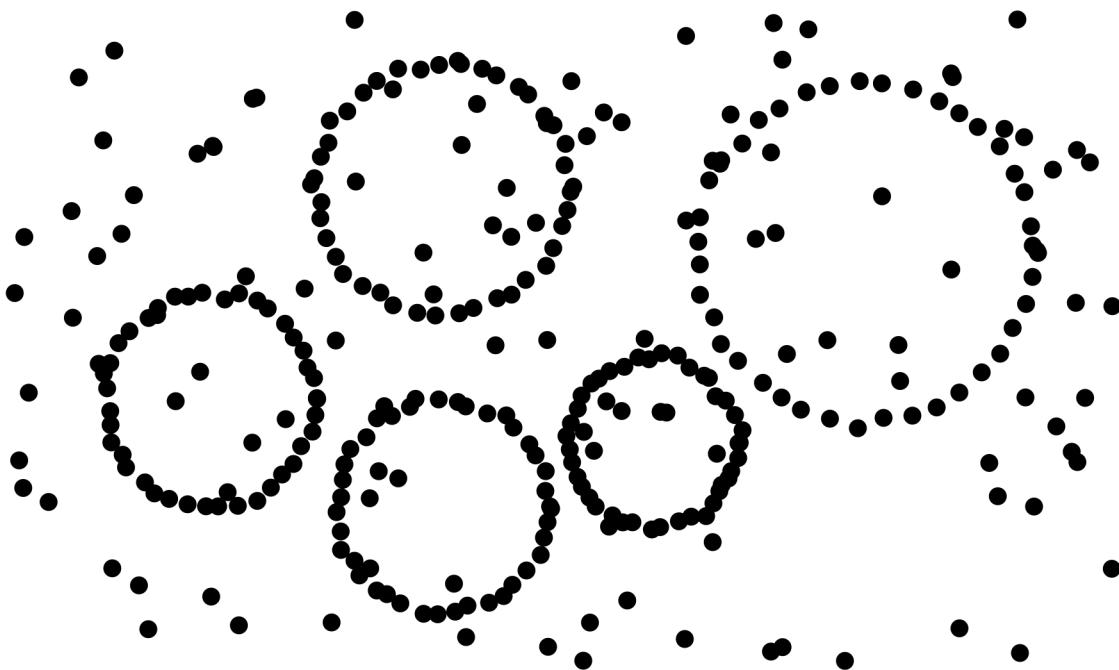
# The problem of multi-model fitting (or structure recovery)

Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



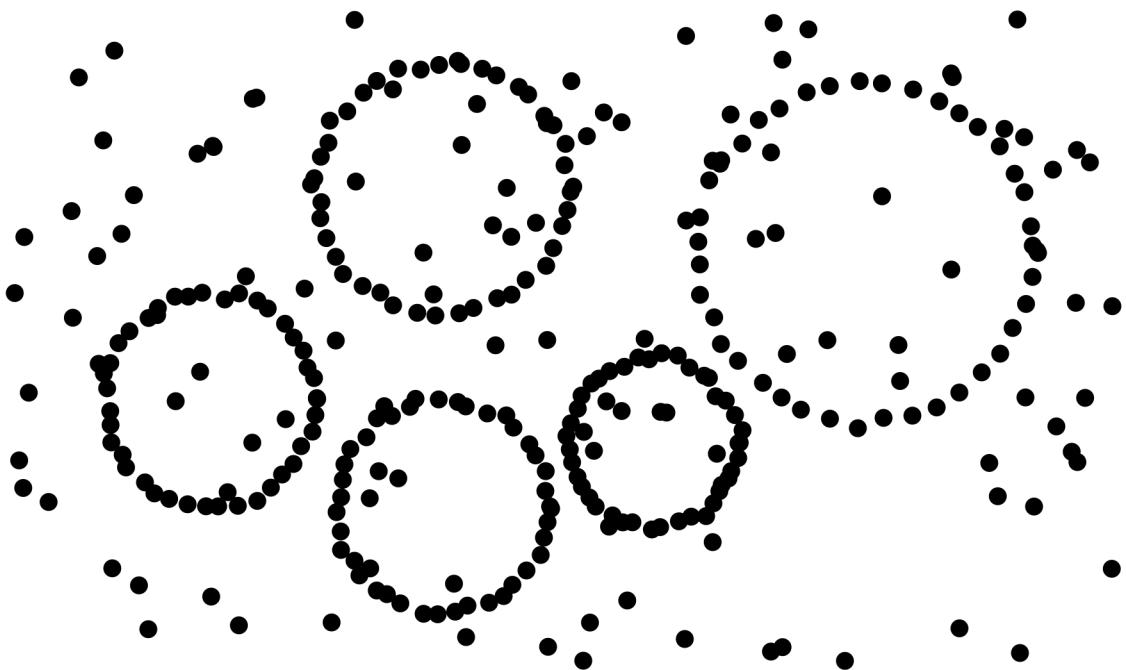
# The problem of multi-model fitting (or structure recovery)

Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



# The problem of multi-model fitting (or structure recovery)

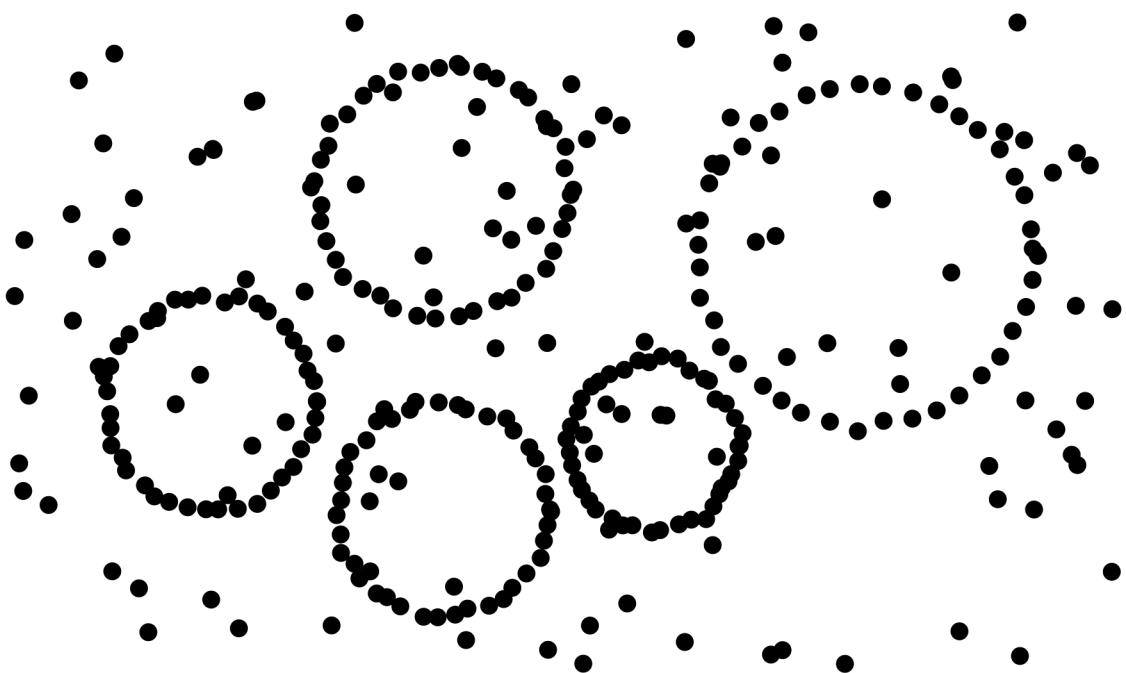
Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



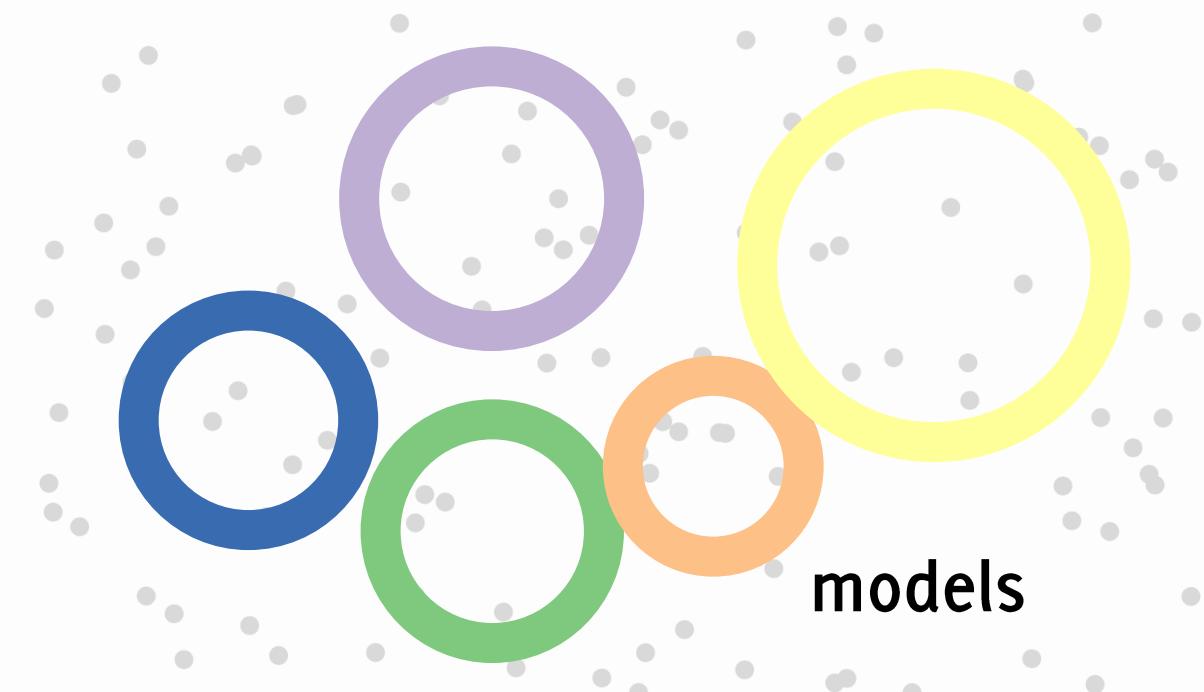
Goal: automatically estimate the models that best explain the data/discover the structures hidden in the data

# The problem of multi-model fitting (or structure recovery)

Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



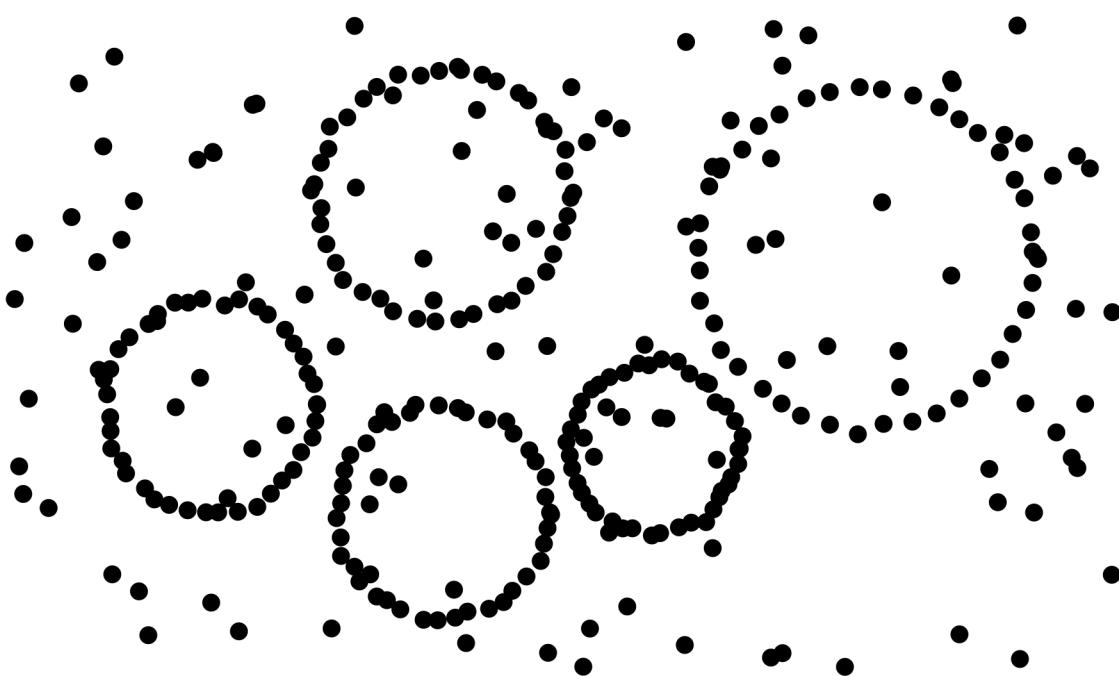
Goal: automatically estimate the models that best explain the data/discover the structures hidden in the data



“in the eye of the beholder”, mathematical descriptions of the data that an observer fits

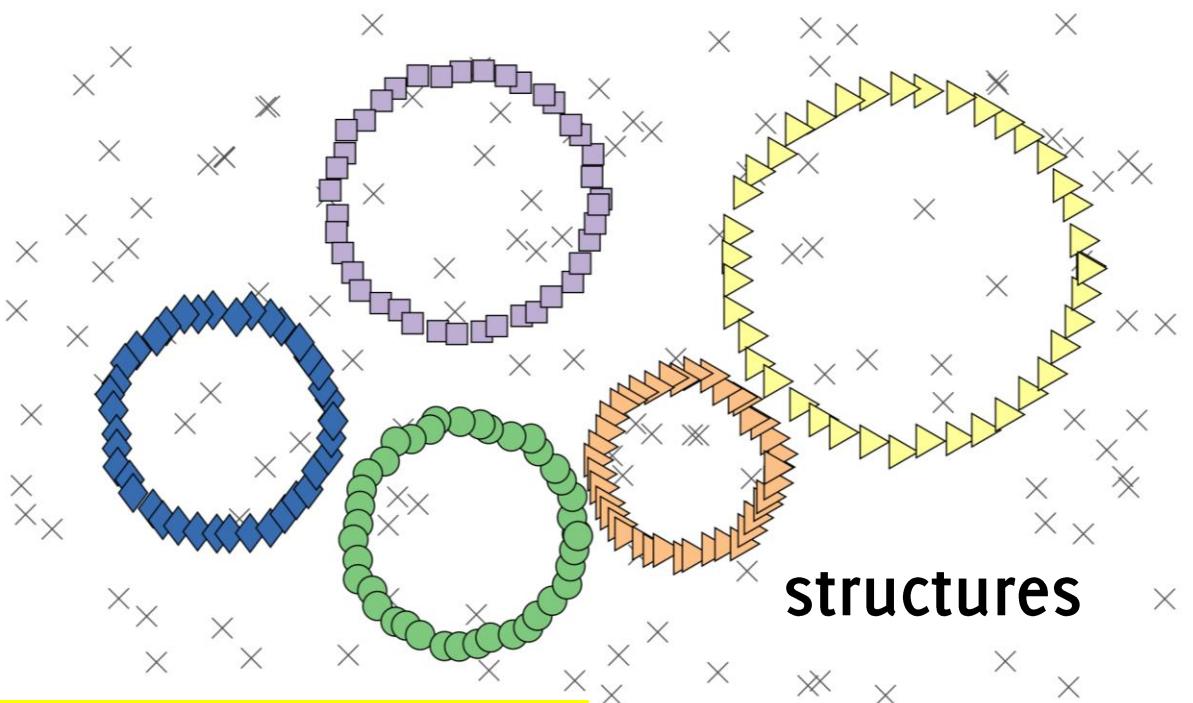
# The problem of multi-model fitting (or structure recovery)

Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$

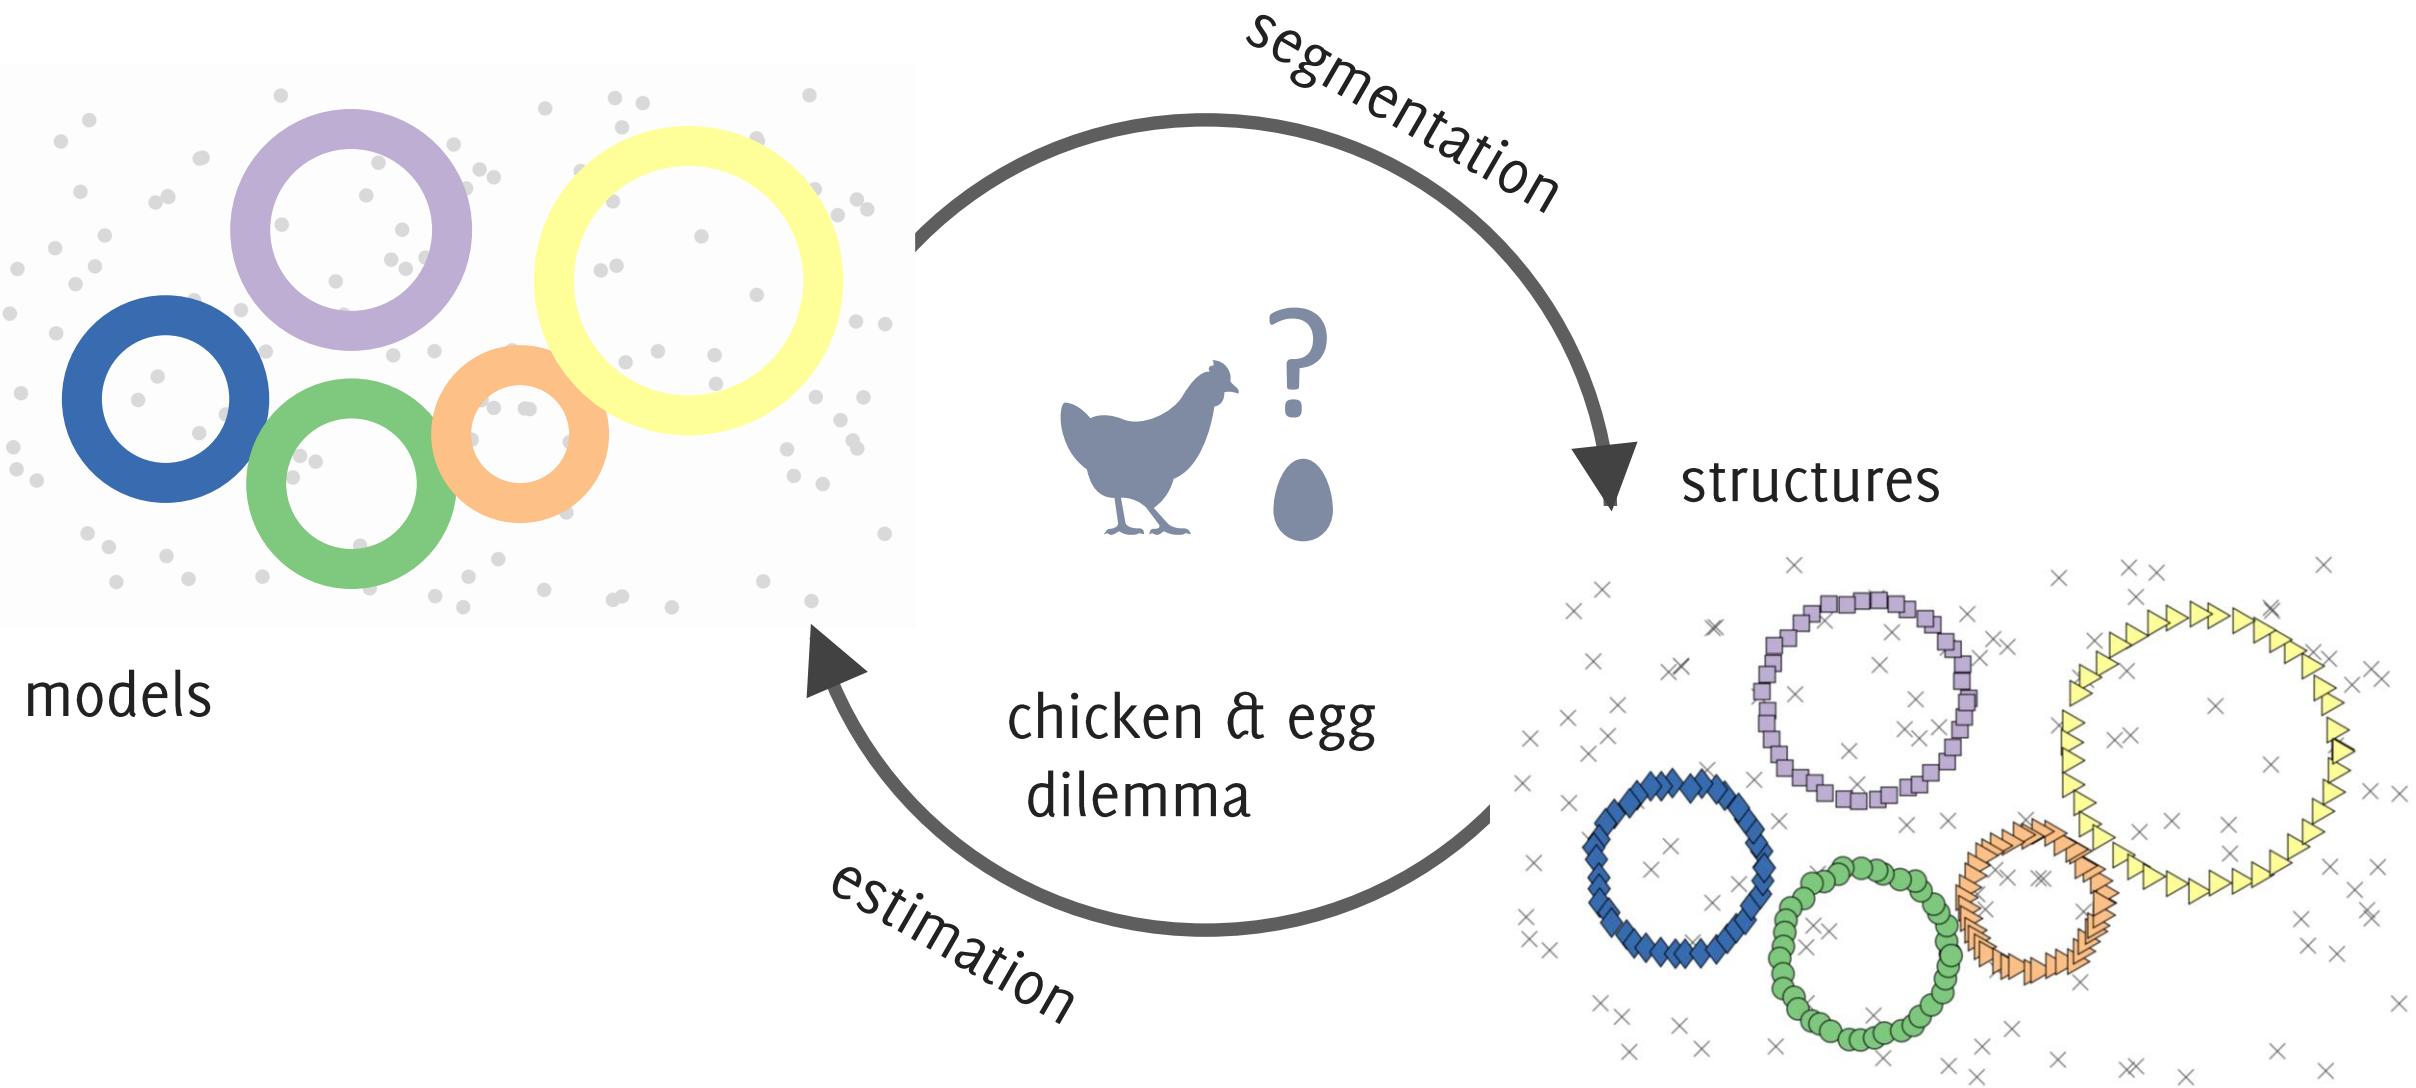


relations among the data, intrinsic to data

Goal: automatically estimate the models that best explain the data/discover the structures hidden in the data

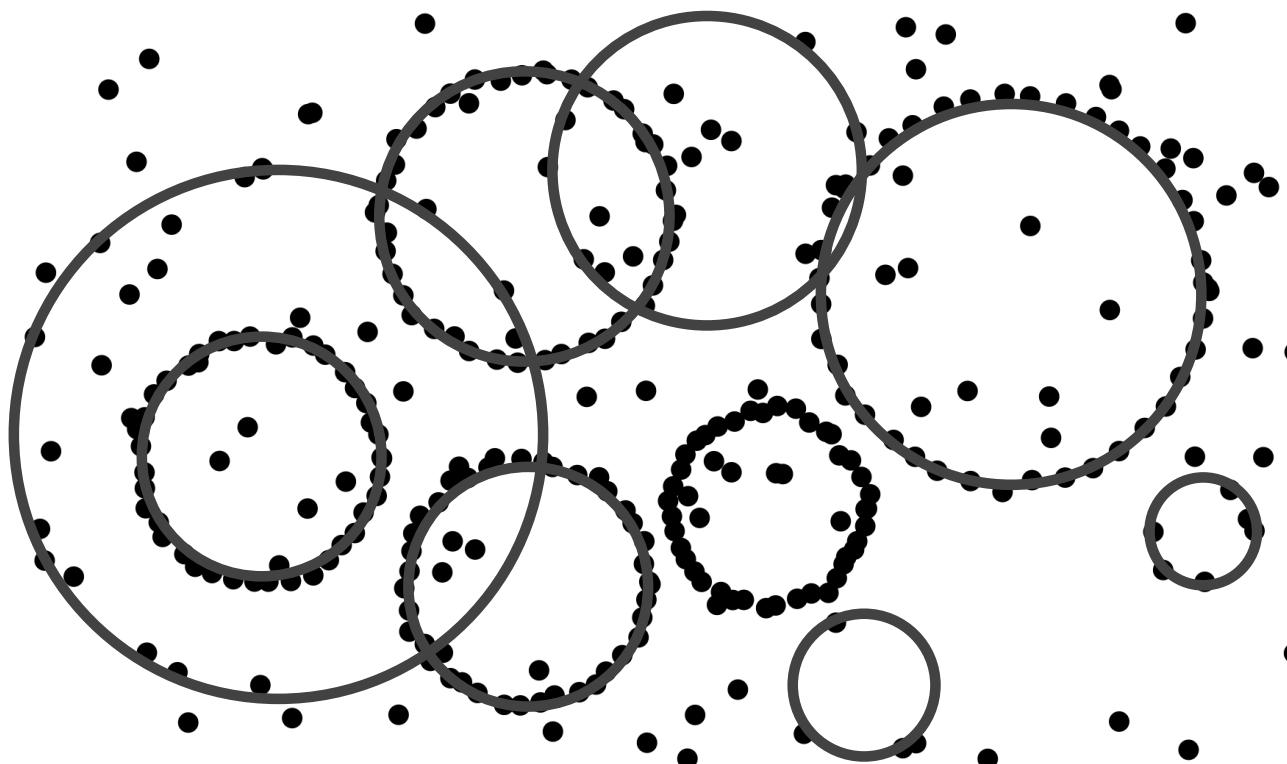


# The Challenges of multi-model fitting



# The Challenges of multi-model fitting

Number of models?



Inliers/outliers?

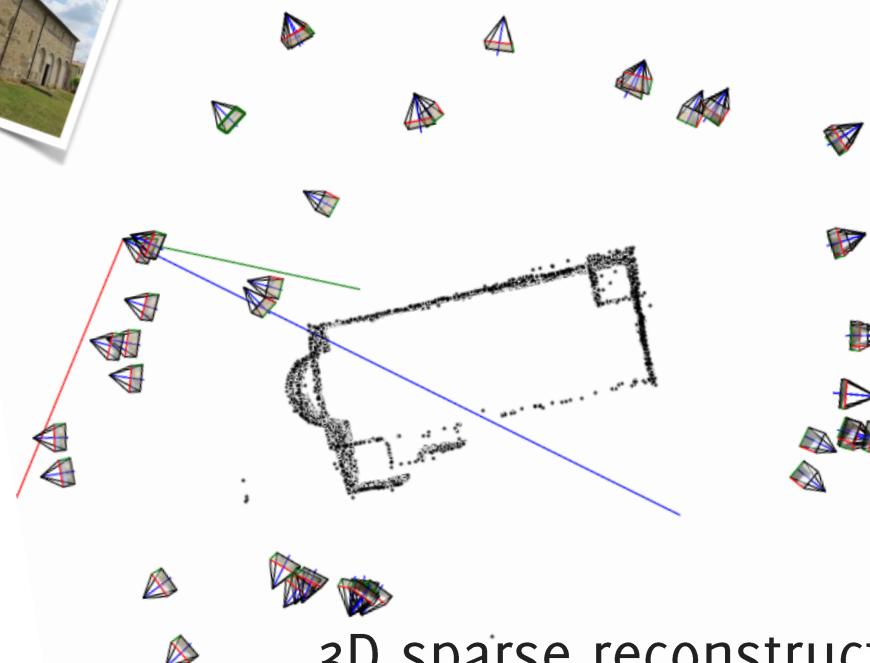
ill posed

# **Multi-Model Fitting**

# Multi model fitting applications: primitive fitting

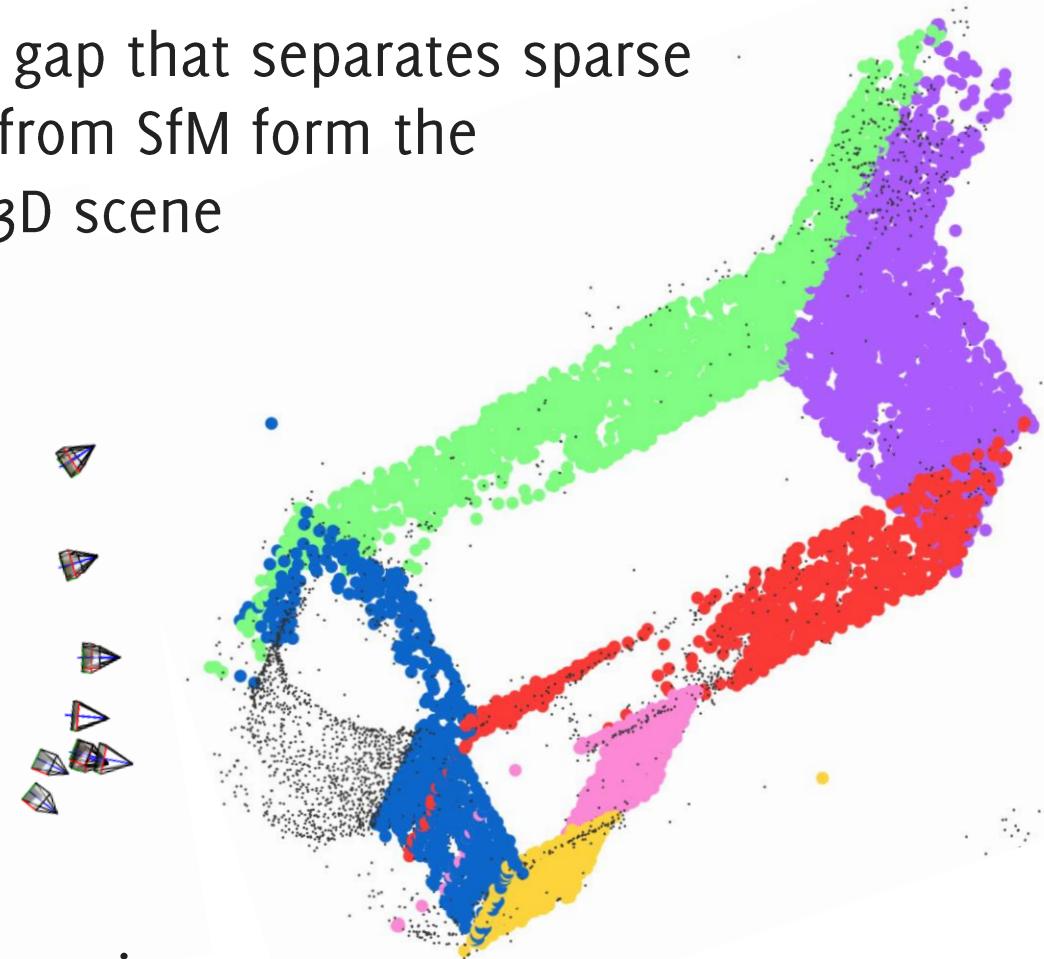


input images



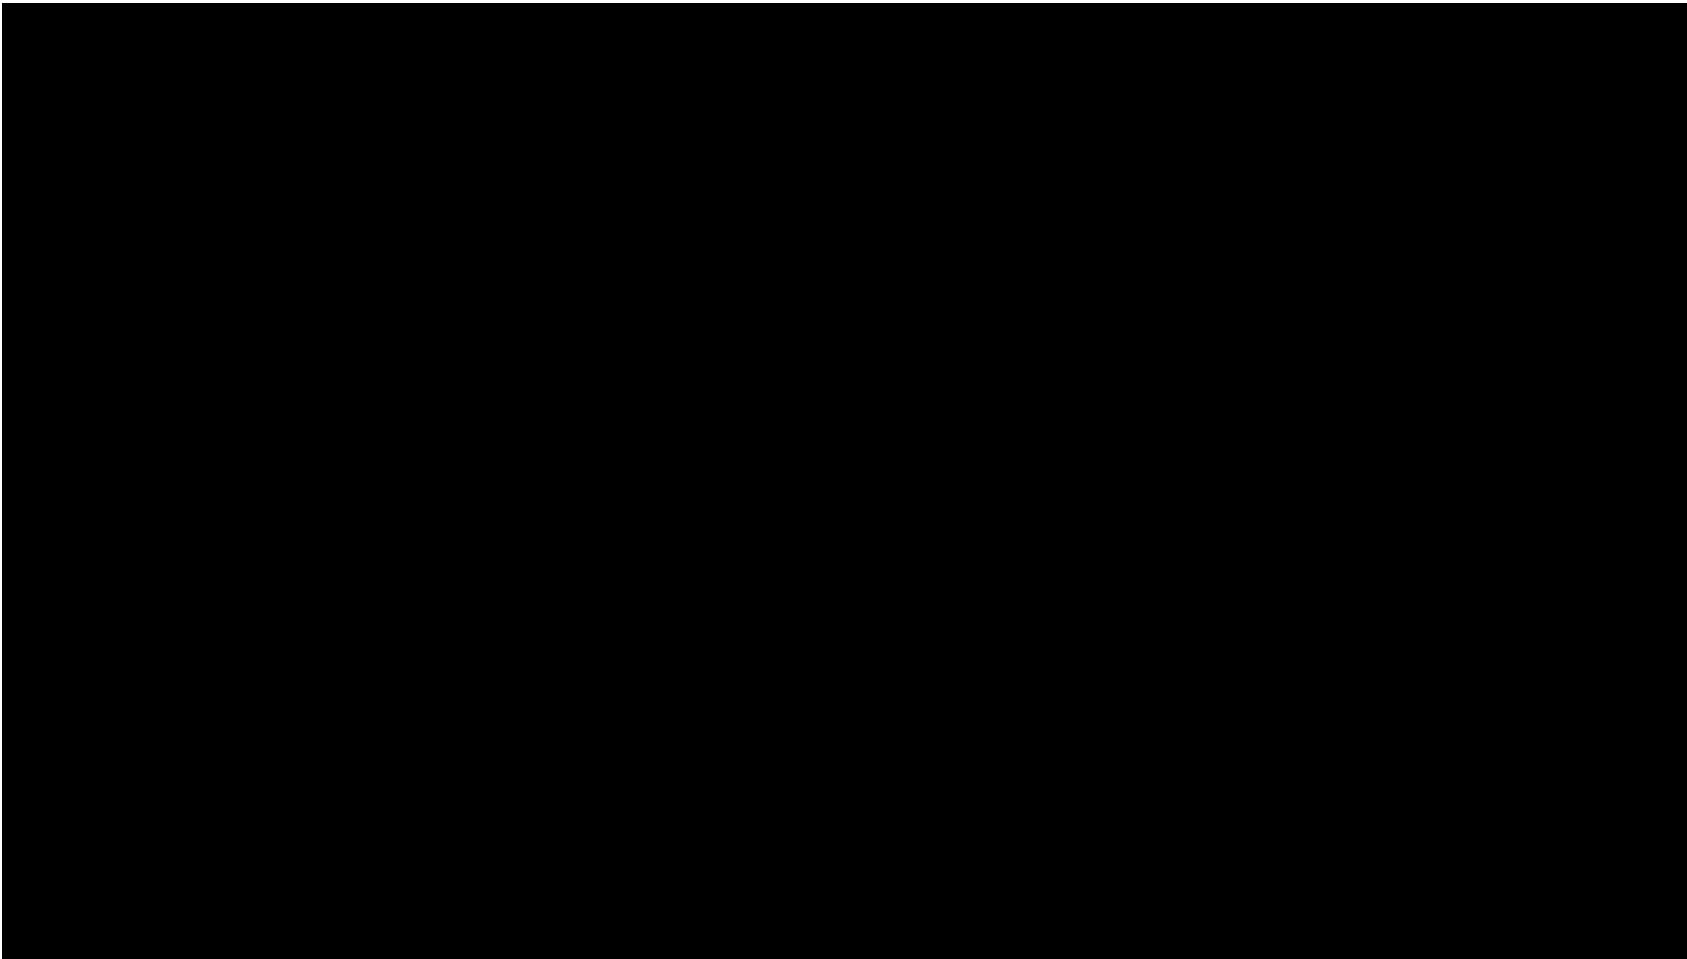
3D sparse reconstruction

Bridge the semantic gap that separates sparse point cloud coming from SfM form the understanding of a 3D scene



$$X \subset \mathbb{R}^3, \Theta = \text{planes}$$

# Multimodel fitting for 3D scattered data



L. Magri, and A. Fusiello. "Reconstruction of interior walls from point cloud data with min-hashed J-linkage." 2018 3DV

L. Magri, and A. Fusiello. "IMPROVING AUTOMATIC RECONSTRUCTION OF INTERIOR WALLS FROM POINT CLOUD DATA." International Archives of the Photogrammetry, Remote Sensing & Spatial Information Sciences (2019).

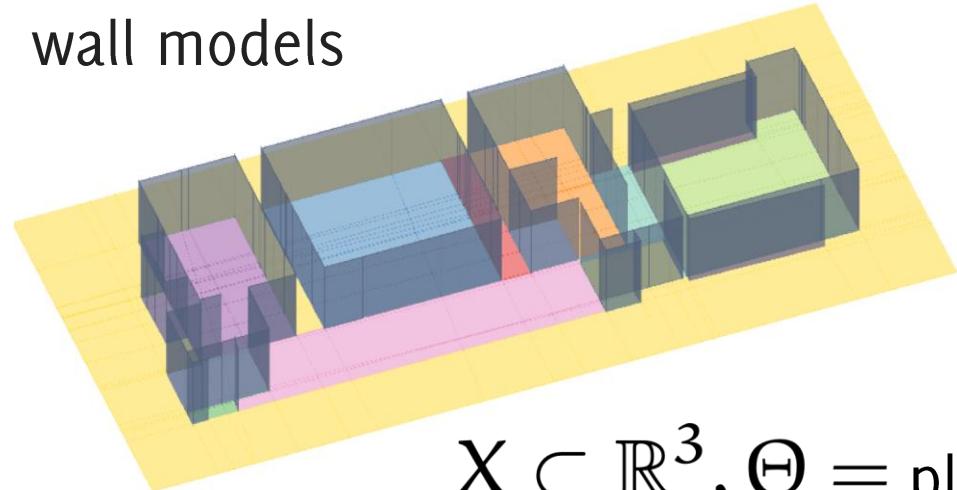
L. Magri, and Andrea Fusiello. "T-linkage: A continuous relaxation of j-linkage for multi-model fitting." CVPR 2014

# Multi model fitting applications: scan2bim

scanned point cloud

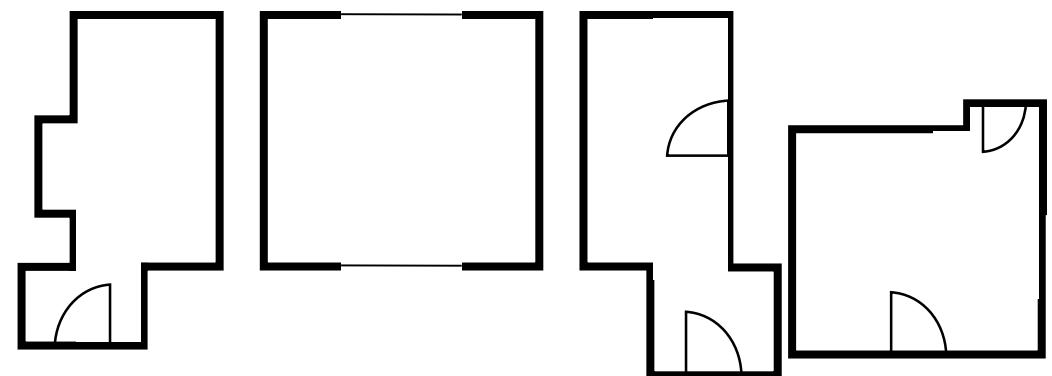


wall models



$$X \subset \mathbb{R}^3, \Theta = \text{planes}$$

floor-plan



Given a scanned point cloud of an interior environment, detect its primary facility surfaces – such as floors, walls, and ceilings.

$$X \subset \mathbb{R}^2, \Theta = \text{lines}$$

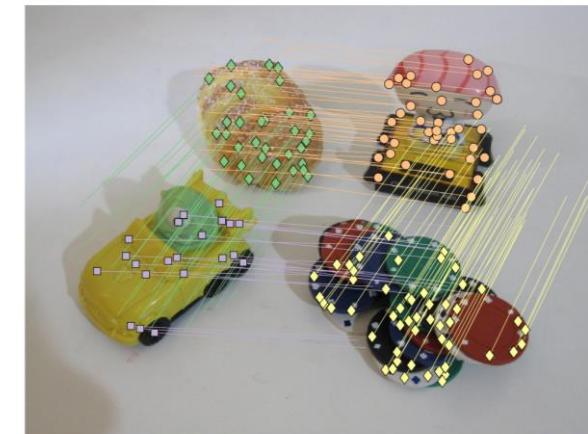
# Multi model fitting applications: two view geometry

Geometric fit on corresponding matches across two images

plane detection



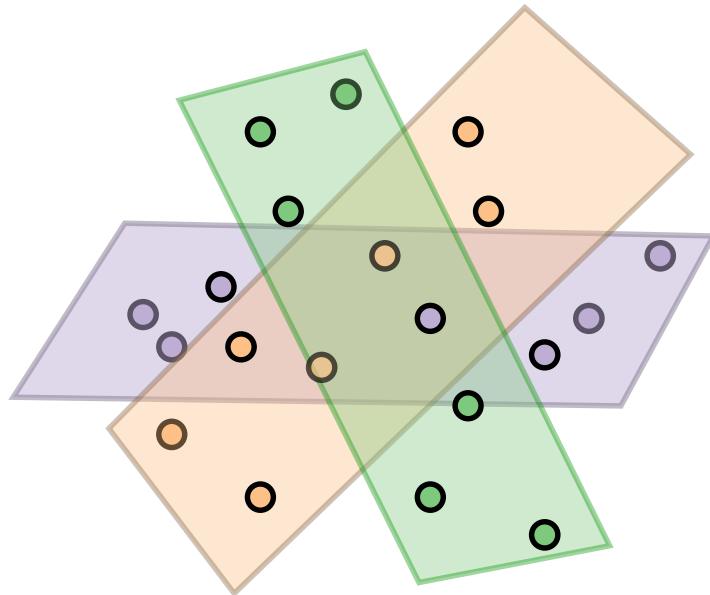
epipolar geometry



$X \subset \mathbb{R}^4$ ,  $\Theta = \text{homographies}$

$X \subset \mathbb{R}^4$ ,  $\Theta = \text{fundamental matrices}$

# Multi model fitting applications: subspace clustering



$$X \subset \mathbb{R}^d, \Theta = \text{subspaces}$$

3D Video segmentation



Face clustering

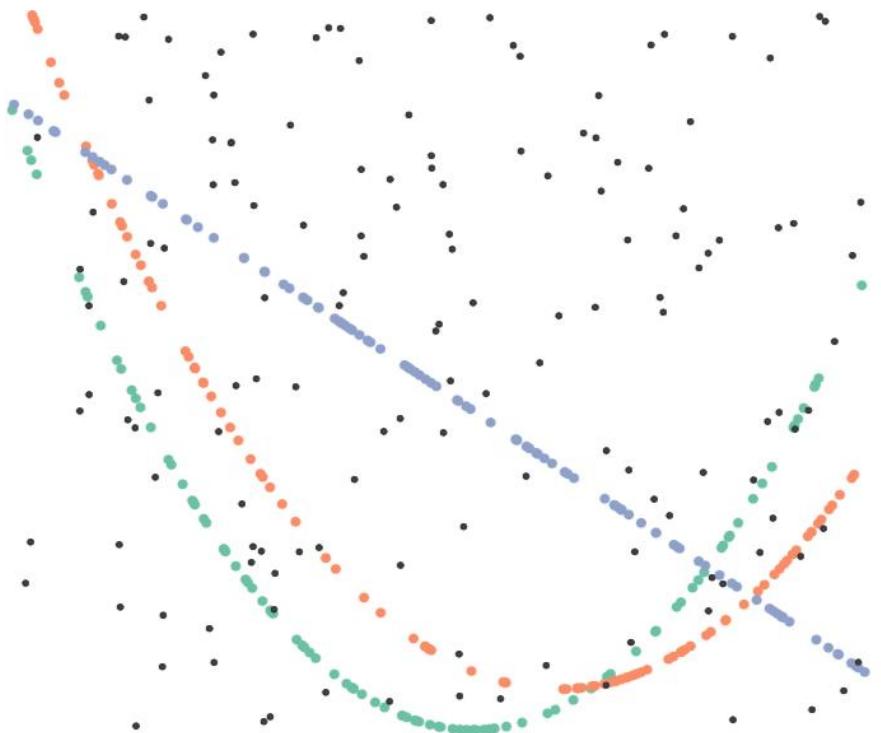


# Template Detection

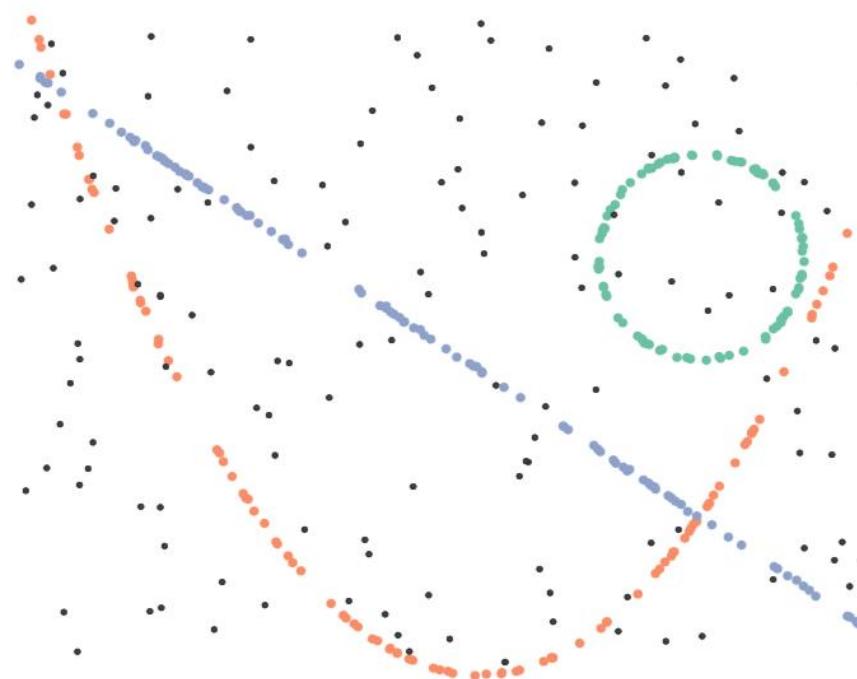


Our Collaboration with an Italian Company T&O

# Multimodel (and multi-class) fitting



(a)



(b)

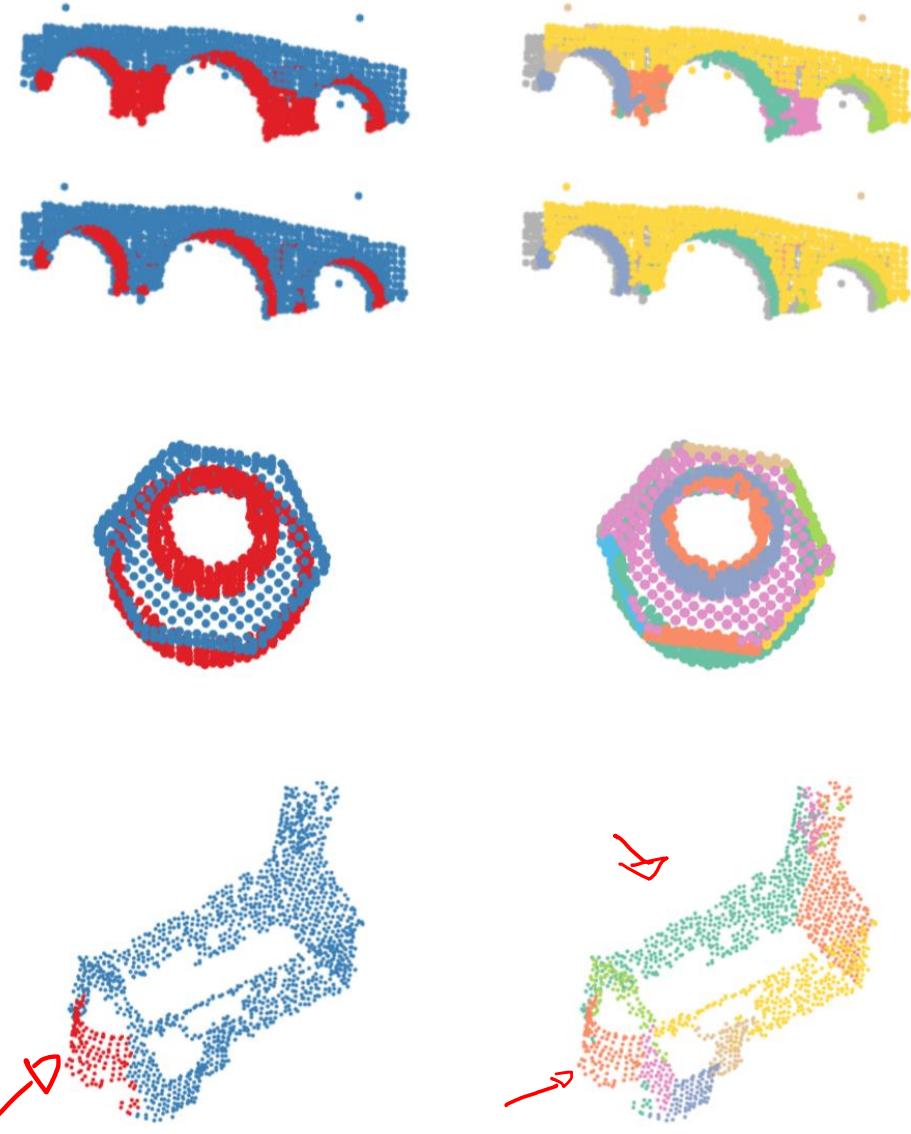


(c)

# Multimodel (and multi-class) fitting



Image of the scene



colour coded class

colour coded model

# MultiLink

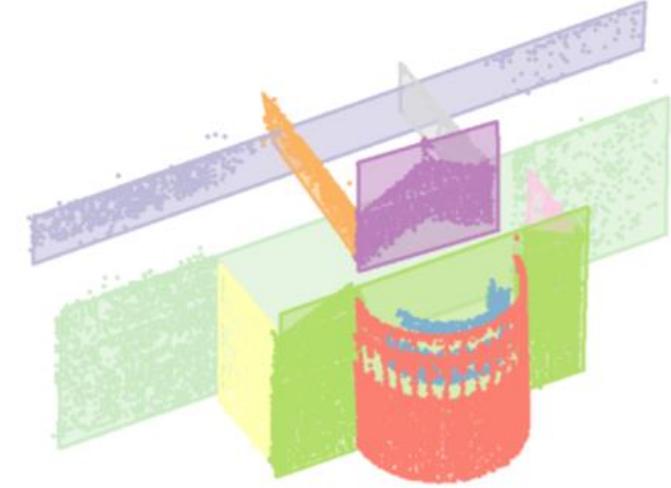
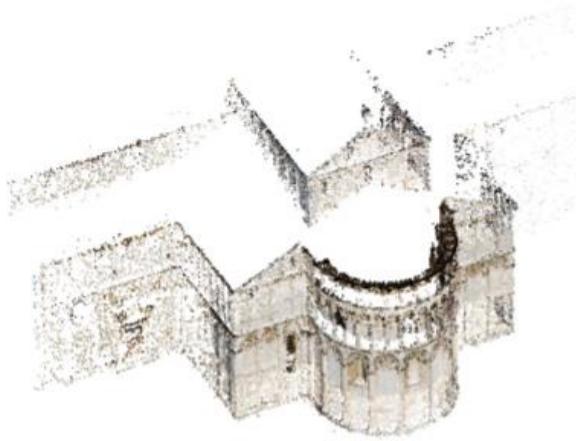
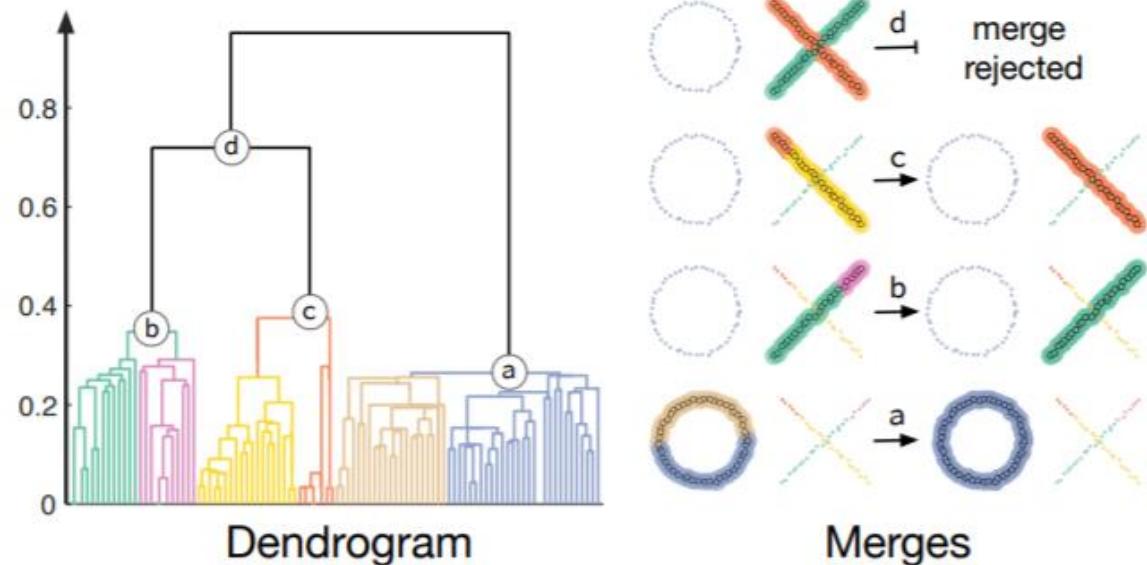
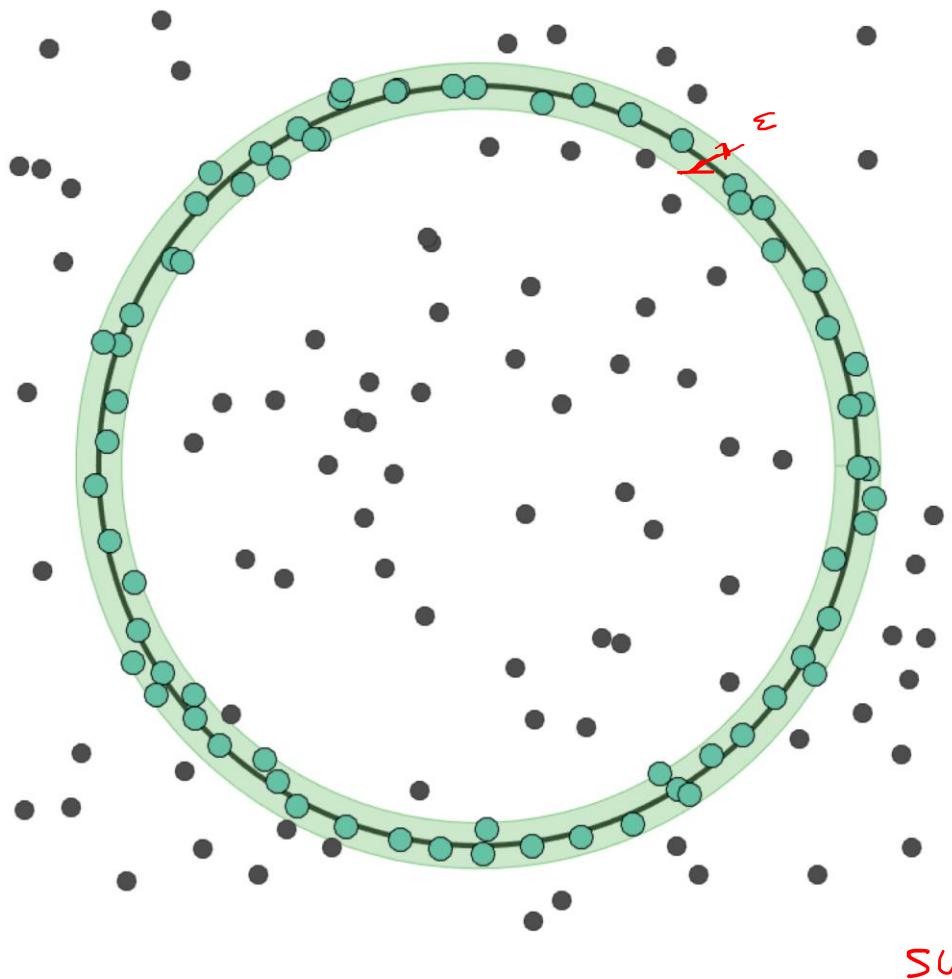


Figure 2: MultiLink combines single-linkage clustering and GRIC. Clusters are merged as long as the GRIC score improves when fitting suitable models on-the-fly. Colors indicate how cluster aggregation proceeds in the dendrogram.

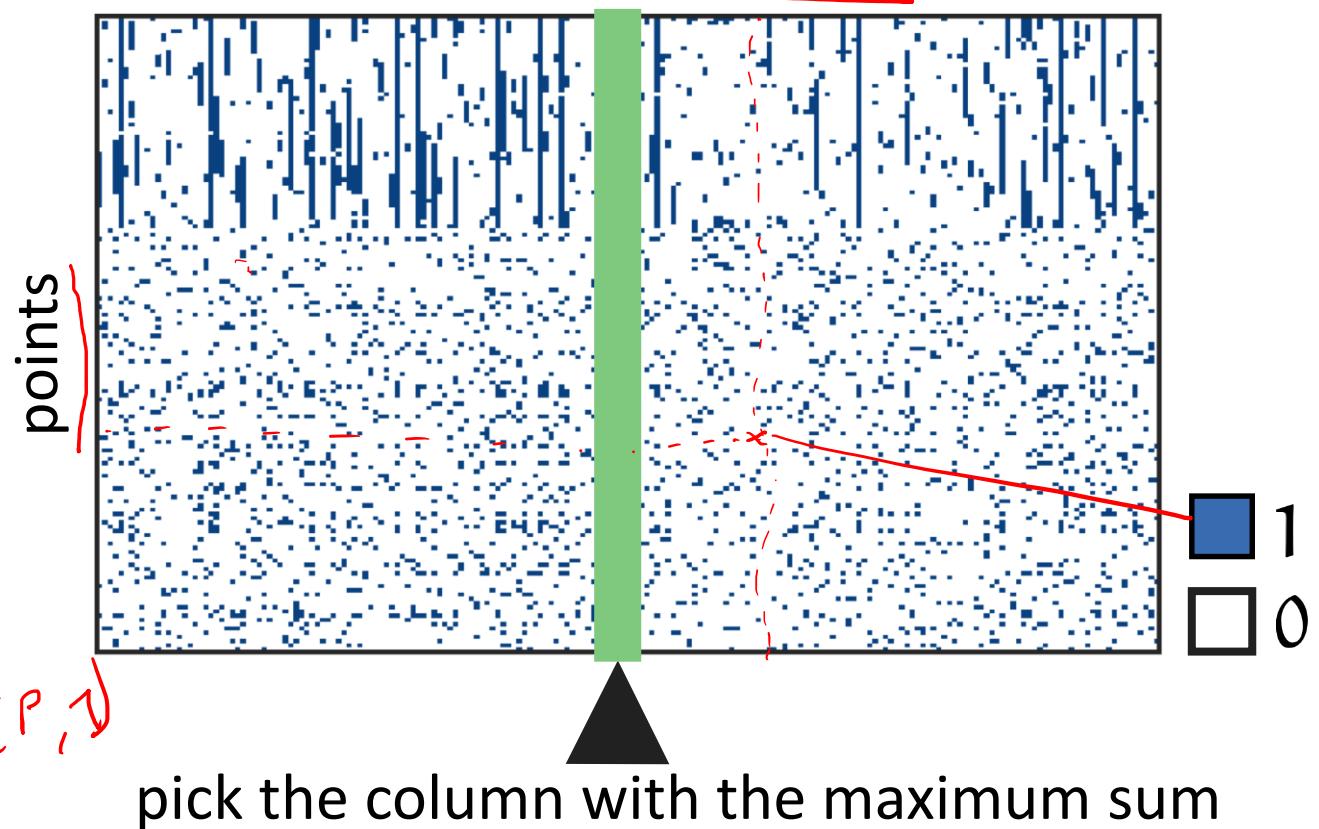
# Multi-model Fitting Solutions

# Let's go back to RanSaC



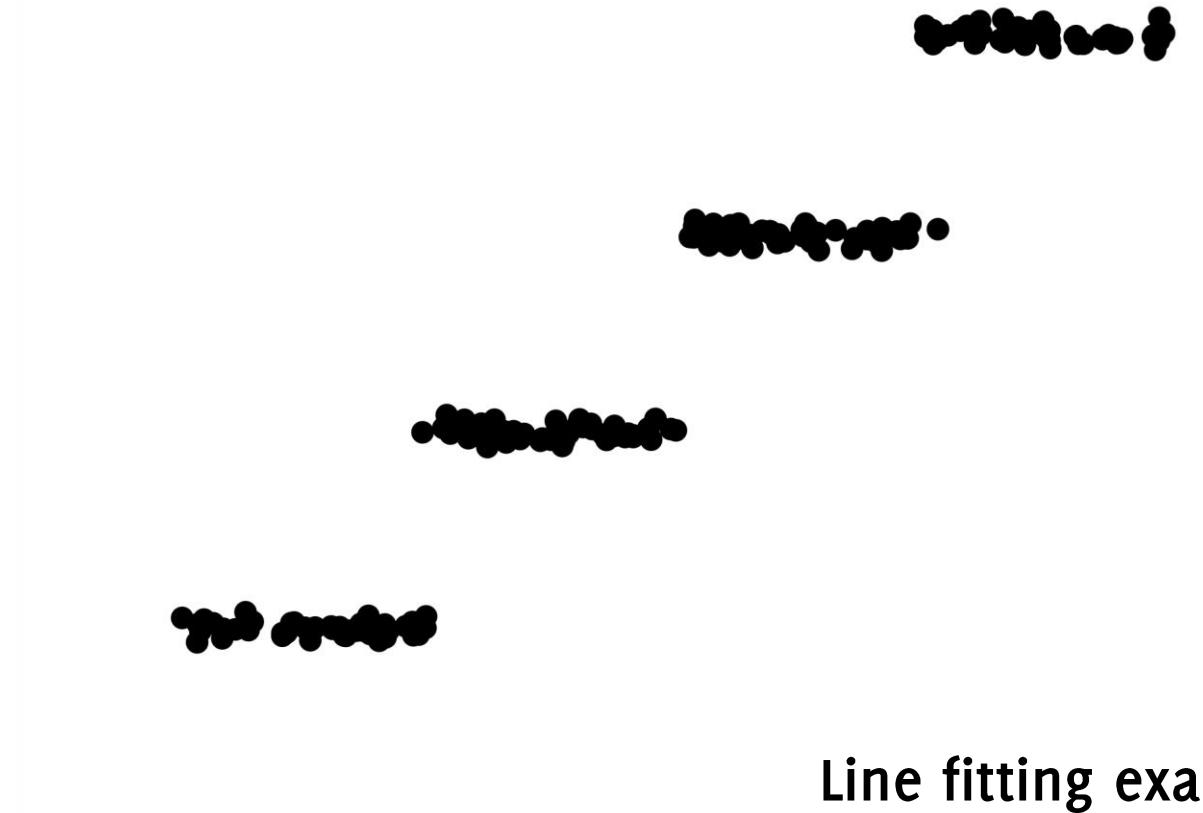
Data driven search of model space  
 $H = \{\theta_1, \theta_2, \dots, \theta_m\} \approx \Theta$

tentative models



# Sequential RanSaC [Zuliani 05]

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

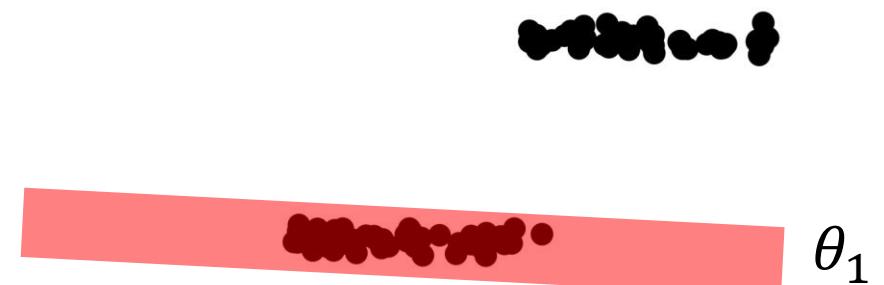


Line fitting example

# Sequential RanSaC [Zuliani 05]

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers



Line fitting example

# Sequential RanSaC [Zuliani 05]

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers from  $X$



Line fitting example

# Sequential RanSaC [Zuliani 05]

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers from  $X$

Iterate through the remaining points



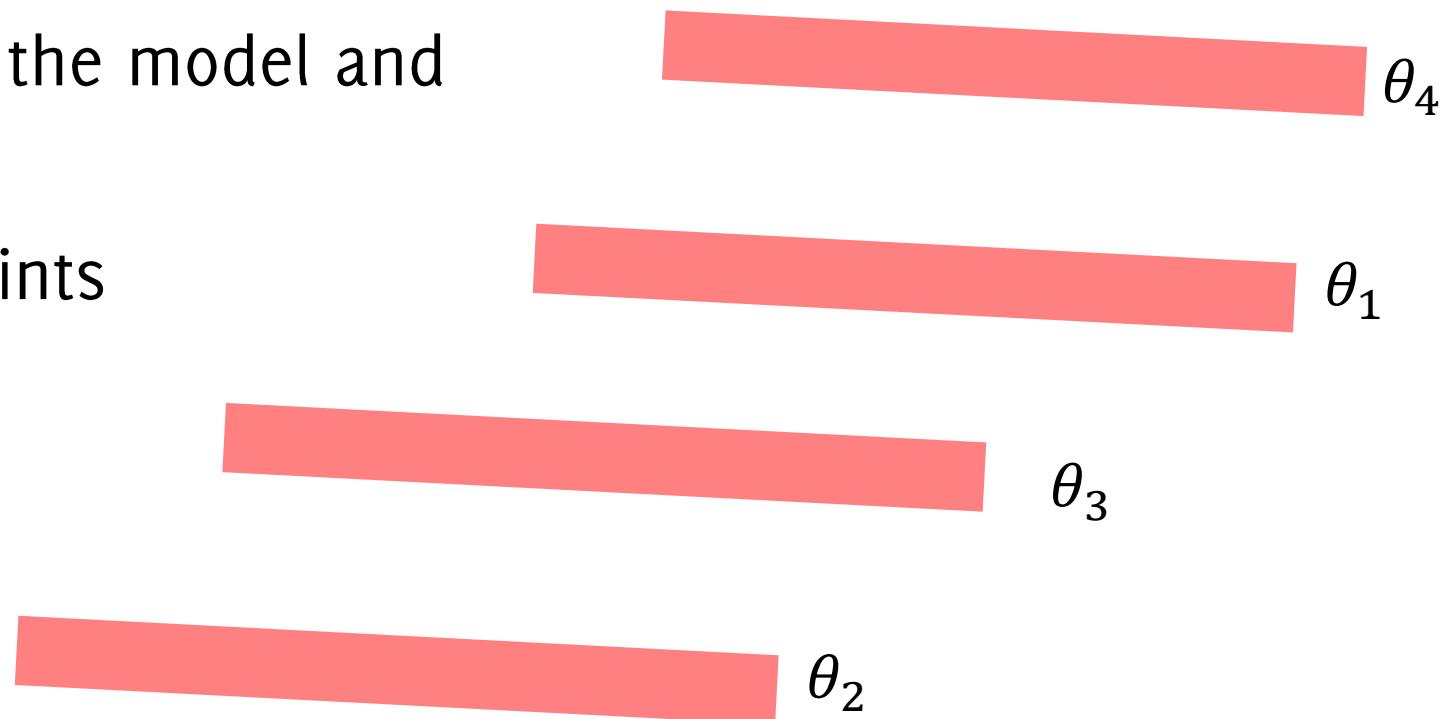
Line fitting example

# Sequential RanSaC [Zuliani 05]

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers from  $X$

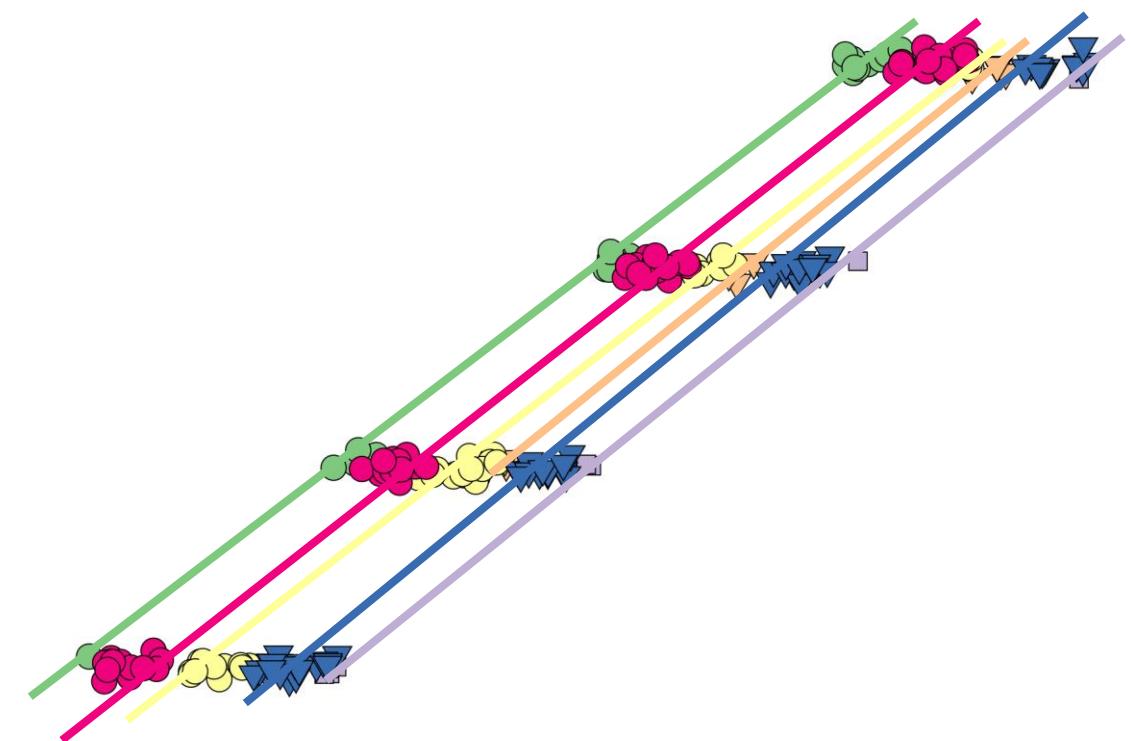
Iterate through the remaining points until there are no models with a sufficiently large consensus



Line fitting example

# Sequential RanSaC [Zuliani 05]

Unfortunately, this does not fit well with the multi-model scenario and the problem becomes even more sever in presence of outliers



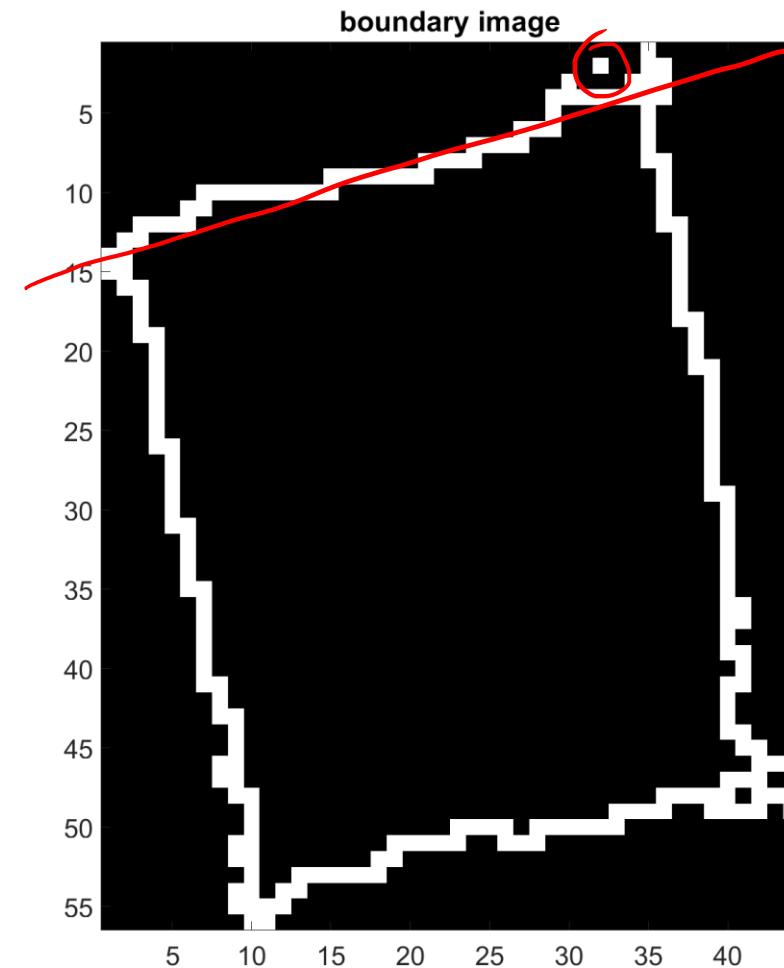
Line fitting example

# Line Detection: Hough Transform

## Extracting Line Equations From Edges

# Line Detection: The problem

Finding all the lines passing through points in (a binary) image

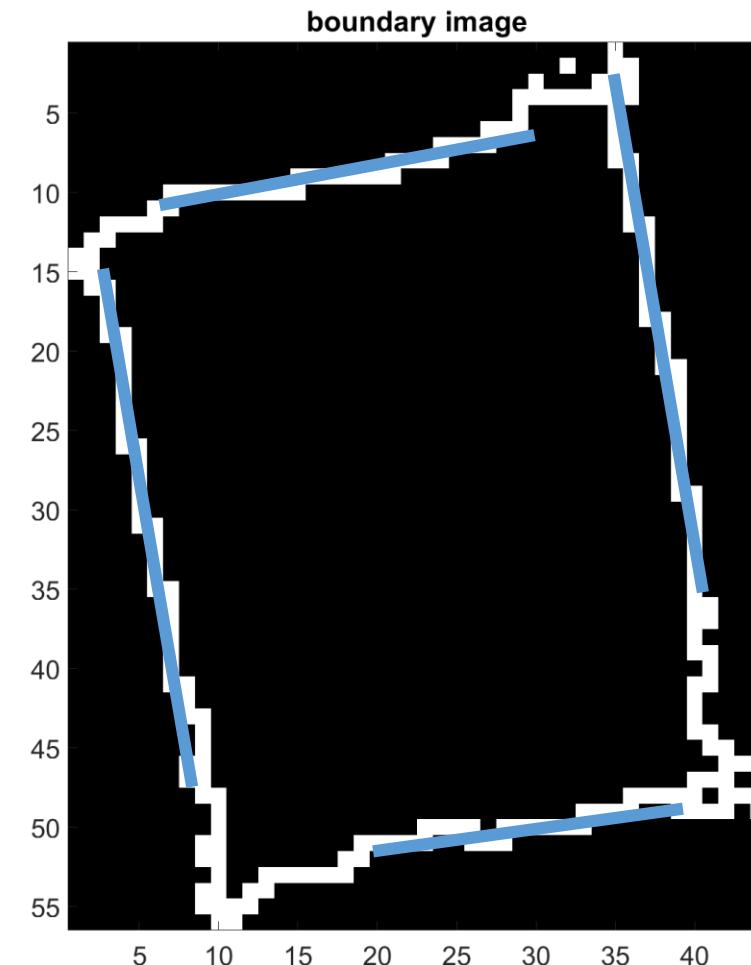


# Line Detection: The problem

Finding all the lines passing through points in (a binary) image

Finding lines means

- Having an analytical expression for each line
- Estimating its direction, length
- Thus, clustering points belonging to the same segment

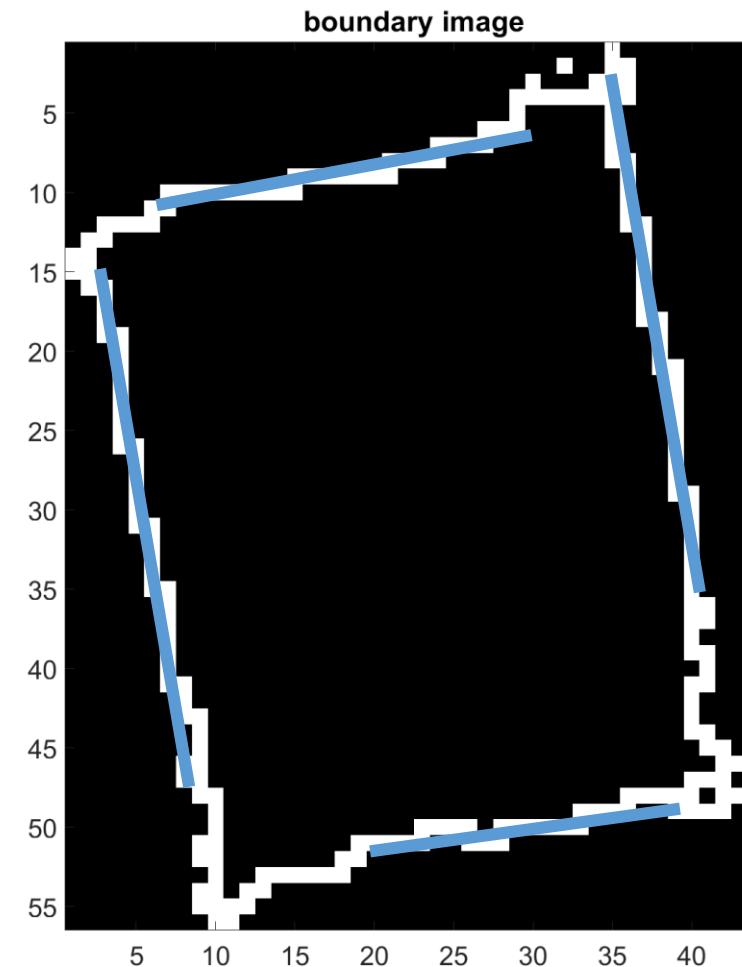


# Line Detection: The problem

Brut-force attempt:

Given  $n$  points in a binary image, find subsets that lie on straight lines

- Compute all the lines passing through any pair of points
- Check subsets of points that belong / are close to these lines

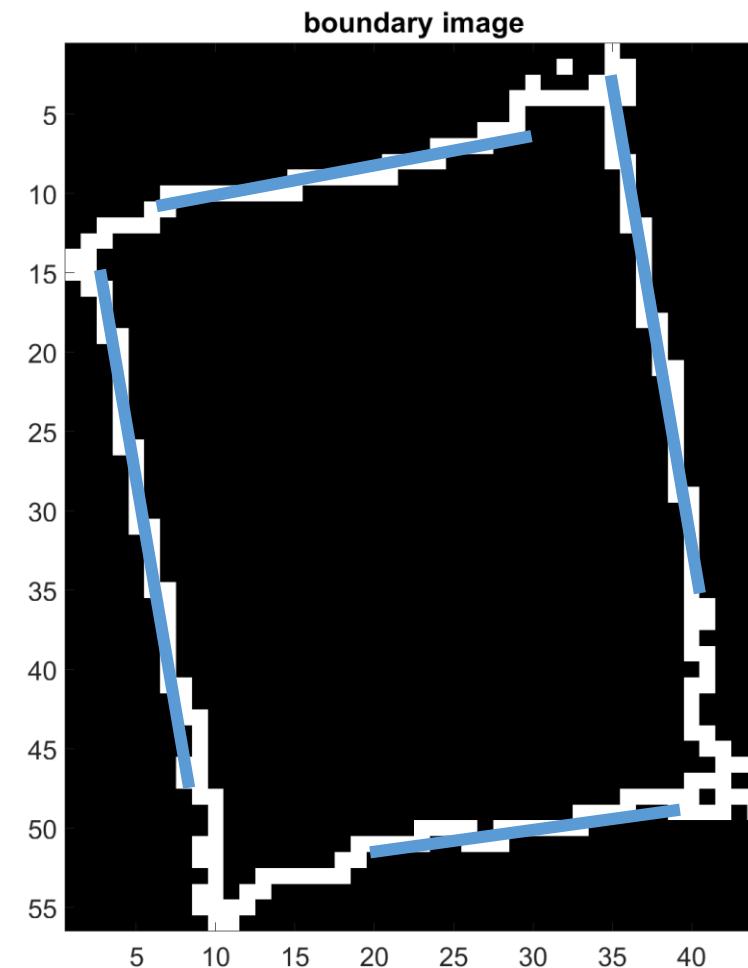


# Line Detection: The problem

Brut-force attempt:

This requires computing

- $\frac{n(n-1)}{2}$  straight lines
- $n \left( \frac{n(n-1)}{2} \right)$  comparisons
- Computationally prohibitive task  
in all but the most trivial  
applications  $\sim n^3$



# Hough Transform

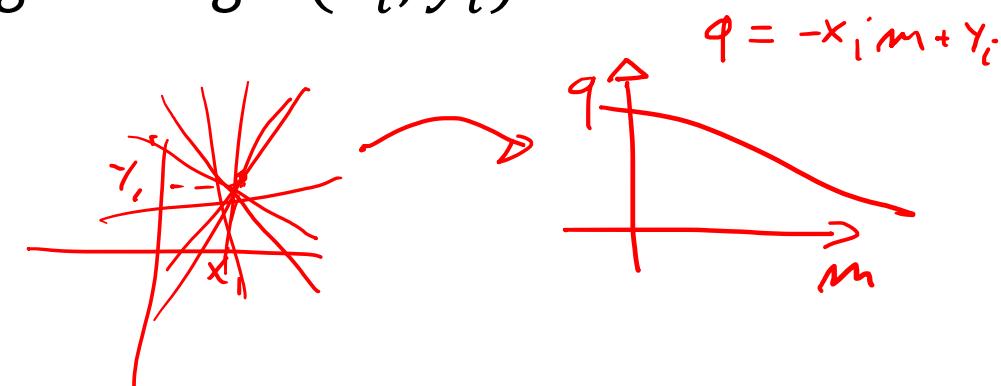
Identify lines in the “*parameter space*” i.e. in the space of the parameters identifying lines  $(m, q)$ . Let a straight line be:

$$y = mx + q \quad y_i = \underline{m} x_i + \underline{q}$$

Now, for a given point  $(x_i, y_i)$ , the equation  $q = -x_i m + y_i$  in the variables  $m, q$  denotes the star of lines passing through  $(x_i, y_i)$

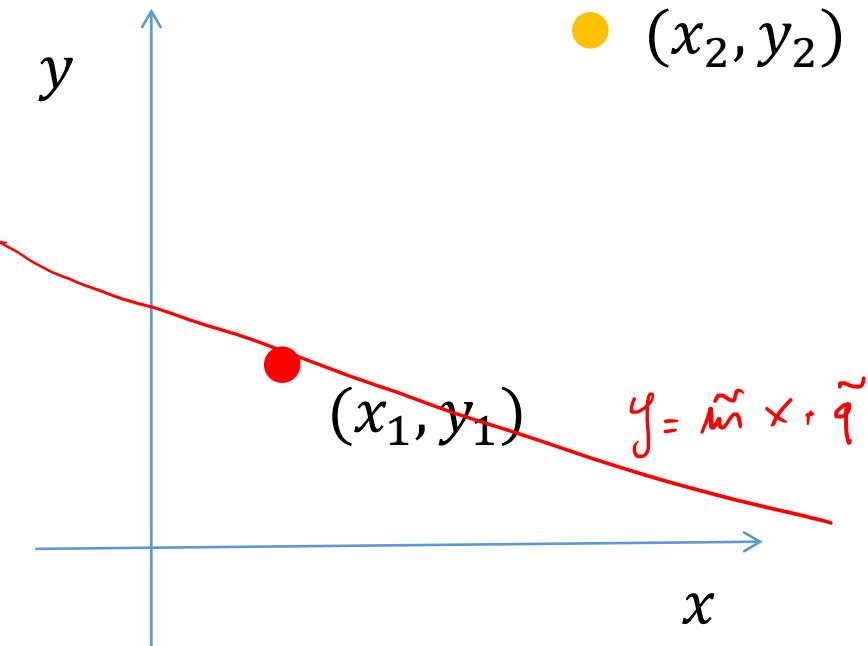
Key intuition:

$$q = -x_i m + y_i$$

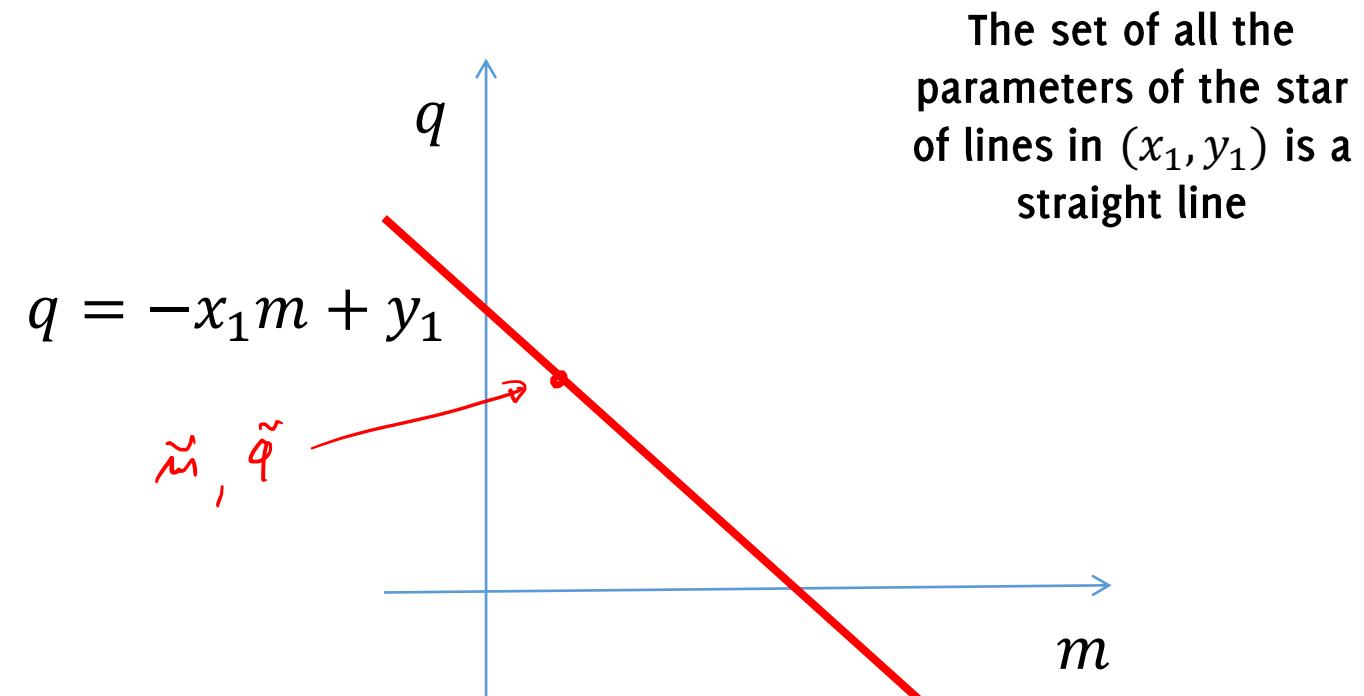


Can be also seen as the equation of a straight line in  $m, q$  in the parameter space

# Line Intersections in the parameter space



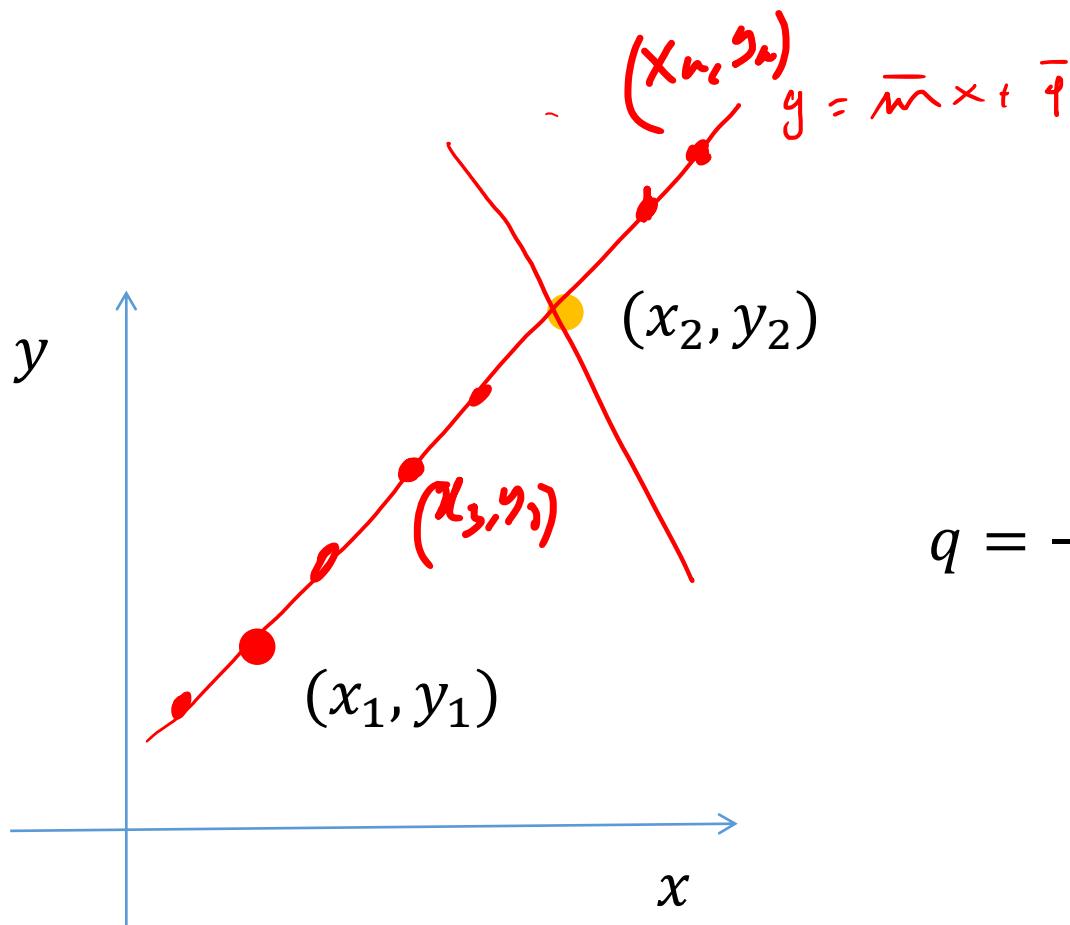
Point space



Parameter space

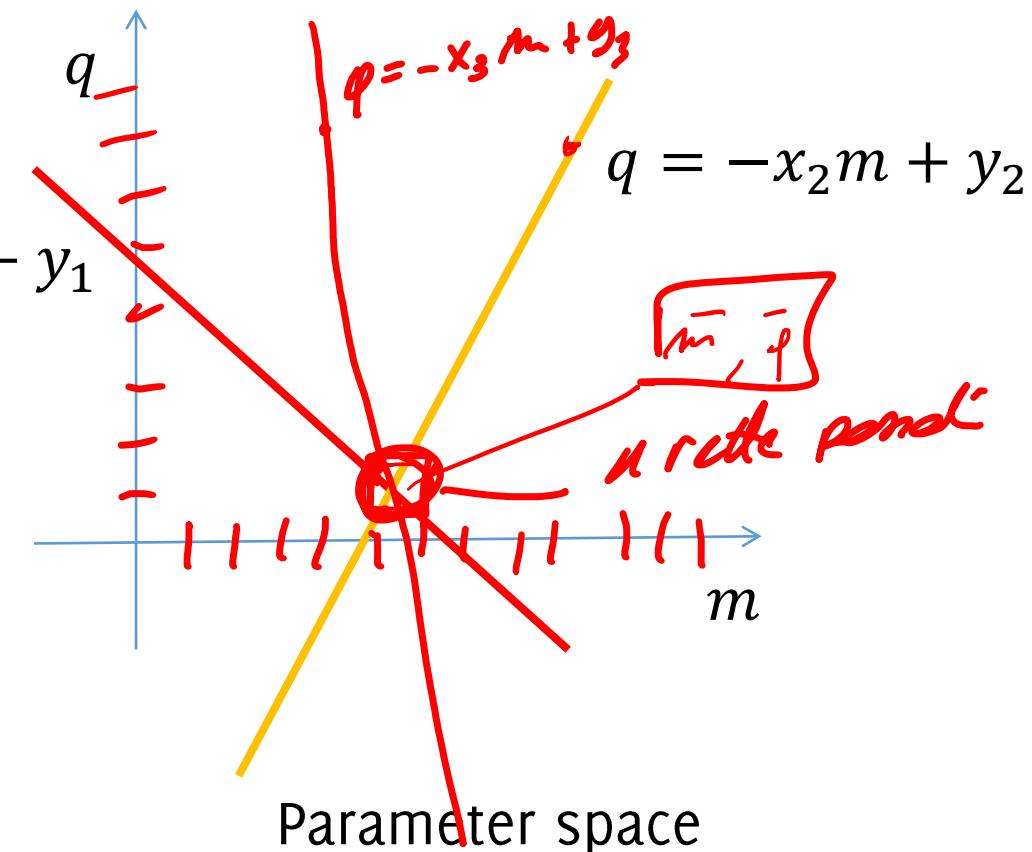
The set of all the parameters of the star of lines in  $(x_1, y_1)$  is a straight line

# Line Intersections in the parameter space

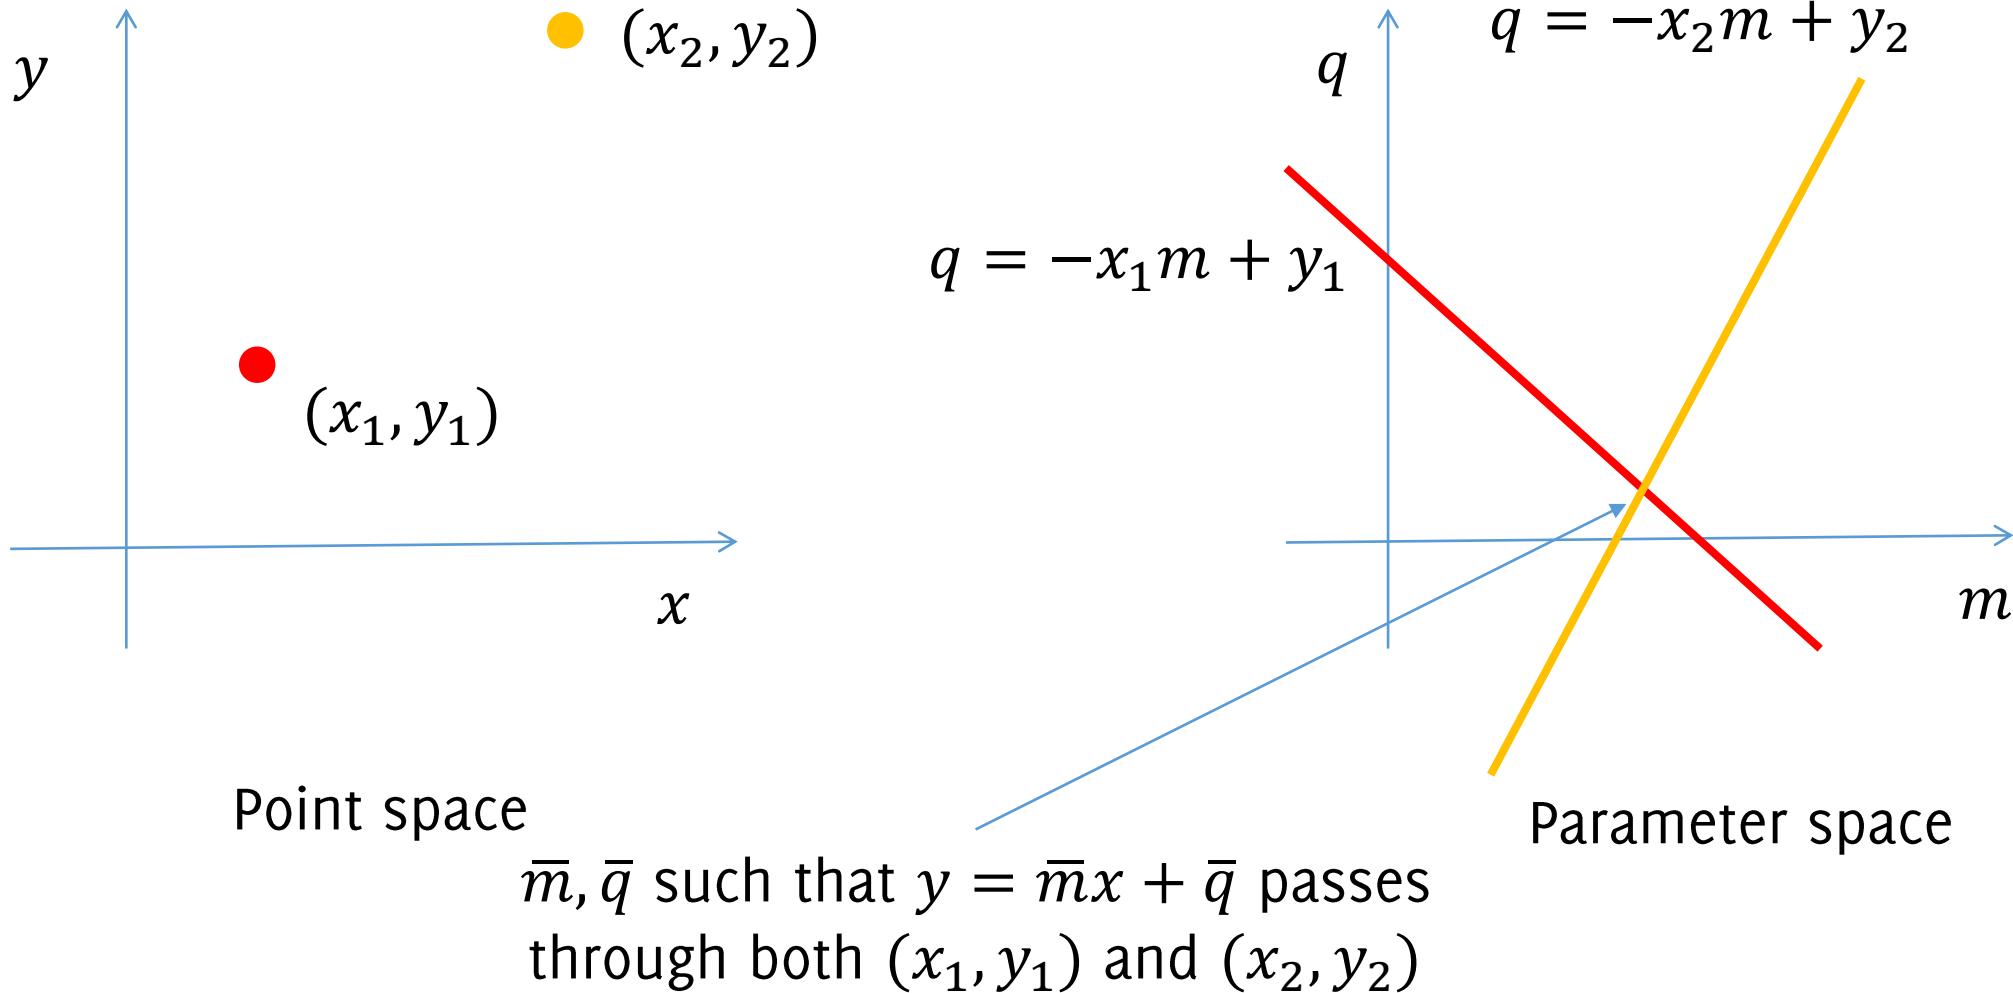


Point space

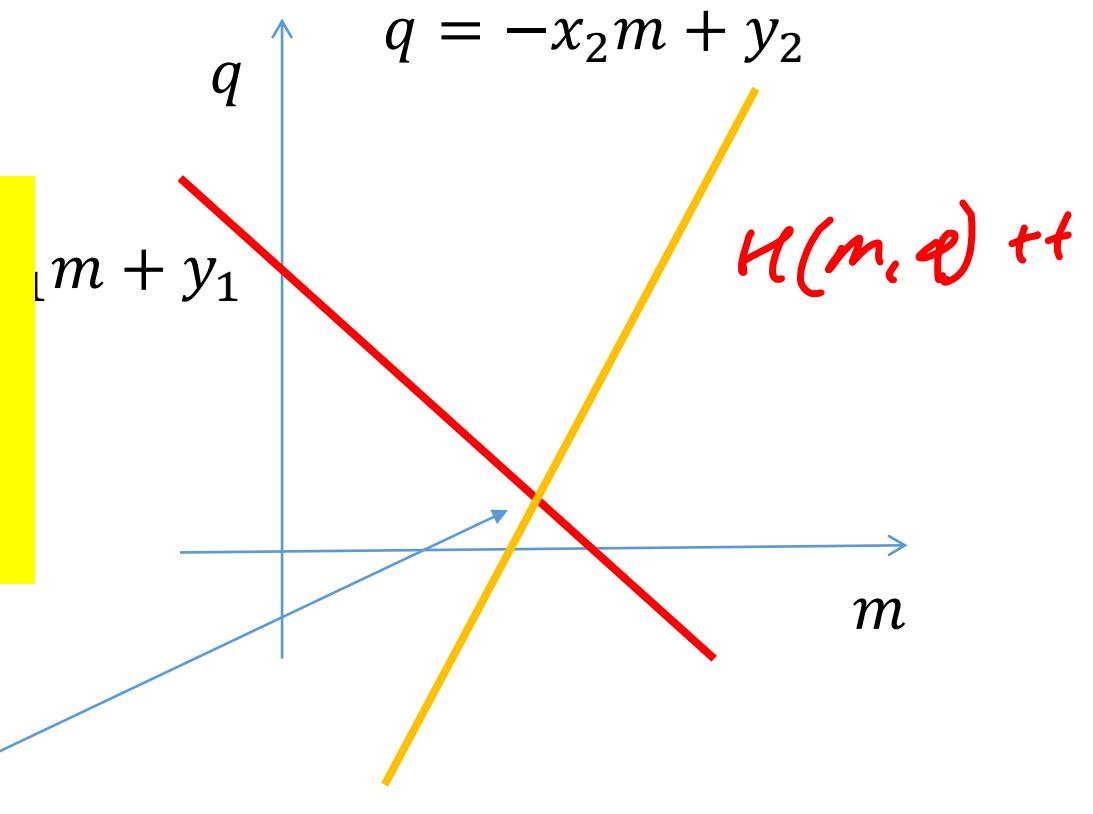
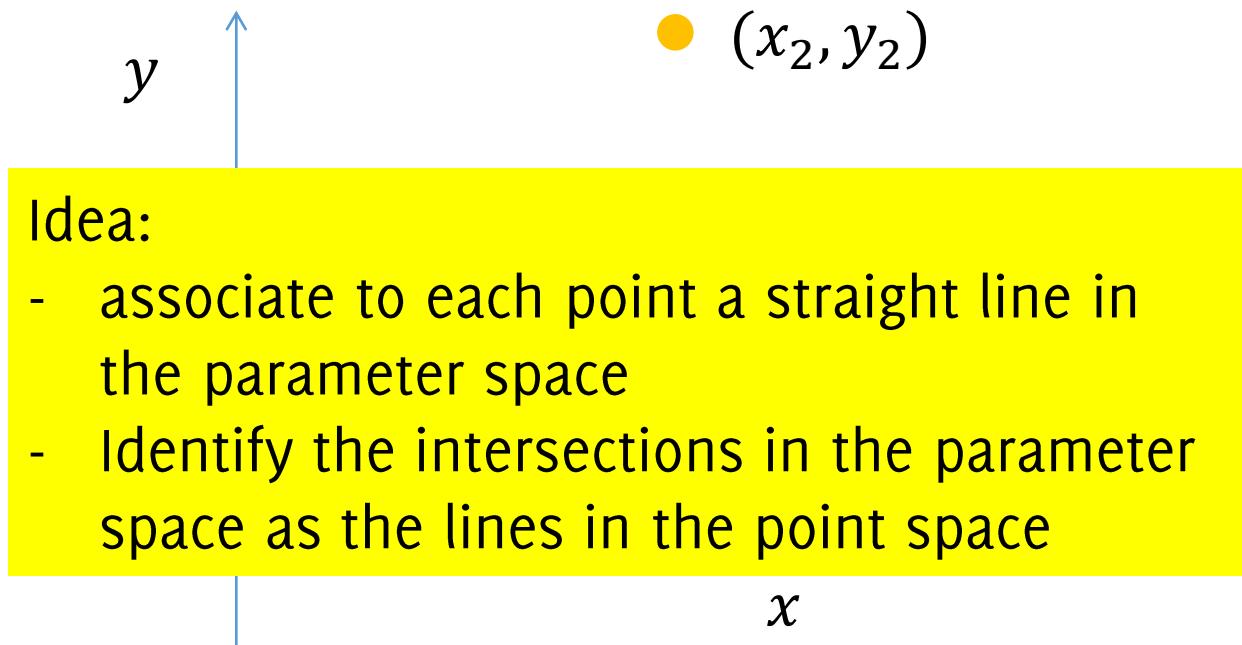
The two straight lines in the parameter space intersect in a point, corresponding to a line passing to both  $(x_1, y_1)$  and  $(x_2, y_2)$



# Line Intersections in the parameter space

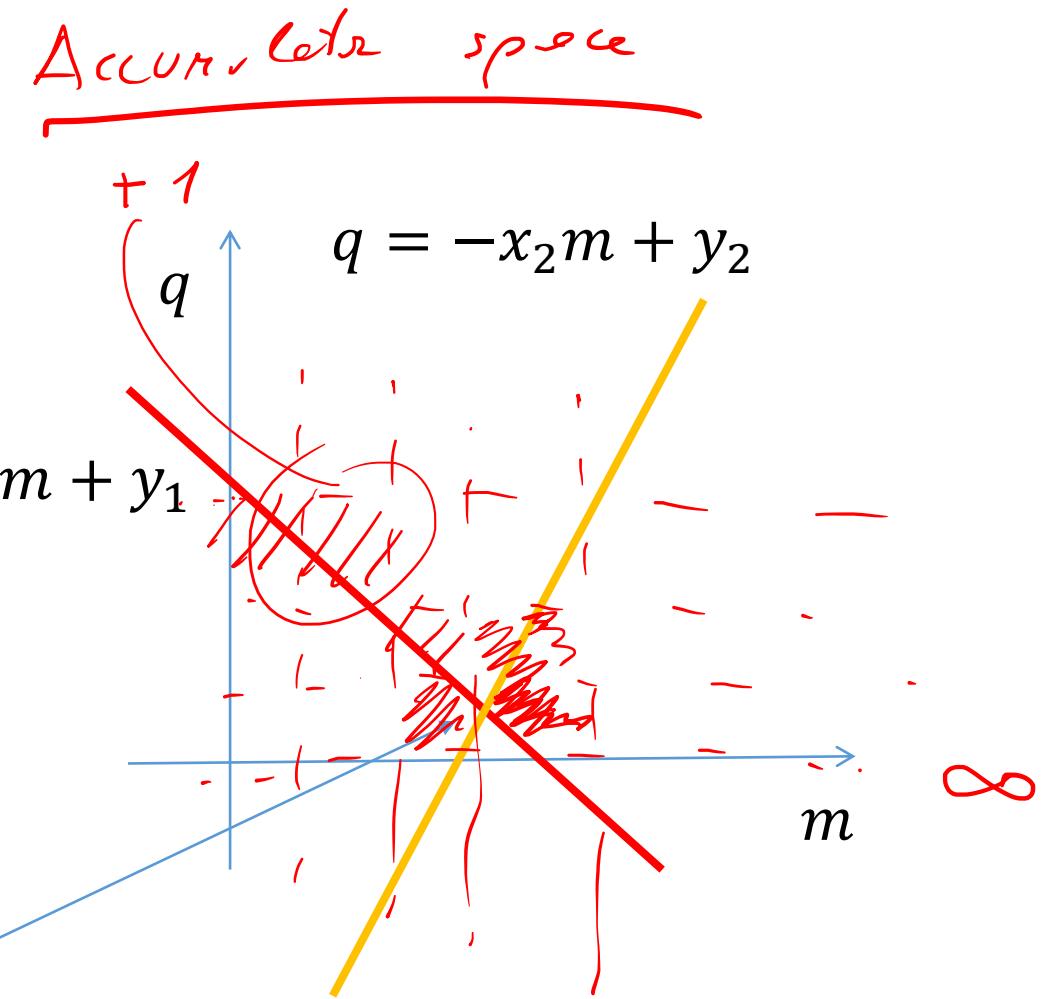
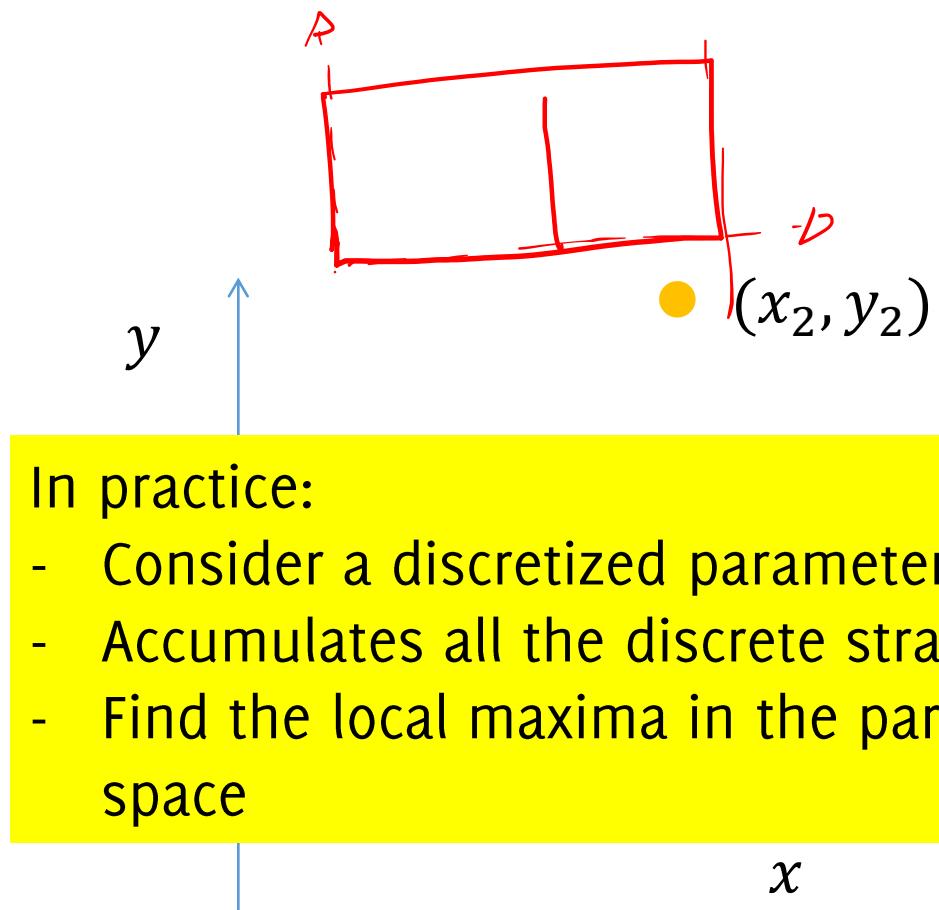


# Intersections in the parameter space



$\bar{m}, \bar{q}$  such that  $y = \bar{m}x + \bar{q}$  passes through both  $(x_1, y_1)$  and  $(x_2, y_2)$

# Intersections in the parameter space



$\bar{m}, \bar{q}$  such that  $y = \bar{m}x + \bar{q}$  passes through both  $(x_1, y_1)$  and  $(x_2, y_2)$

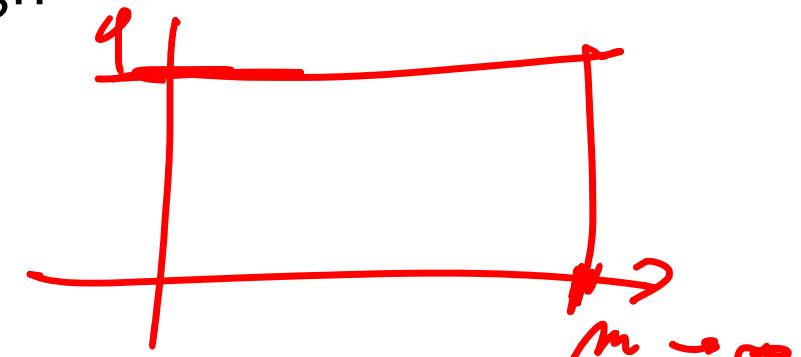
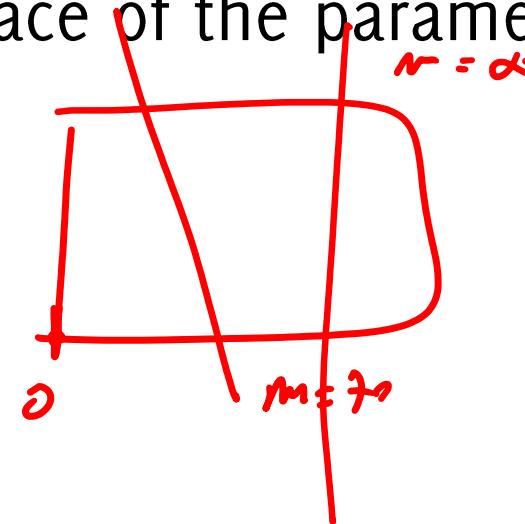
# Hough Transform

Identify lines in the “parameter space” i.e. in the space of the parameters identifying lines.

$$q = -x_i m + y_i, \quad \forall (x_i, y_i)$$

Core Idea:

- Discretize the parameter space where  $m, q$  live
- Accumulate the consensus in the parameter space by summing +1 at those bins where a straight line passess through
- Locate local maxima in the accumulator space



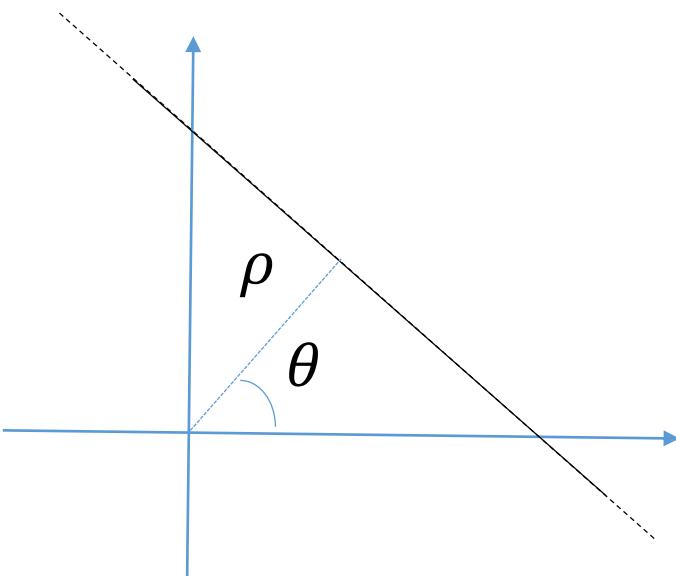
Major issue:  $m$  goes to infinity at vertical lines!

# New Parametrization for Hough Transform

There is a more convenient way of expressing a straight line, for this purpose:

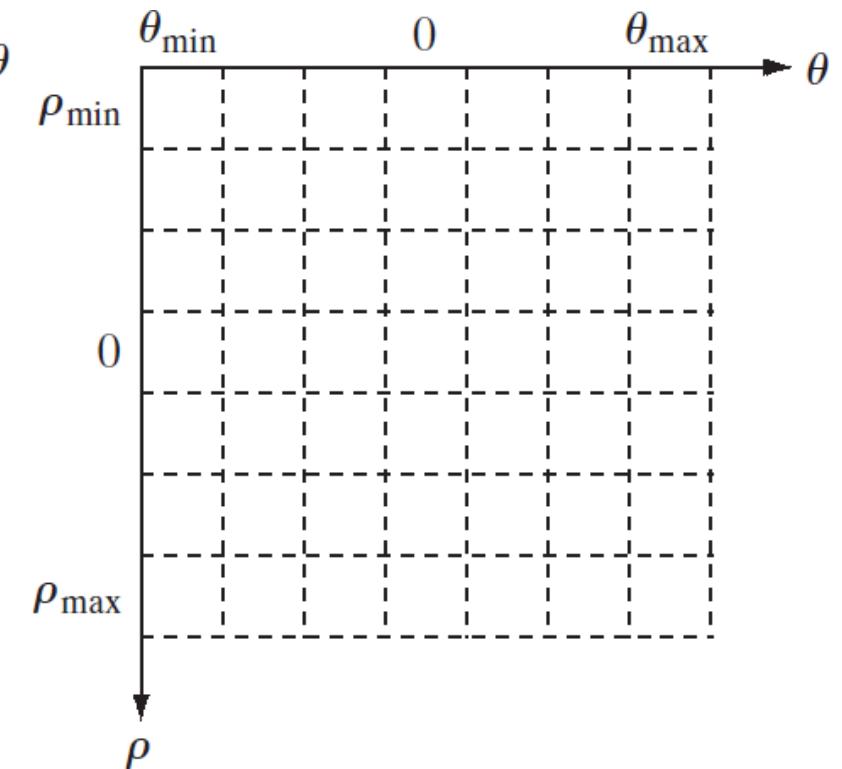
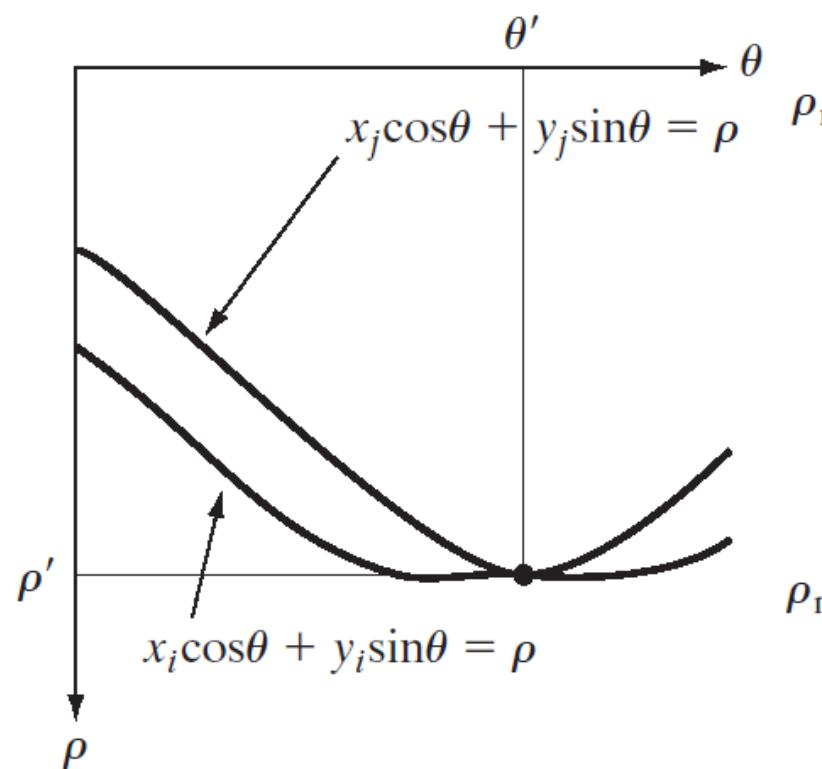
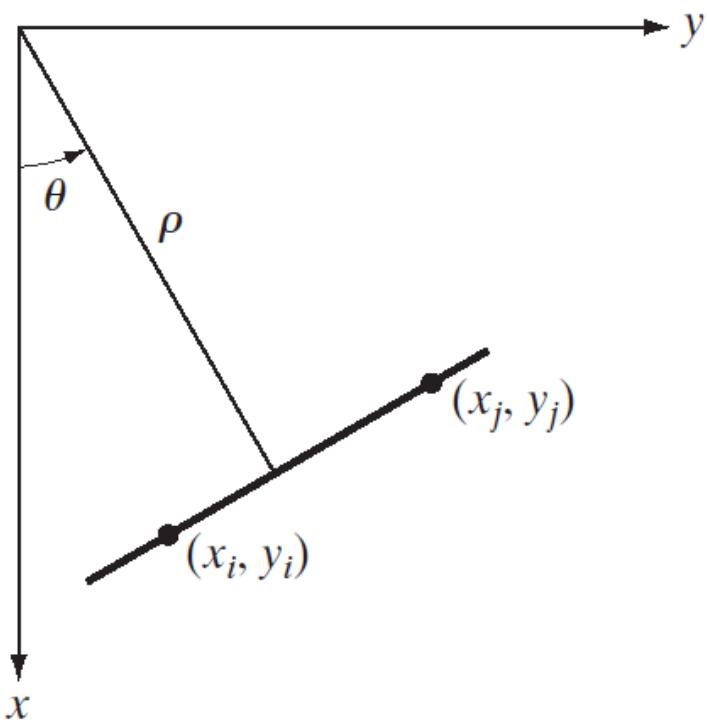
$$x \cos(\theta) + y \sin(\theta) = \rho$$

Where  $\{(\rho, \theta), \rho \in [-L, L], \theta \in \left[-\frac{\pi}{2}, \frac{\pi}{2}\right]\}$  being  $L$  the image diagonal



Same as before: a line in the image space is a point in parameter Hough space.

# New parametrization of straight lines



# Hough Transform

The Hough transform identifies **through an optimized voting procedure** the most represented lines

The voting procedure is performed in the «accumulator space» which is a grid in  $(\rho, \theta)$ -domain, for all the possible values.

From the Accumulator space we then extract local maxima, namely pairs  $(\rho, \theta)$  corresponding to lines passing through most of points

What is the maximum size of the domain?

# Hough Transform: the algorithm

Initialize  $H[\rho, \theta] = 0$

for each edge point  $(x, y)$  in the image:

    for  $\theta$  in range( $\theta_{\min}, \theta_{\max}$ ):

$\rho = x \cos(\theta) - y \sin(\theta)$

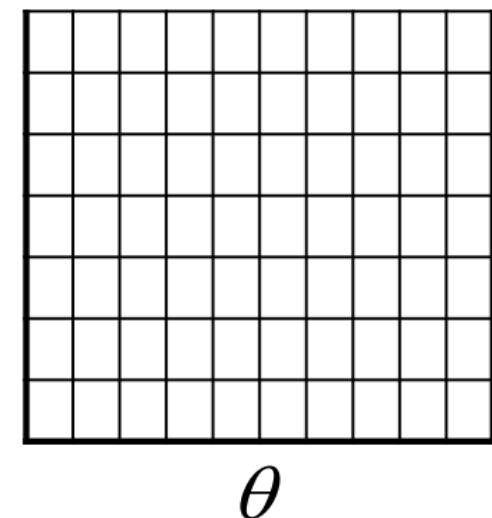
$H[\rho, \theta] += 1$

Find the value(s) of  $(d, \theta)$  where  $H[d, \theta]$  is maximum

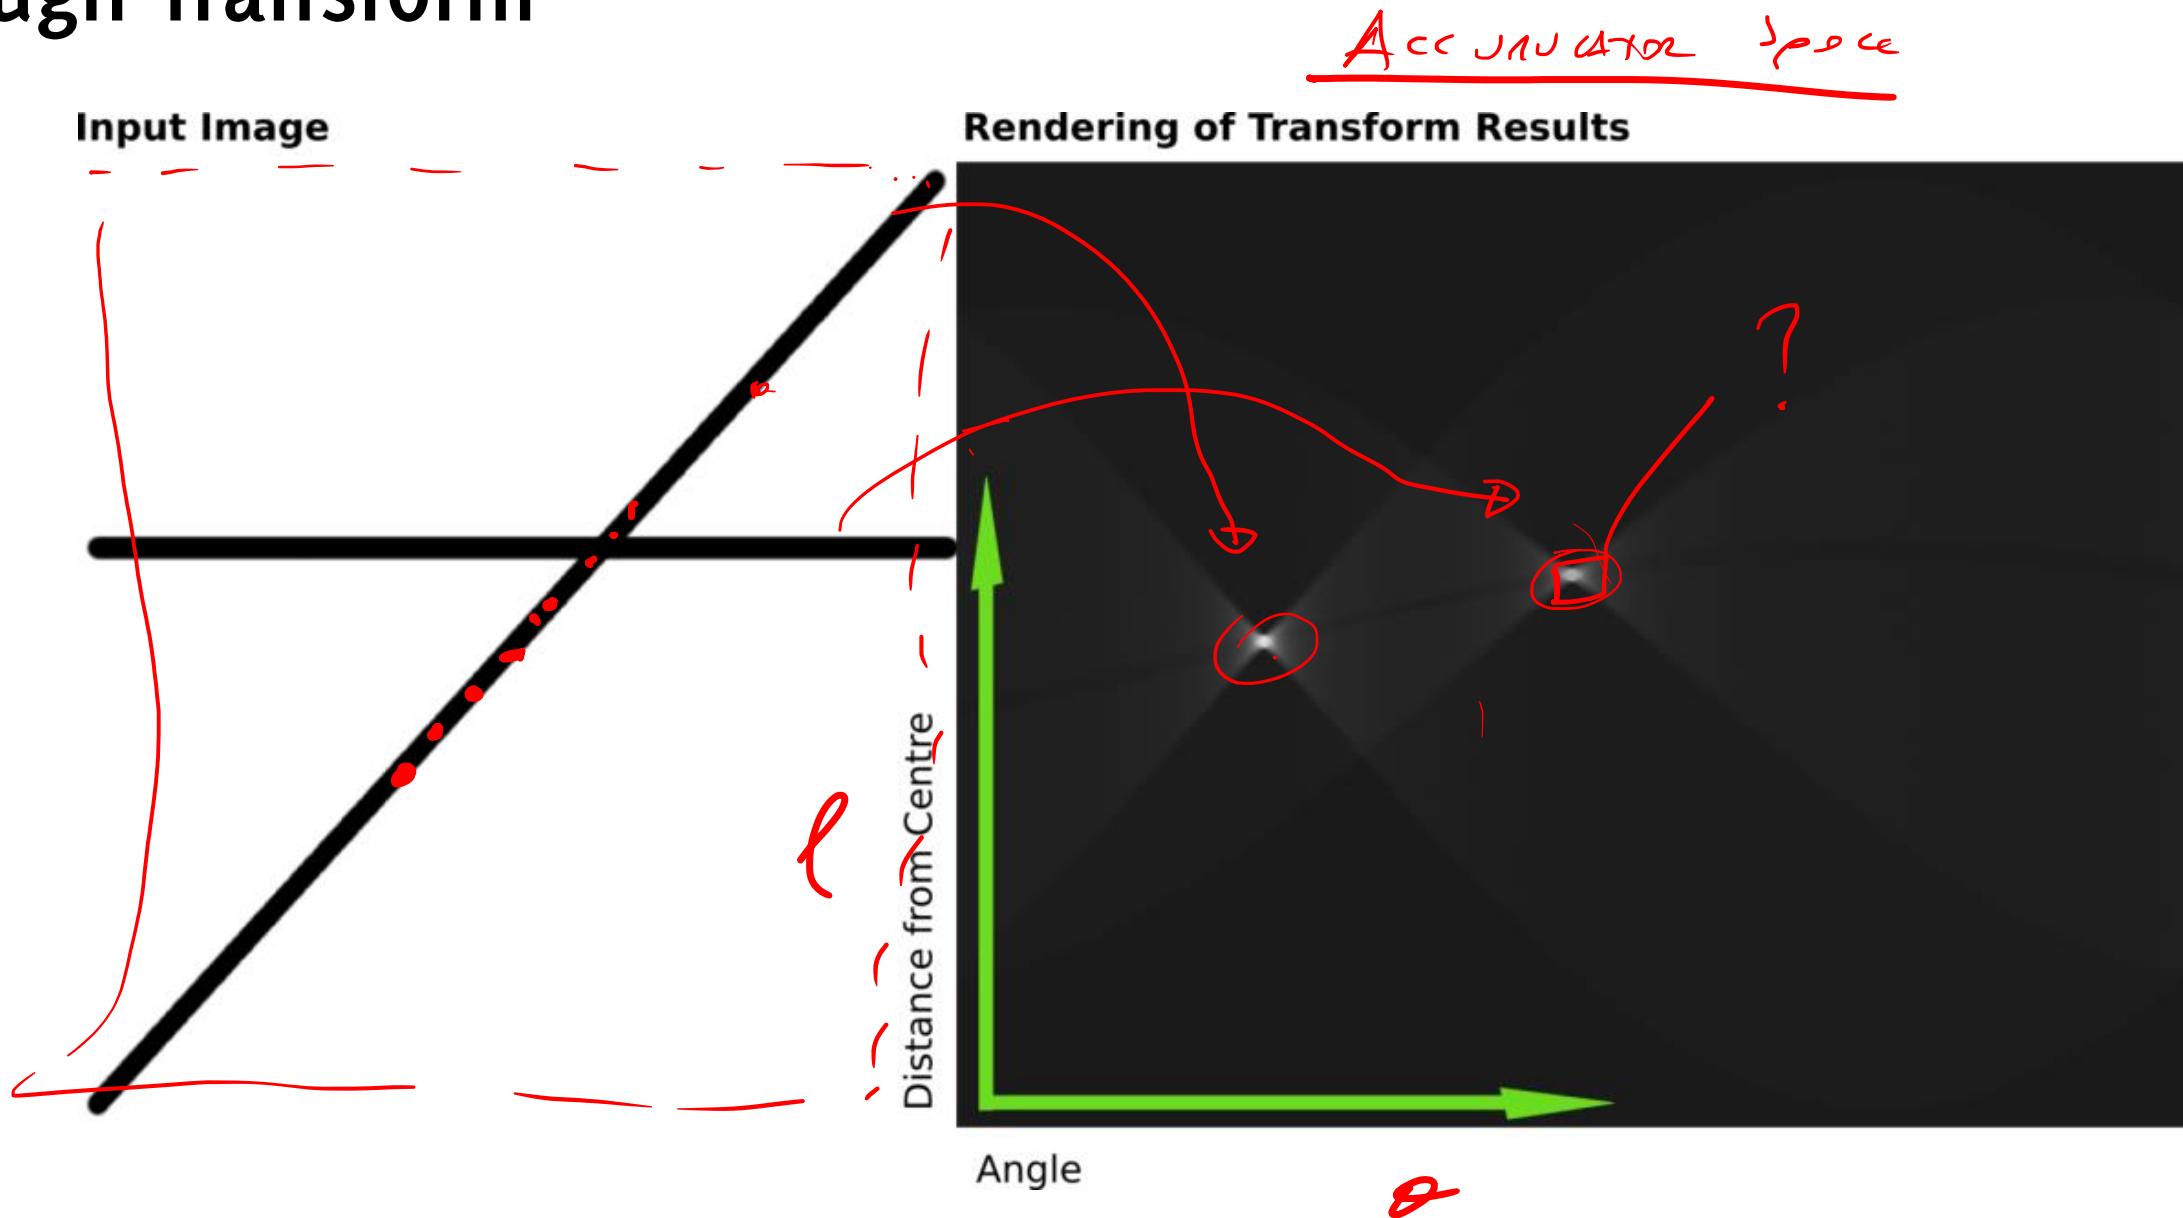
The detected line in the image is given by

$d = x \cos(\theta) - y \sin(\theta)$

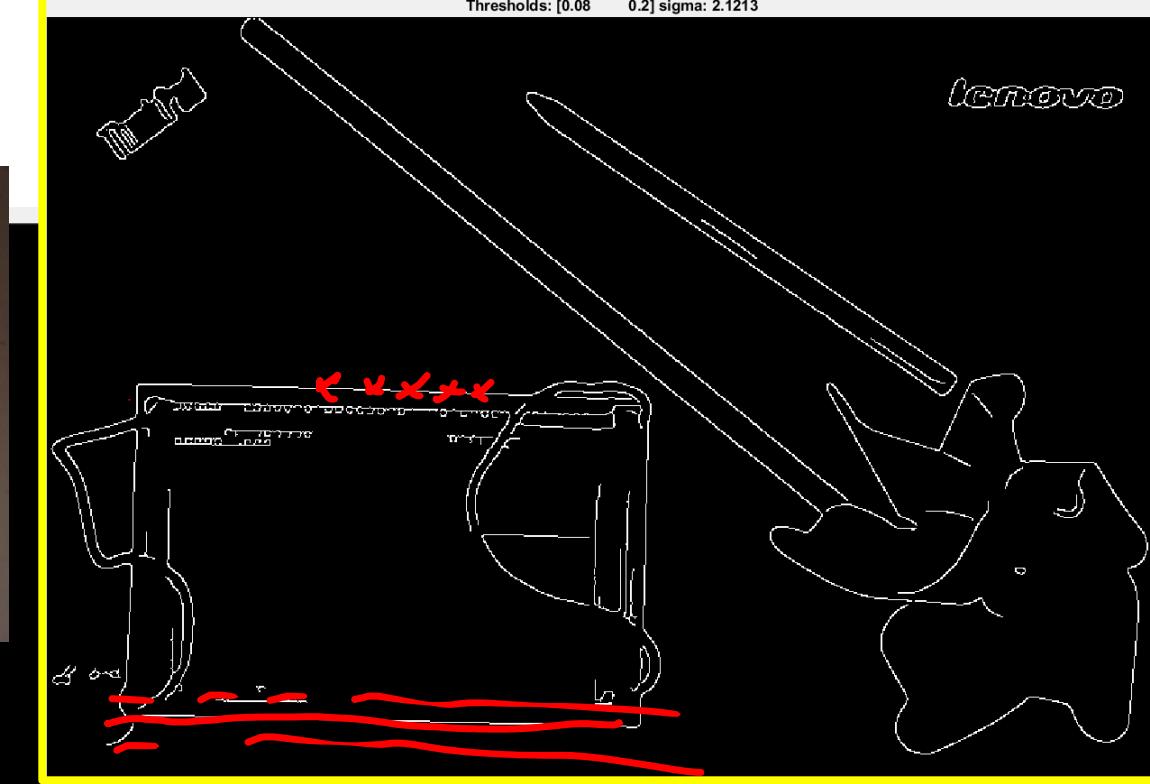
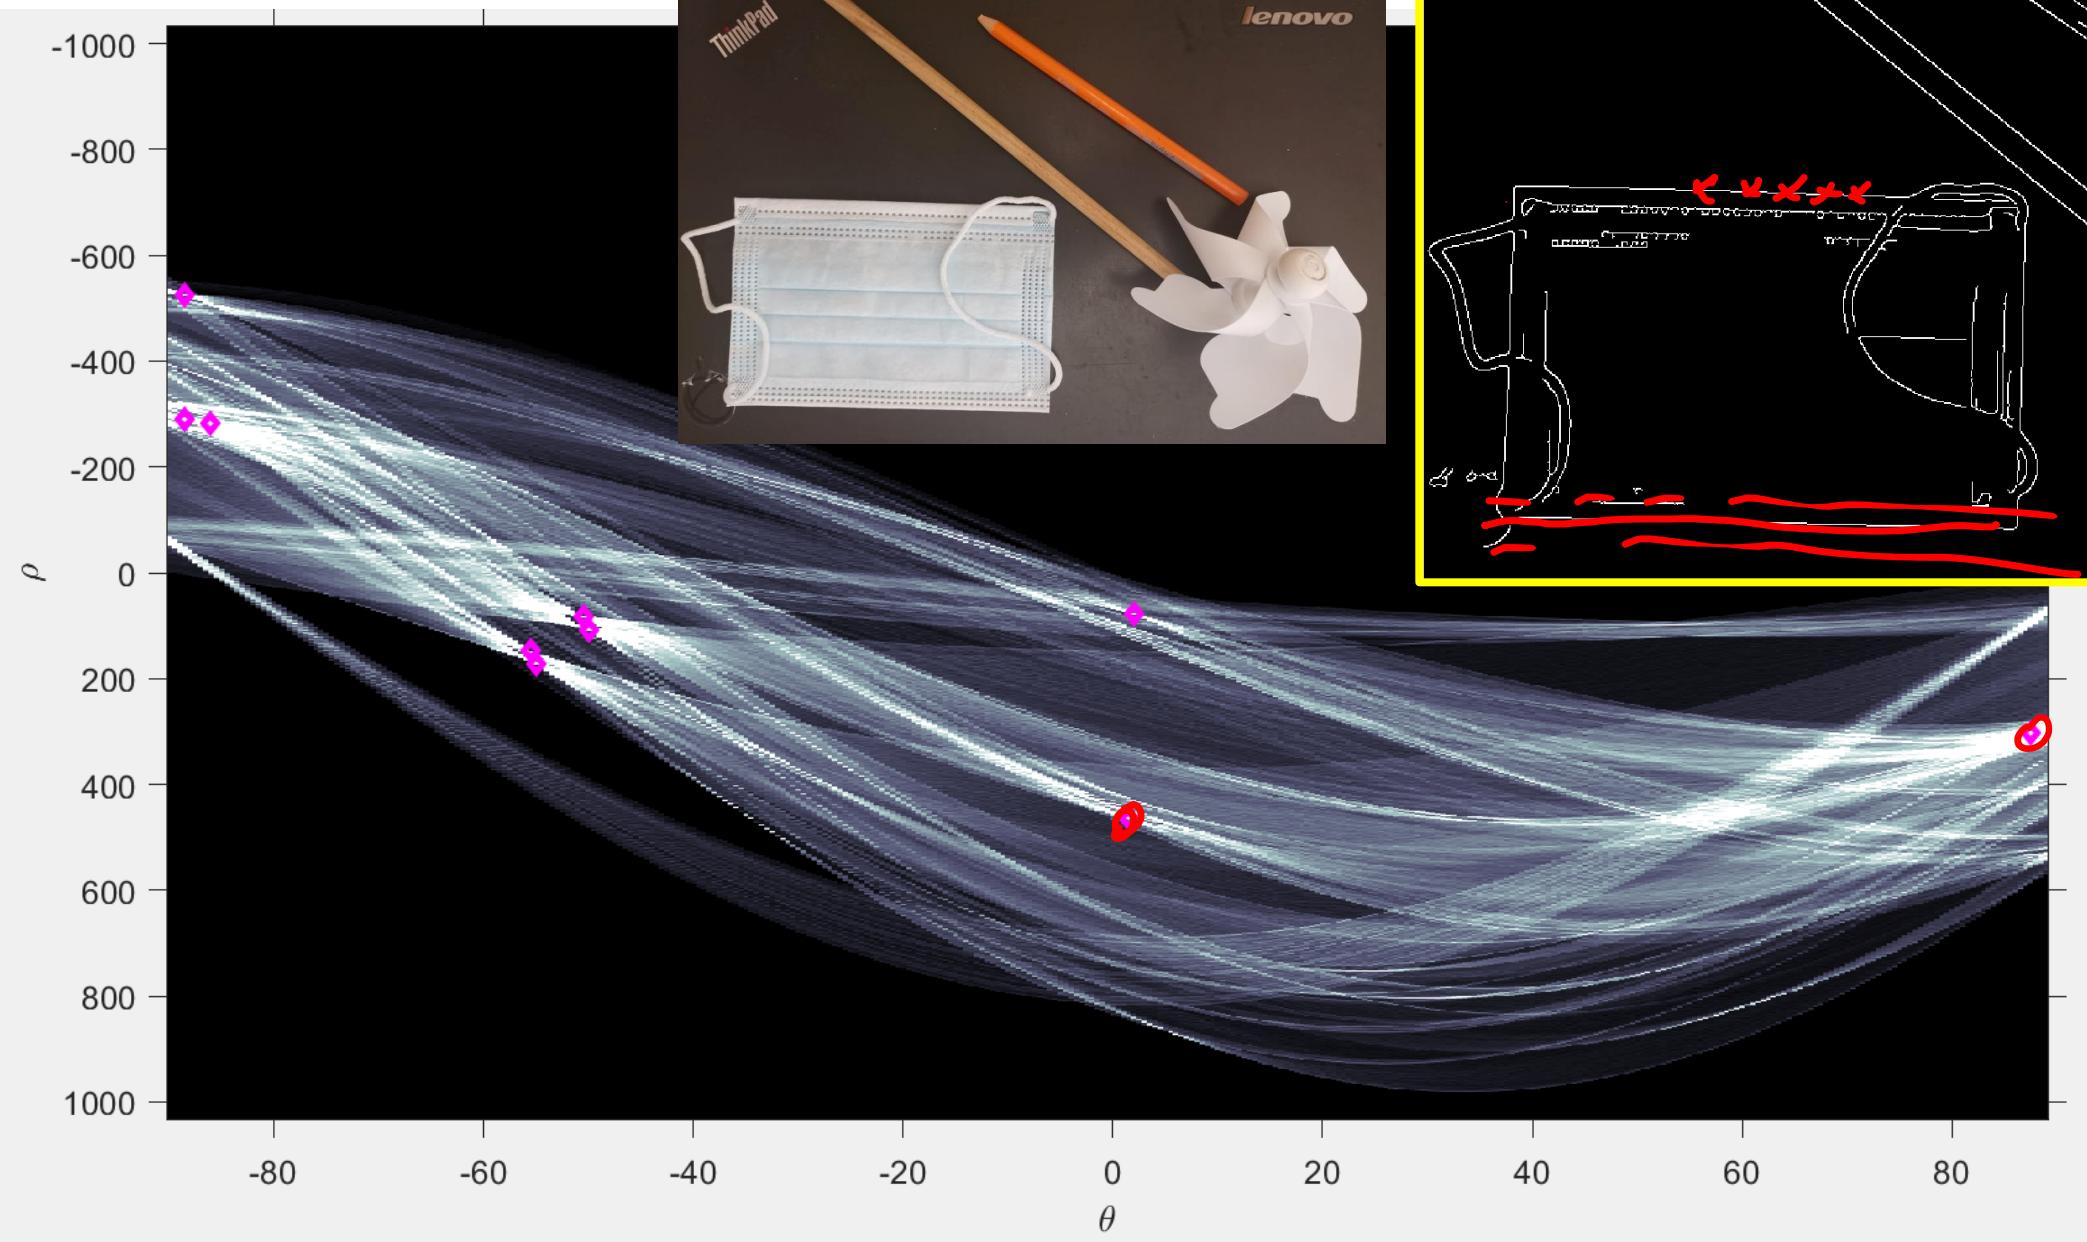
$H$ : accumulator array (votes)



# Hough Transform



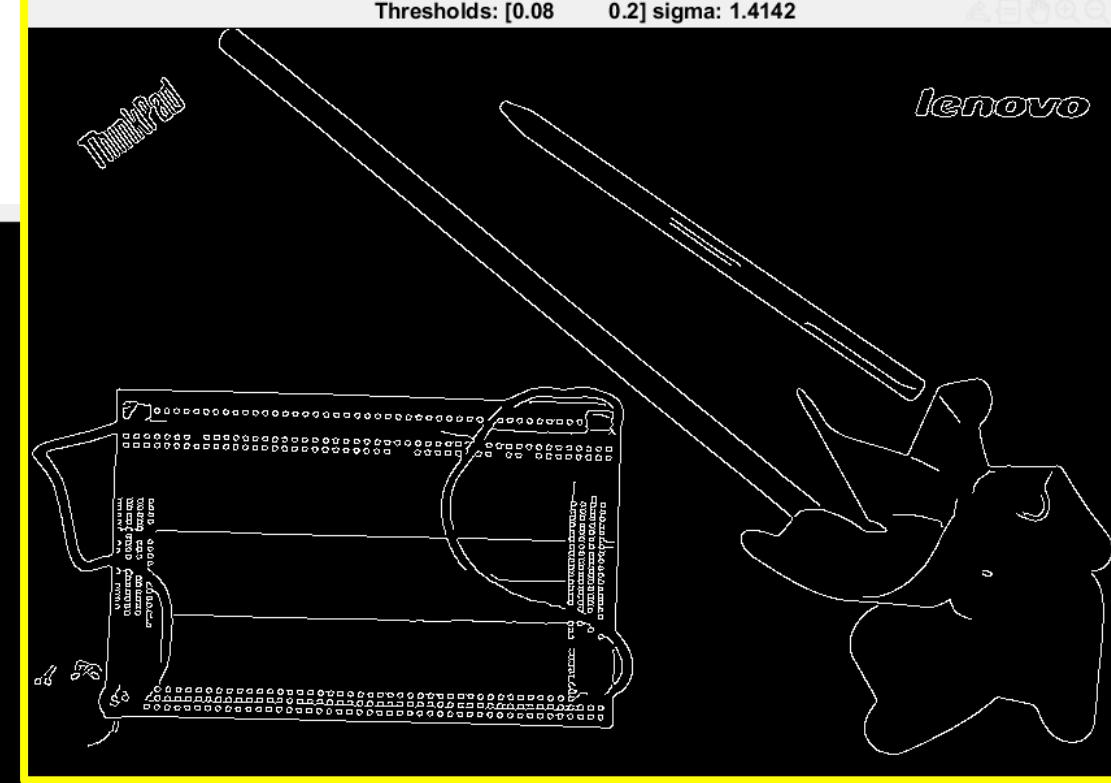
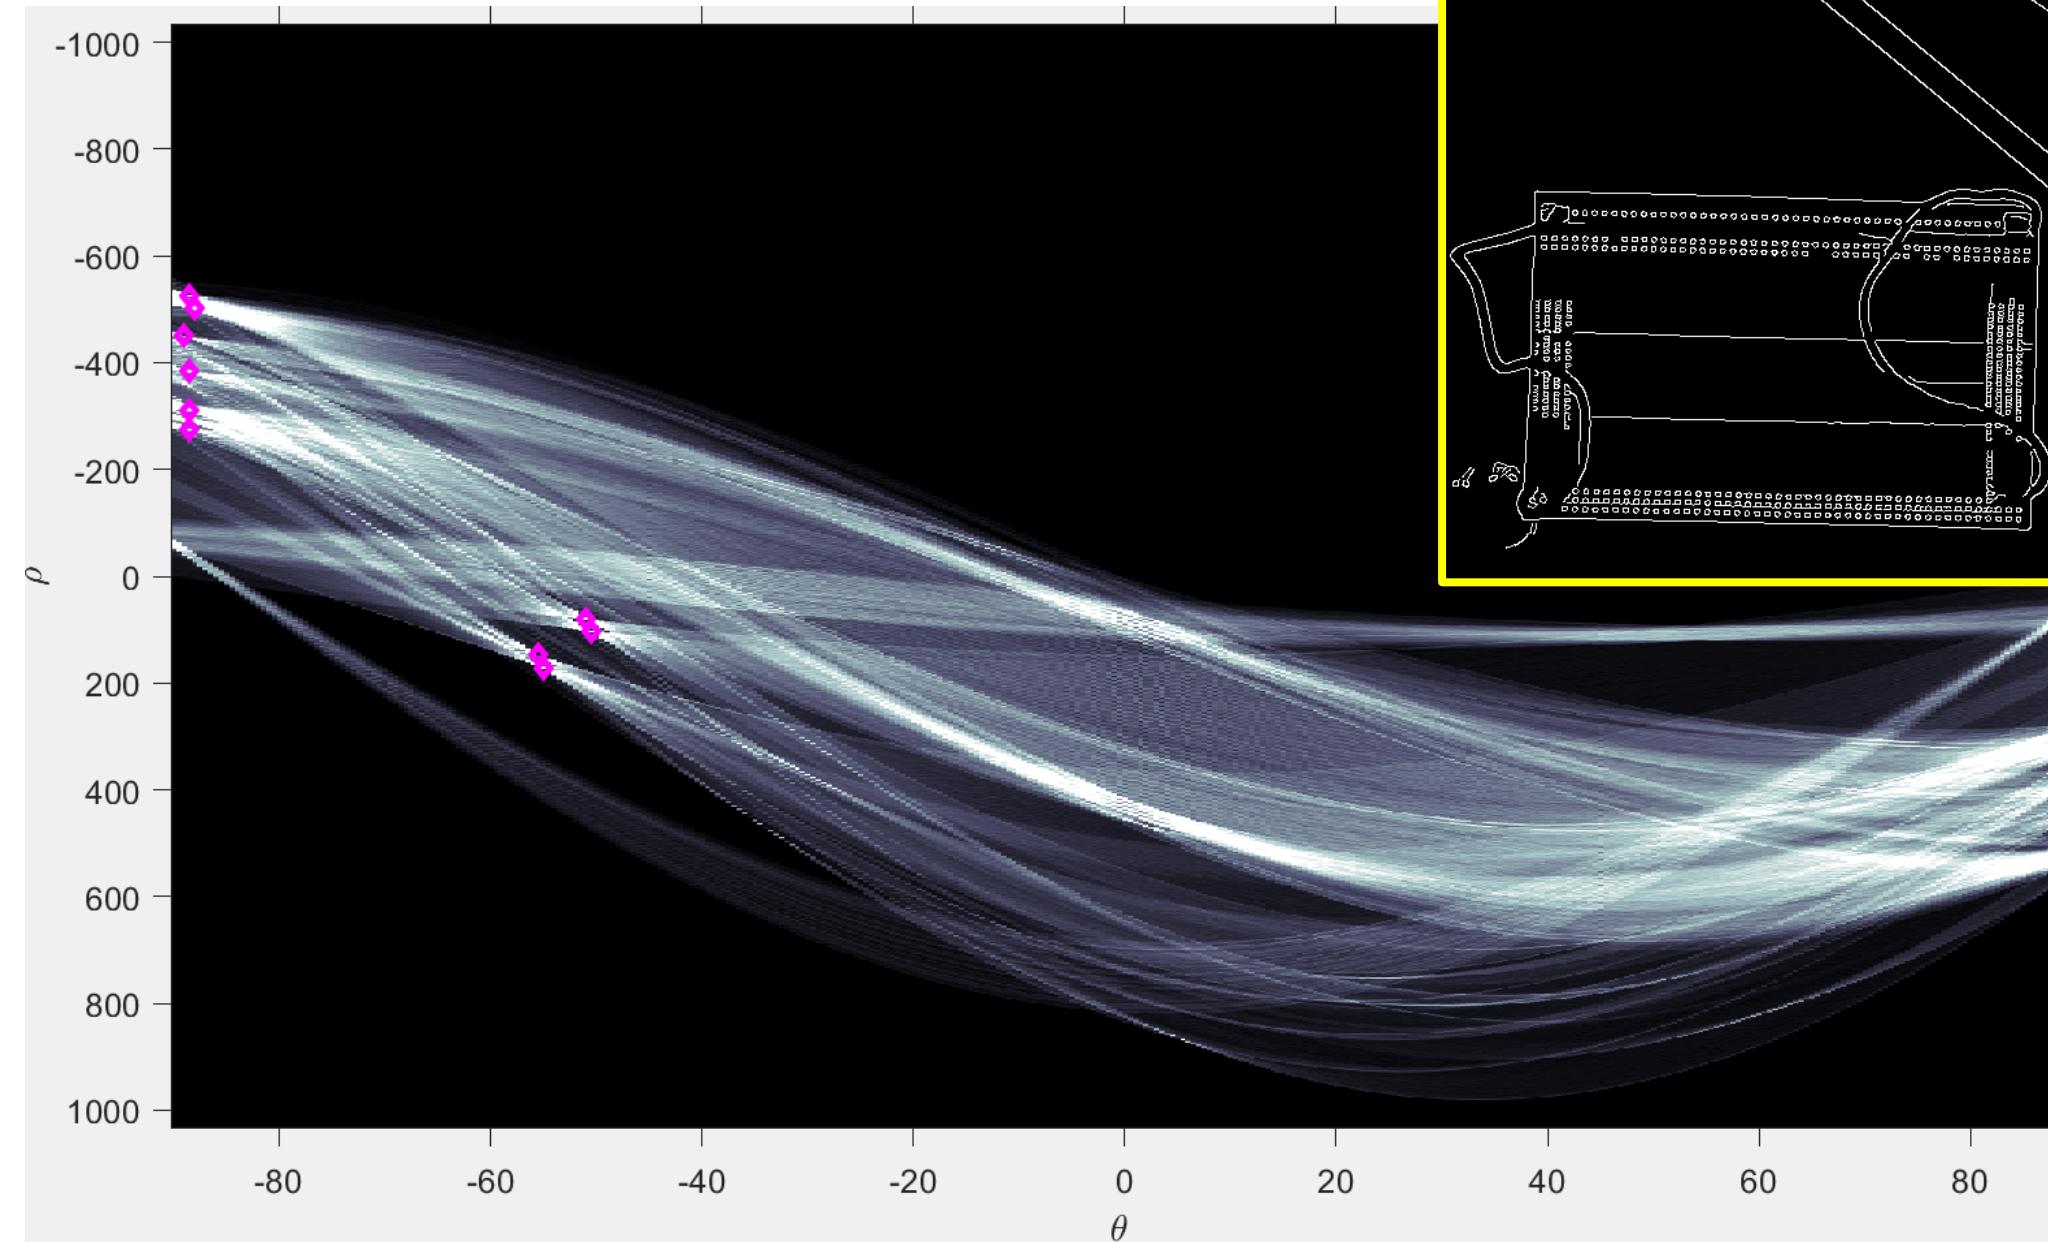
# Where is the facemask in H?



Thresholds: [0.08 0.2] sigma: 1.4142

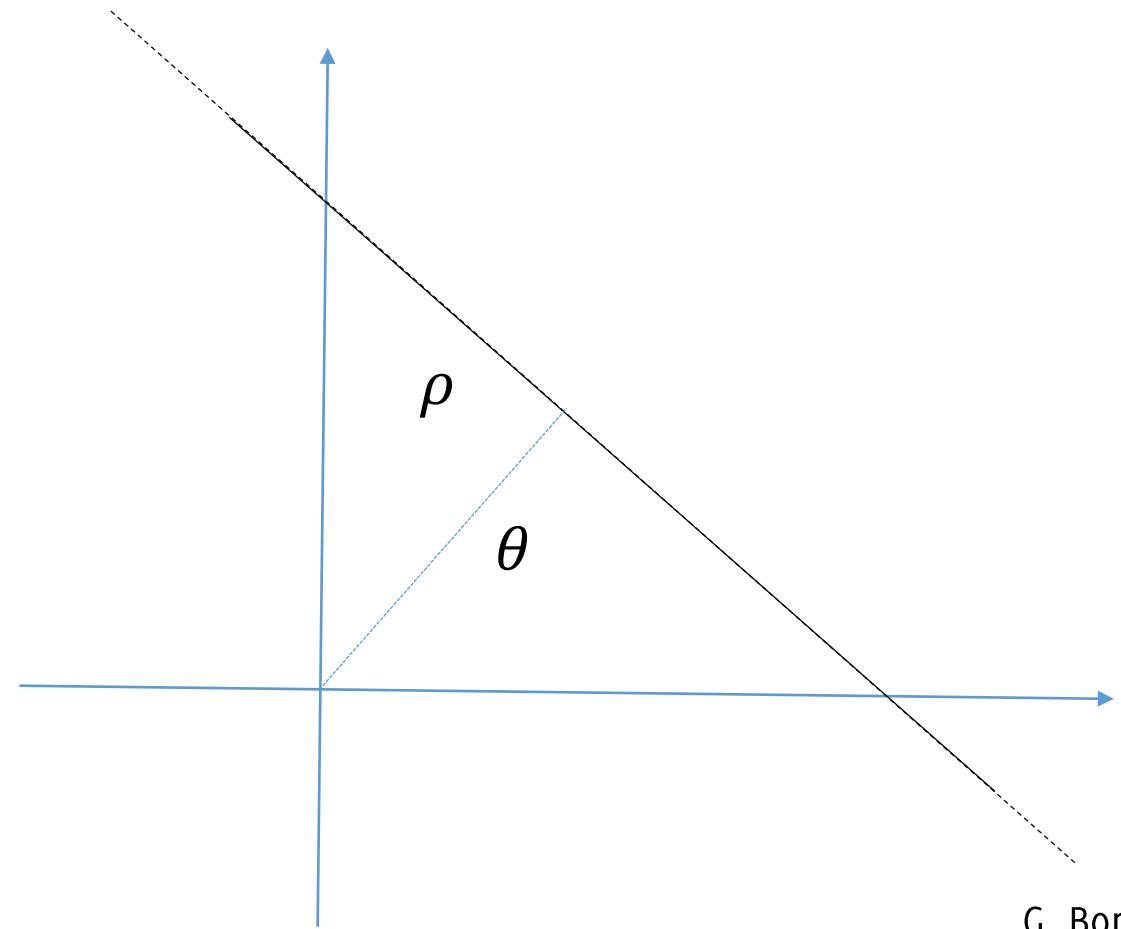
lenovo

# What if we take more edges?



# Size of the Accumulator Space

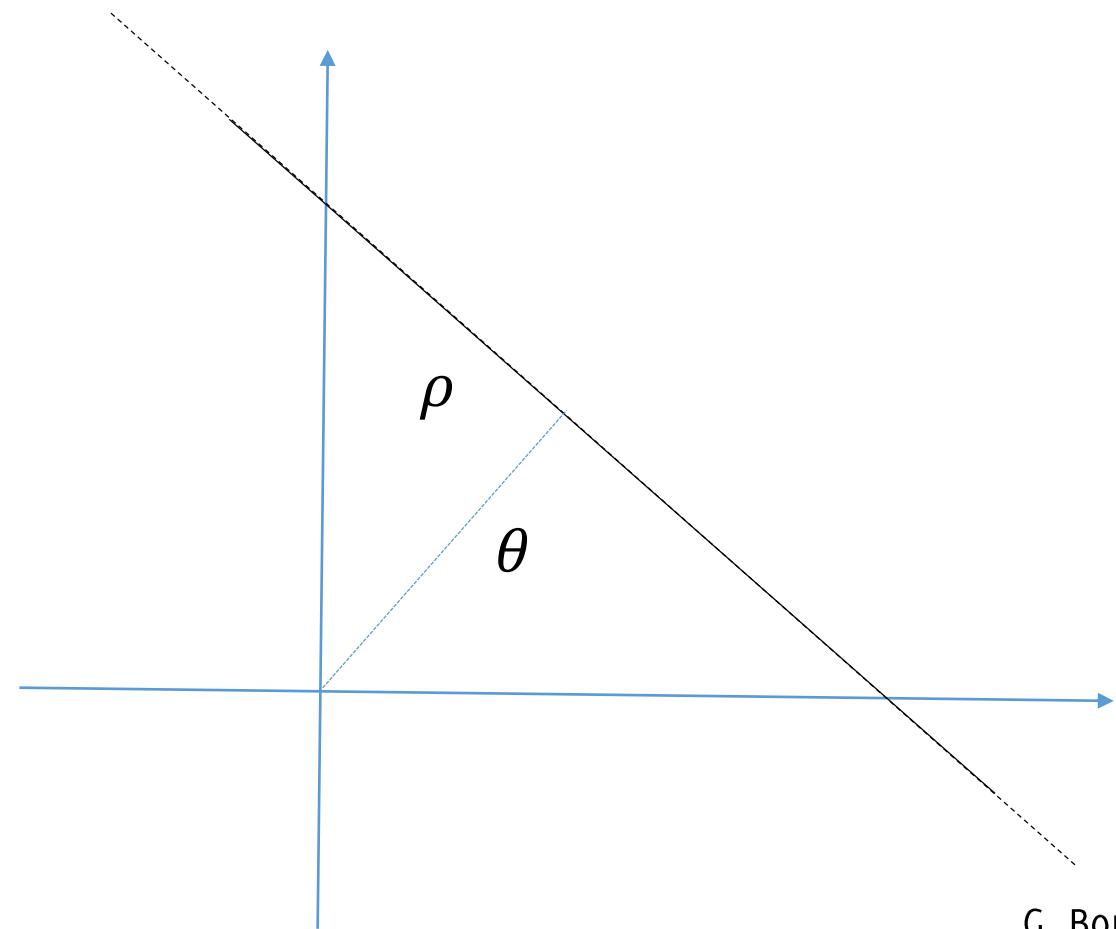
What are the maximum sizes of the accumulator space to represent any line intersecting the  $H \times W$  image?



# Size of the Accumulator Space

What are the maximum sizes of the accumulator space to represent any line intersecting the  $H \times W$  image?

It is the diagonal, so  $\sqrt{H^2 + W^2}$



# Bin size in the accumulator: an important parameter

How large are the bins in the accumulator?

- Too small: many weak peaks due to noise
- Just right: one strong peak per line, despite noise
- Too large:
  - poor accuracy in locating the line
  - many votes from clutter might end up in the same bin

A solution:

- keep bin size small but also vote for neighbors in the accumulator  
(this is the same as “smoothing” the accumulator image)

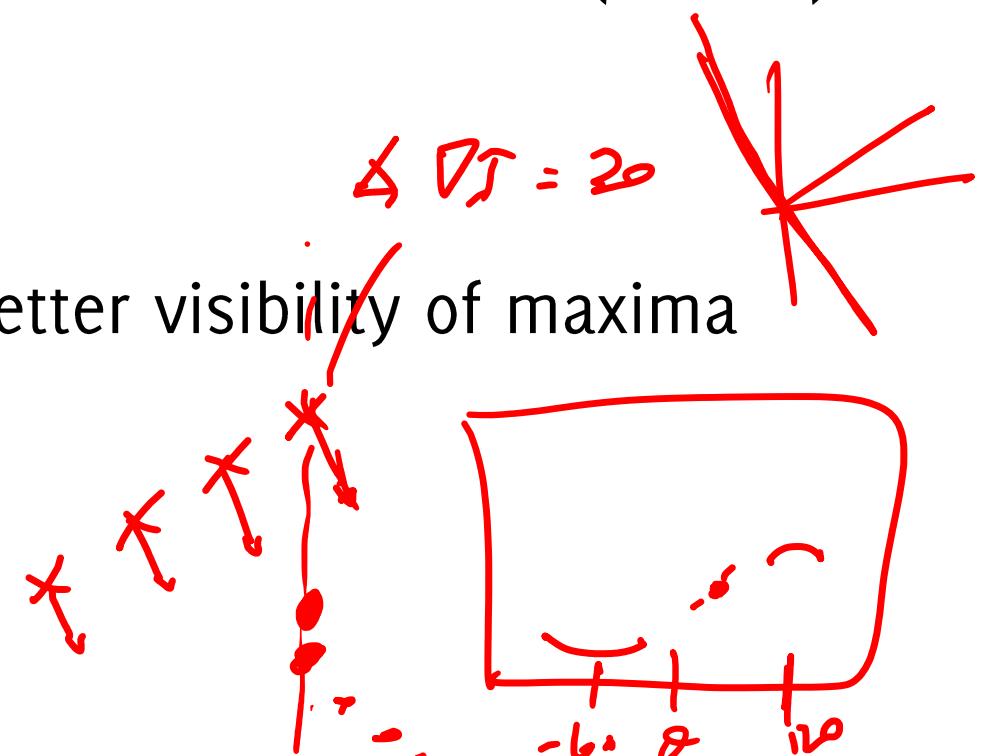
# Extension

From the edge detection algorithm, we know the direction of the gradient for each edge pixel

Remember how that **edge direction is orthogonal to gradient direction**

We can make sure an edge pixel only votes for lines that have (almost) the direction of the edge!

- Reduces the computation time
- Reduces the number of useless votes (better visibility of maxima corresponding to real lines)



# Hough Transform

The approach is not only limited to lines, but rather to any parametric model that we are able to fit

- Circles can be fit in a 3d accumulator space

It is quite robust to noise

# Hough Transform For Circles

slide Credits Alessandro Giusti, USI

# Hugh Transform for Circles

1. Every **edge point** casts votes for all **circles** that are compatible with it
2. We choose **circles** that accumulated a lot of votes

# How do we parametrize circles?

$$(x - a)^2 + (y - b)^2 = r^2$$

Center ( $x = a, y = b$ ) and radius  $r$  : 3 degrees of freedom

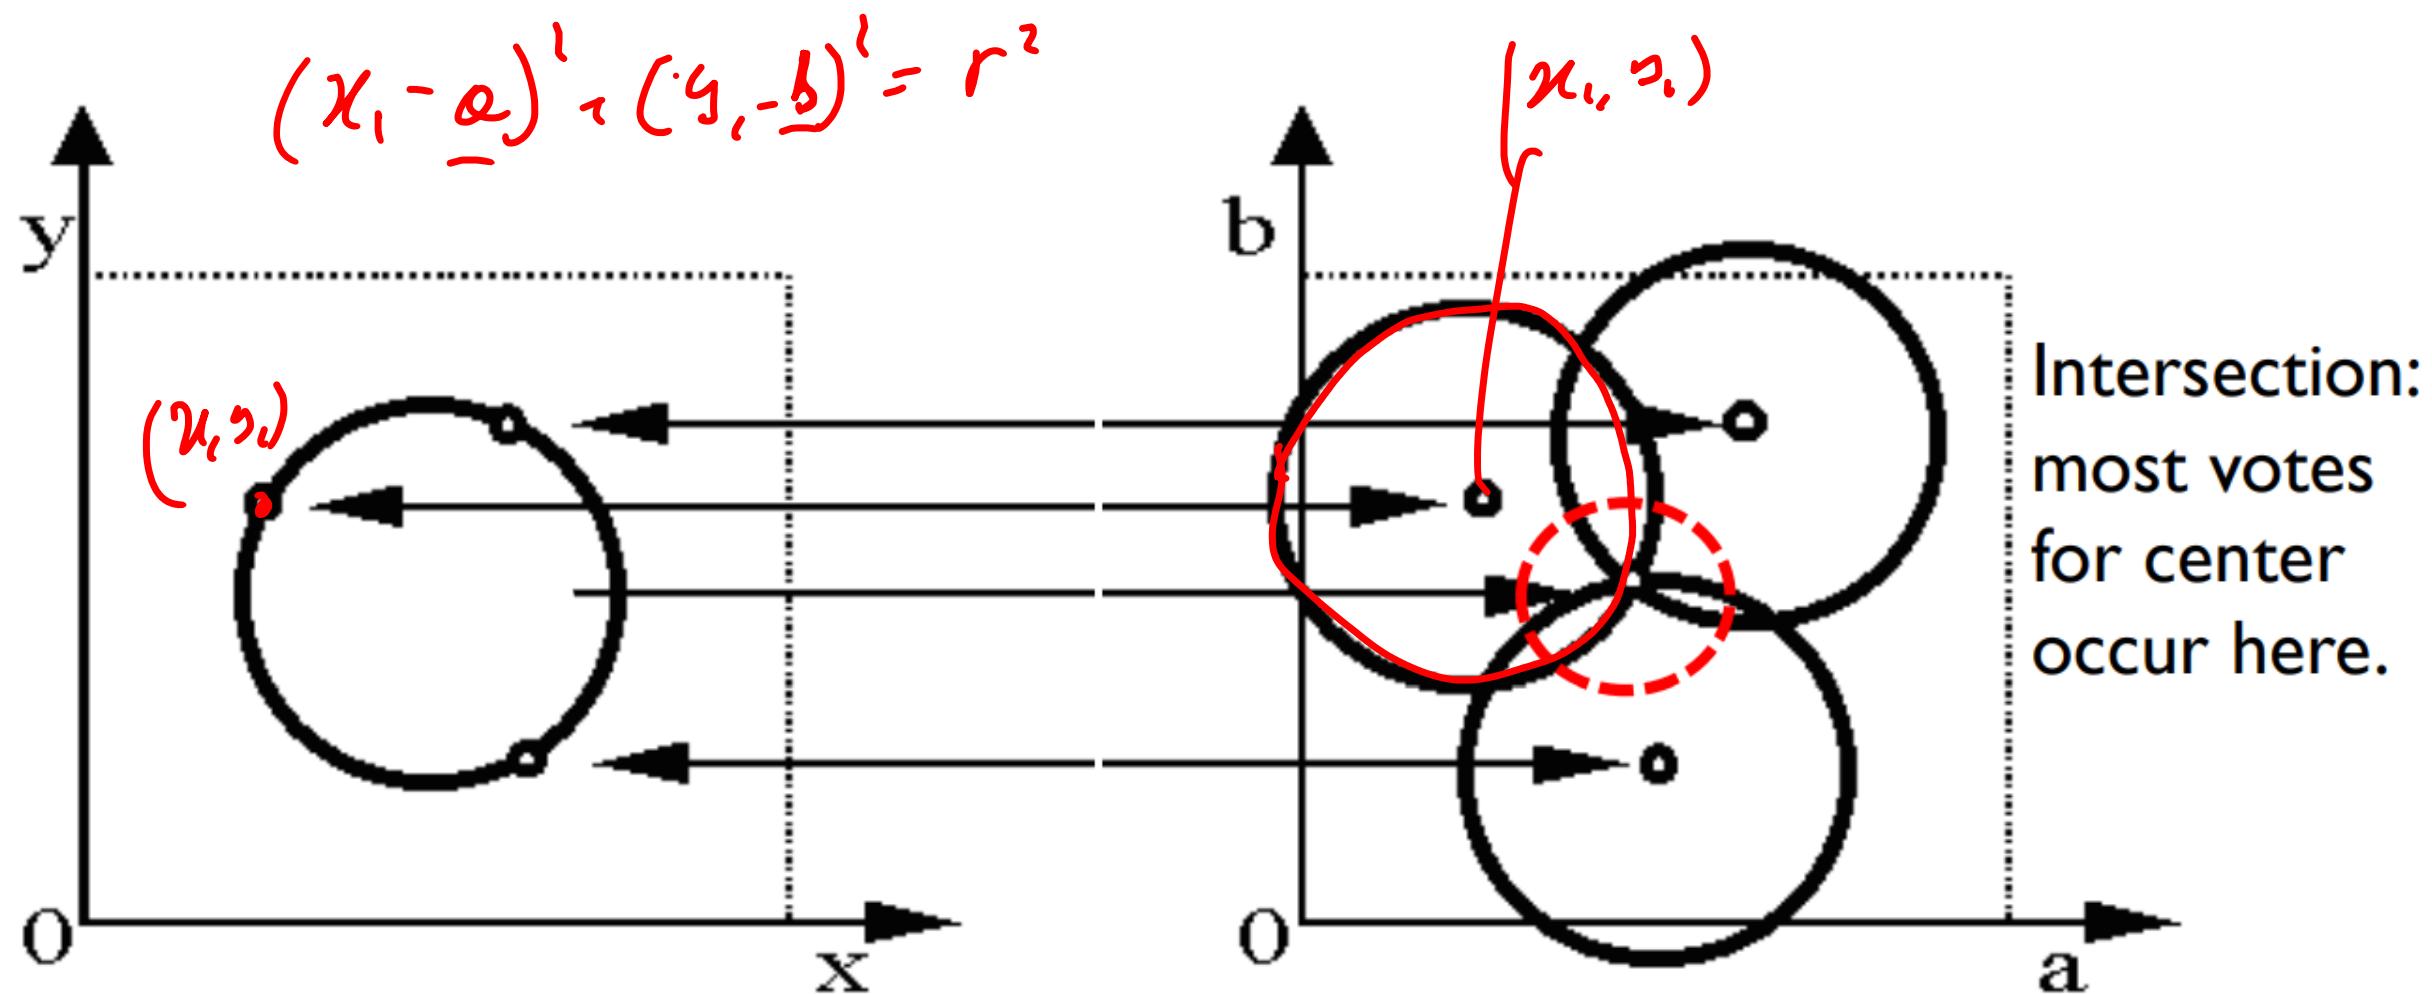
If we assume  $r$  known, the Hough space is 2D:

- $a$ :  $x$  coordinate of circle center
- $b$ :  $y$  coordinate of circle center

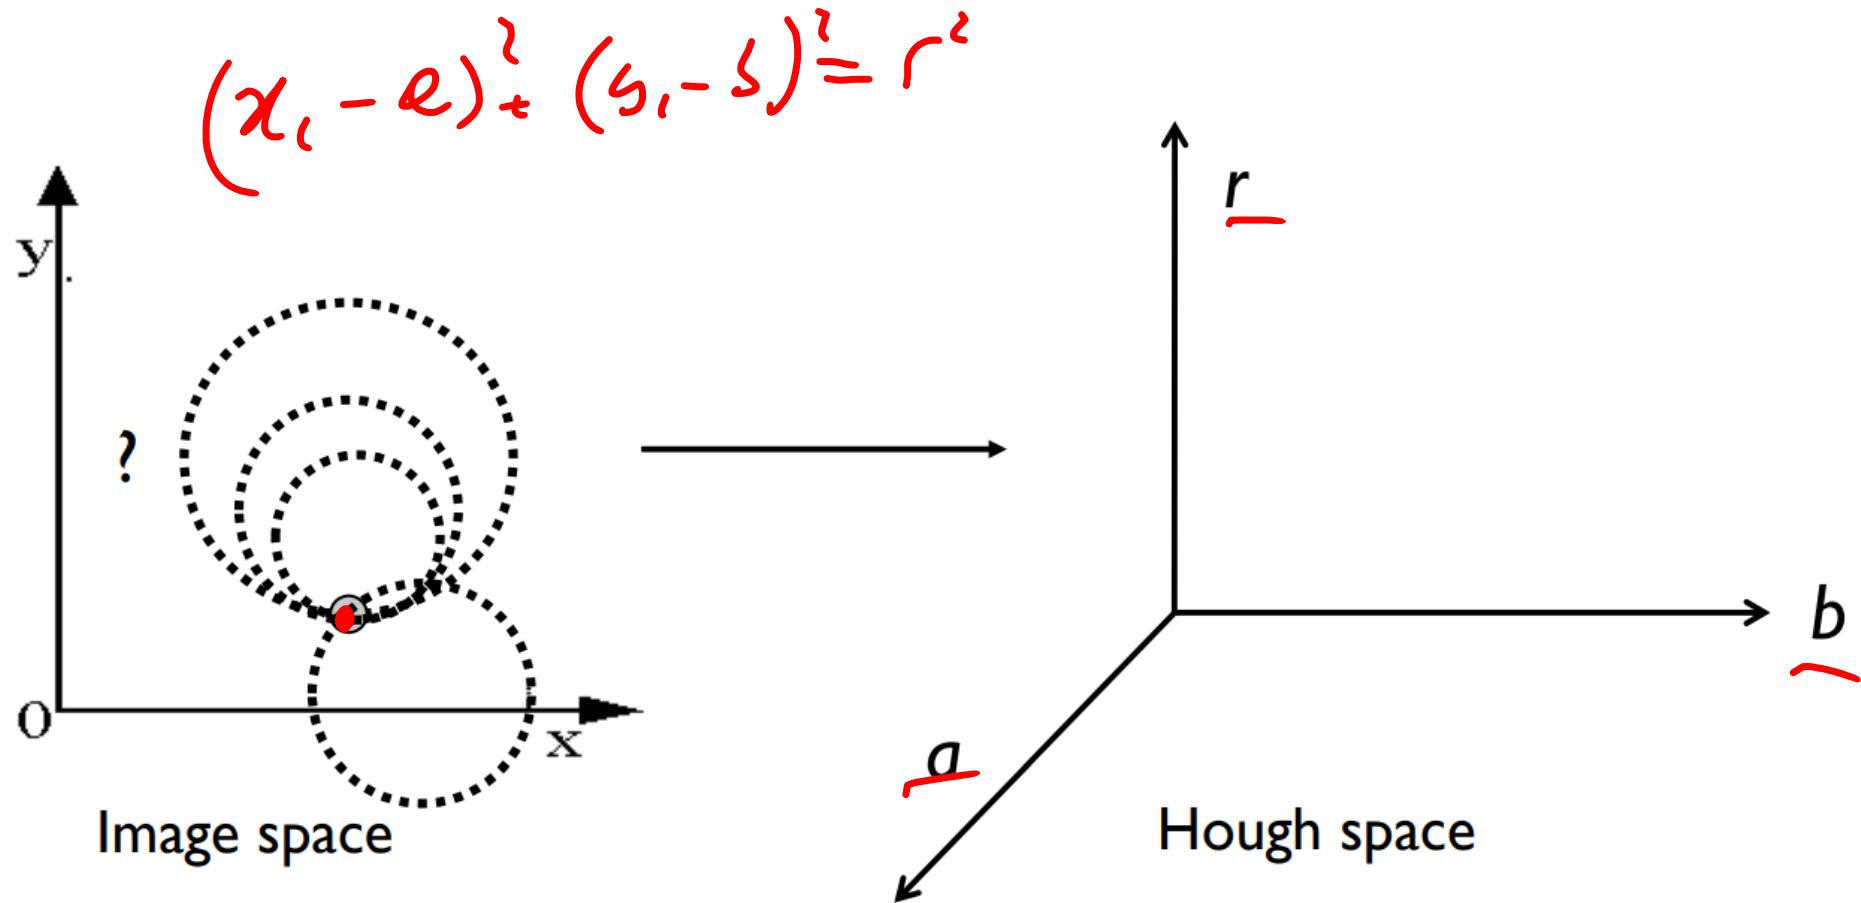
The role of  $(a, b)$  and  $(x, y)$  are interchangeable, thus:

One point in image space maps to a circle in Hough space

# Hough space for circles with known radius

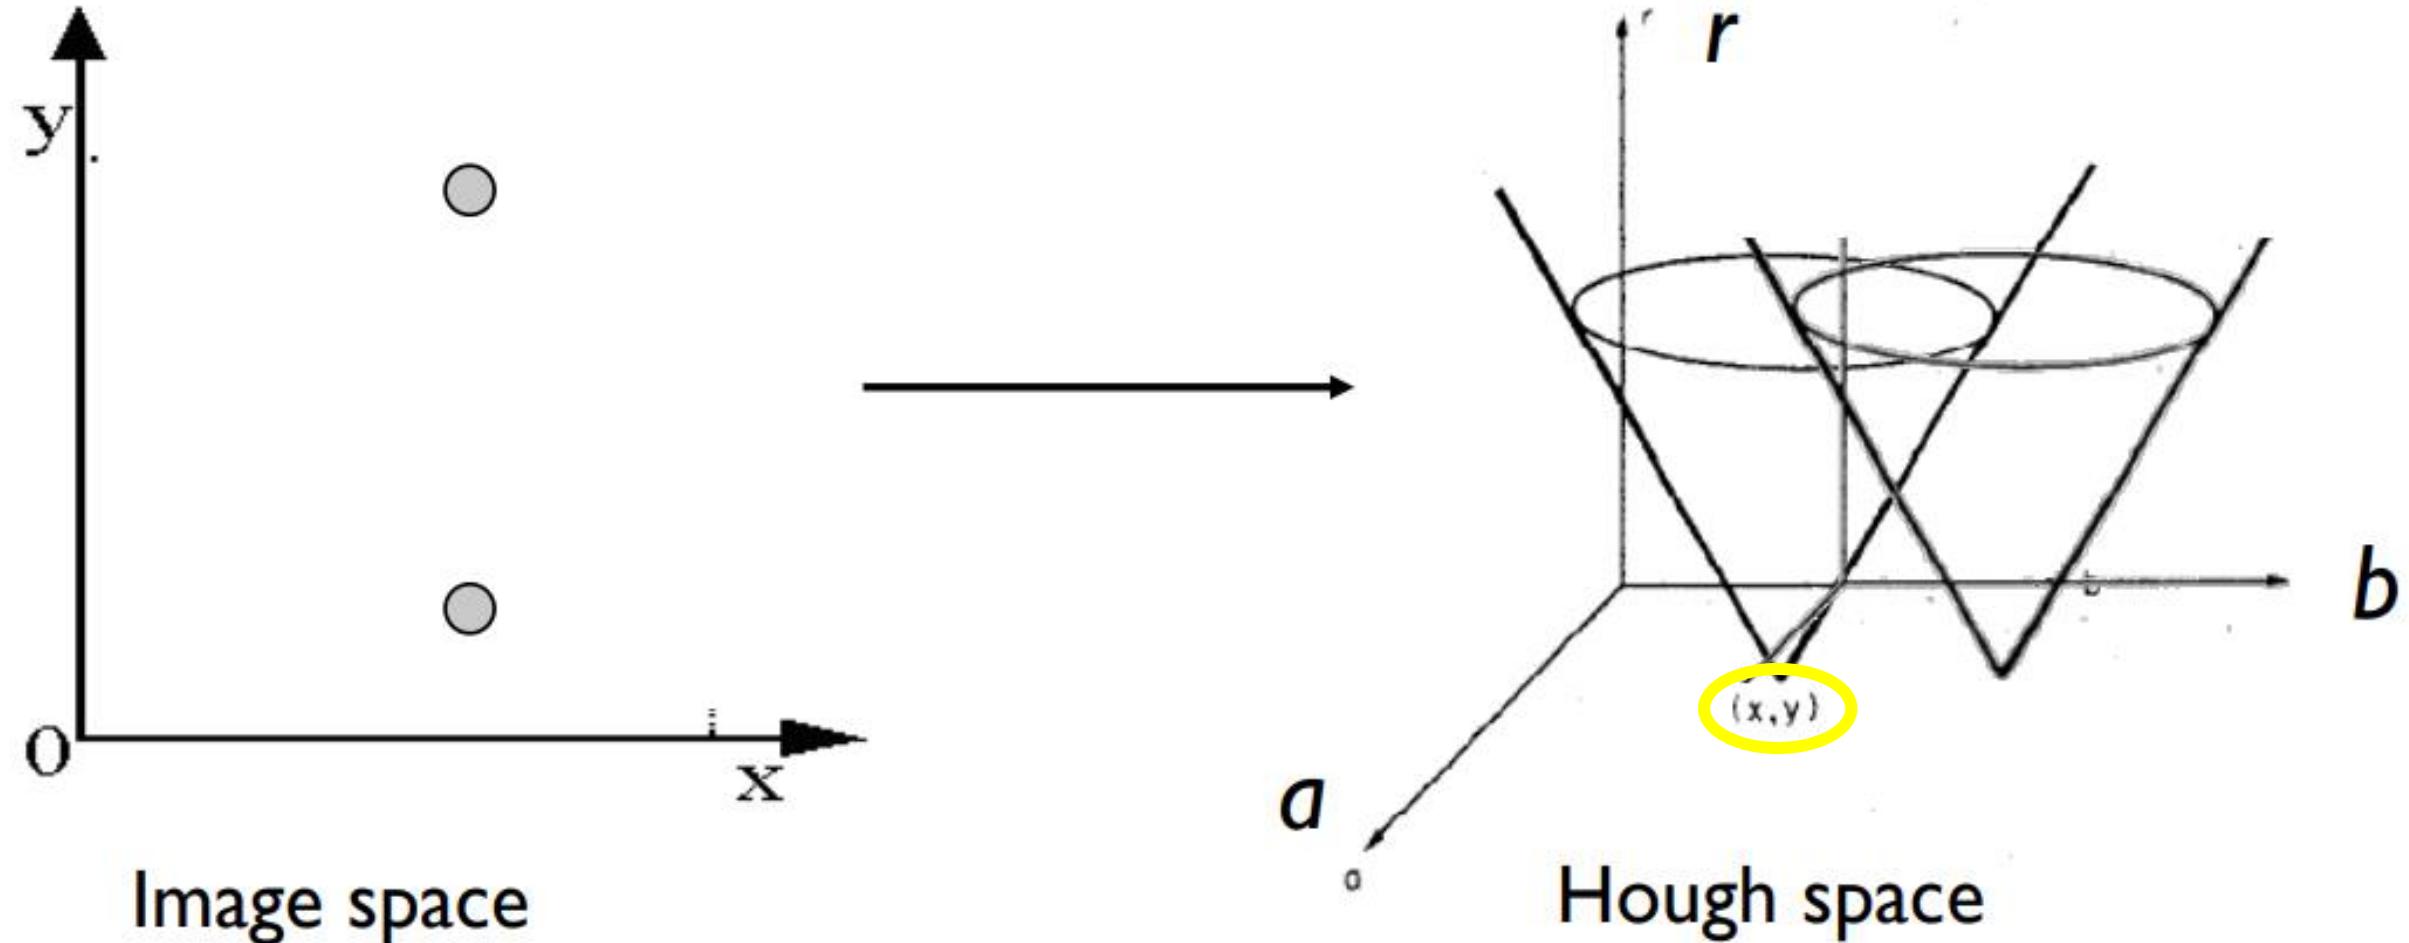


# Hough space for circles with unknown radius



One point in image space maps to...  
a cone in Hough space

# Hough space for circles with unknown radius



When the radius is zero  $(a, b) = (x, y)$

# If we know the gradient direction...

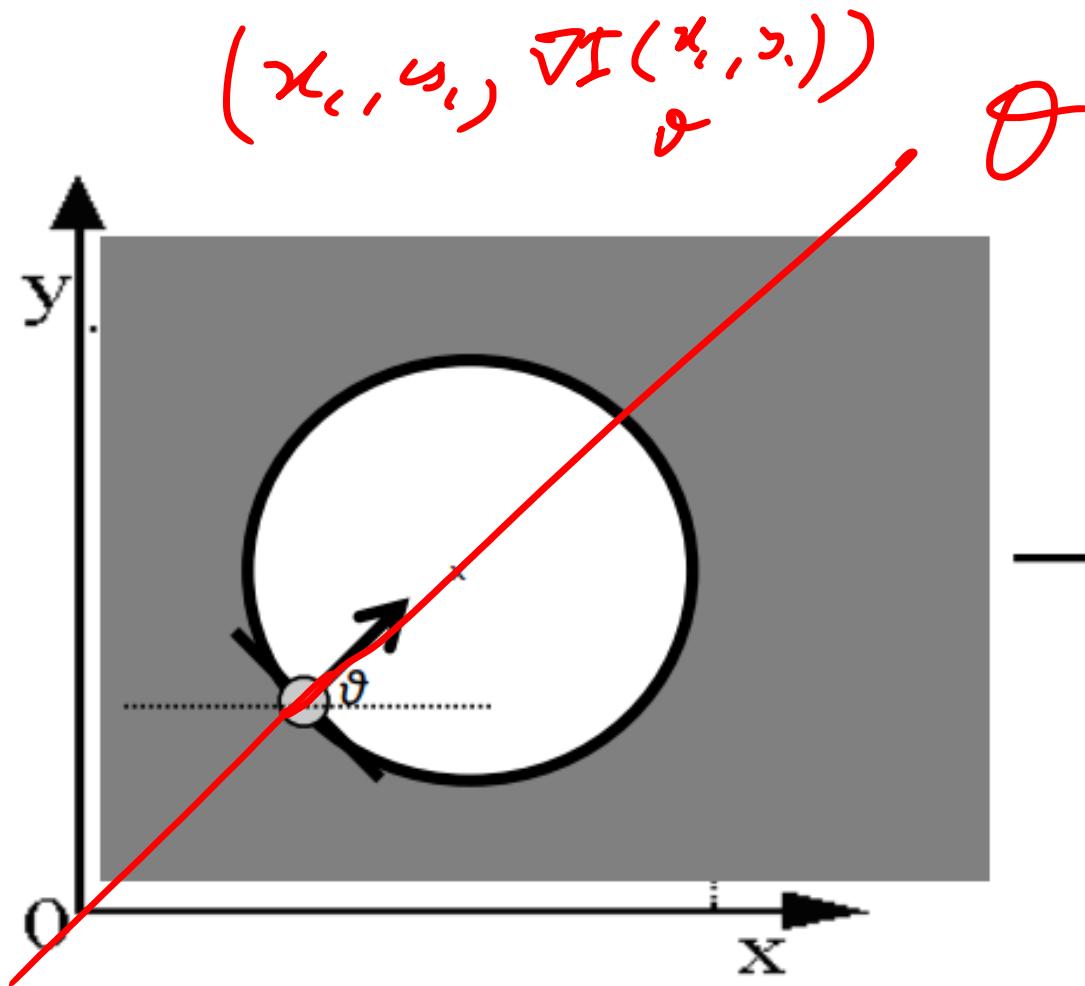
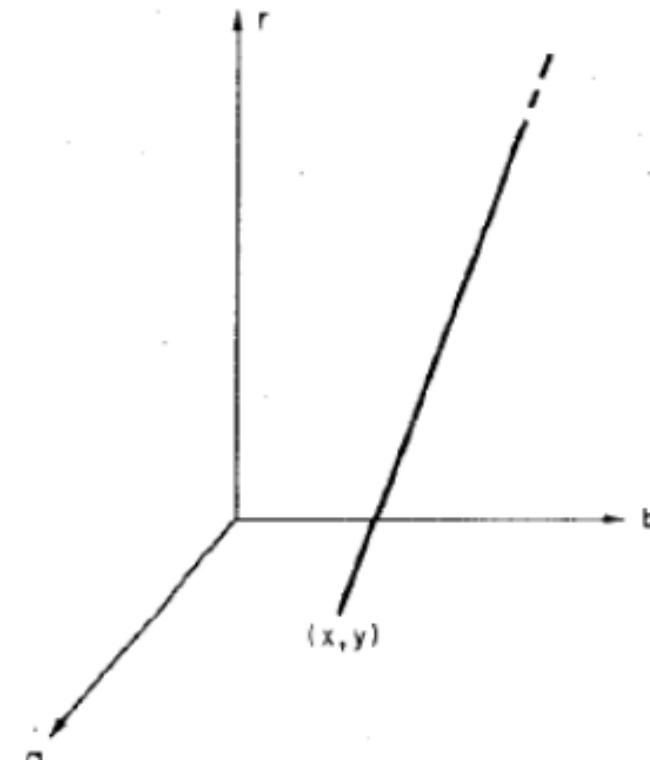


Image space

When increasing the radius, the center can only live in a line, thus the linear relation between  $a, b$



Hough space

# Hugh Transform for Circles

Initialize H accumulator to zeros

For every edge pixel  $(x, y)$ :

    For each possible radius value  $r$ :

        For each possible gradient direction  $\theta$ :

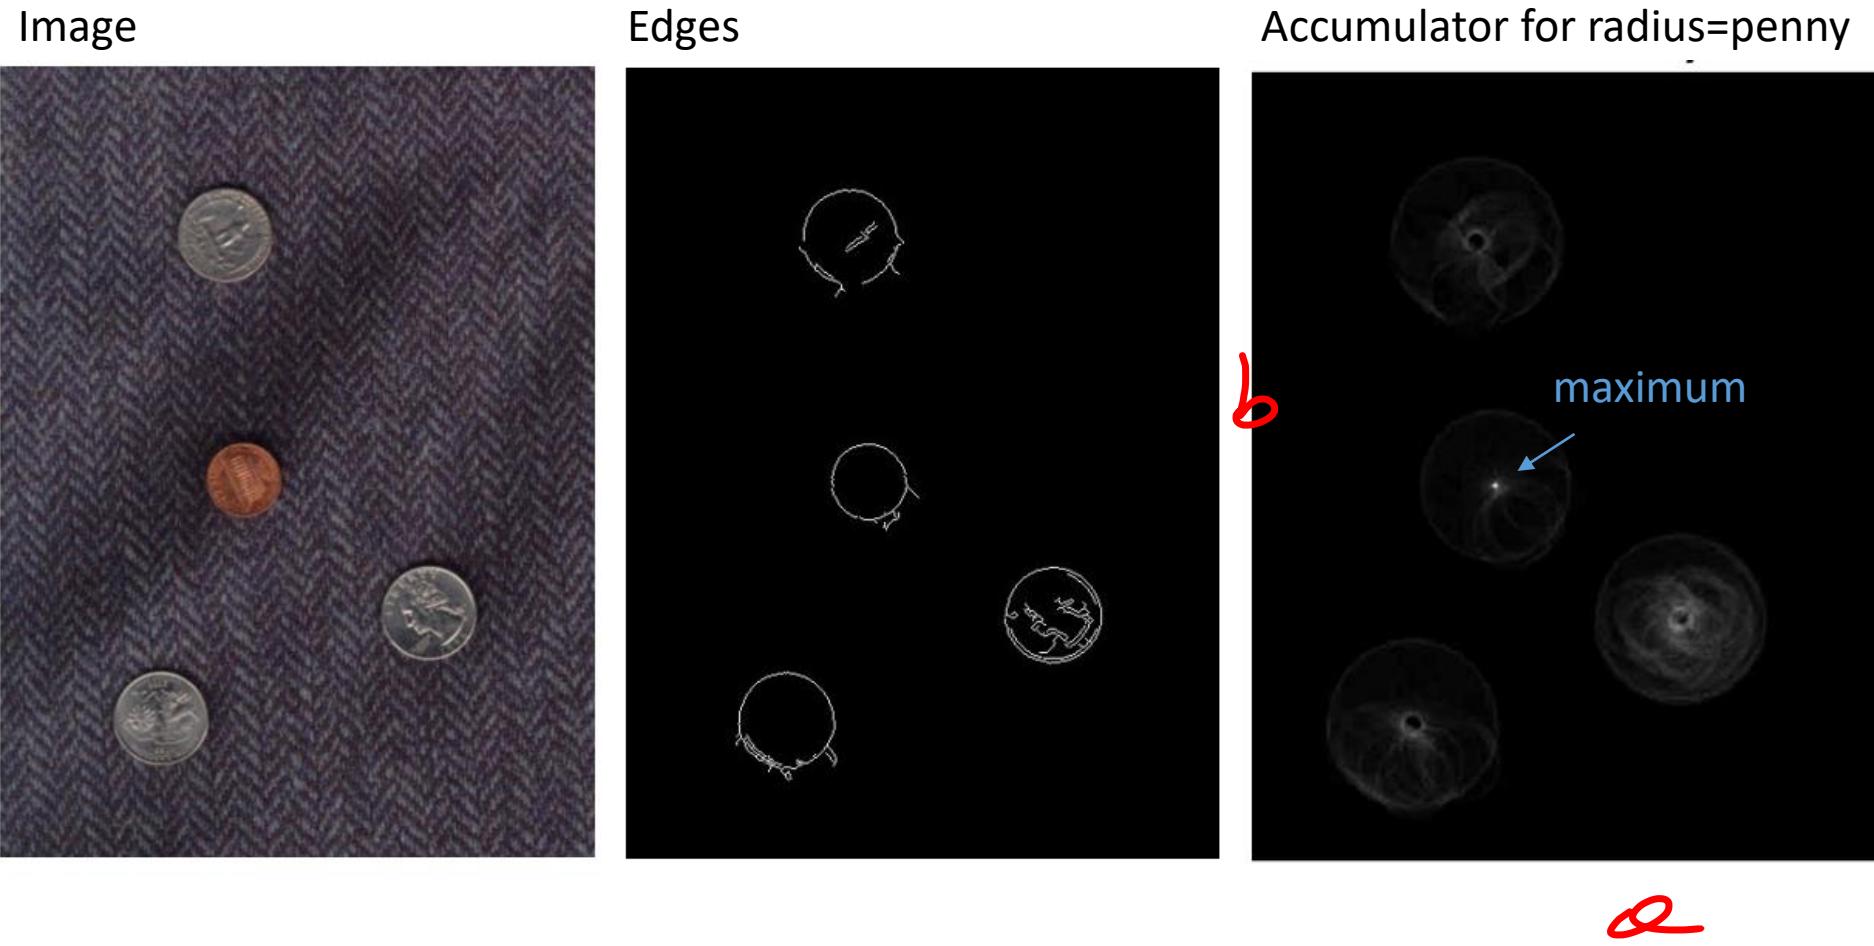
$a = x - r \cos(\theta)$  // column

$b = y + r \sin(\theta)$  // row

$H[a, b, r] += 1$

# An example

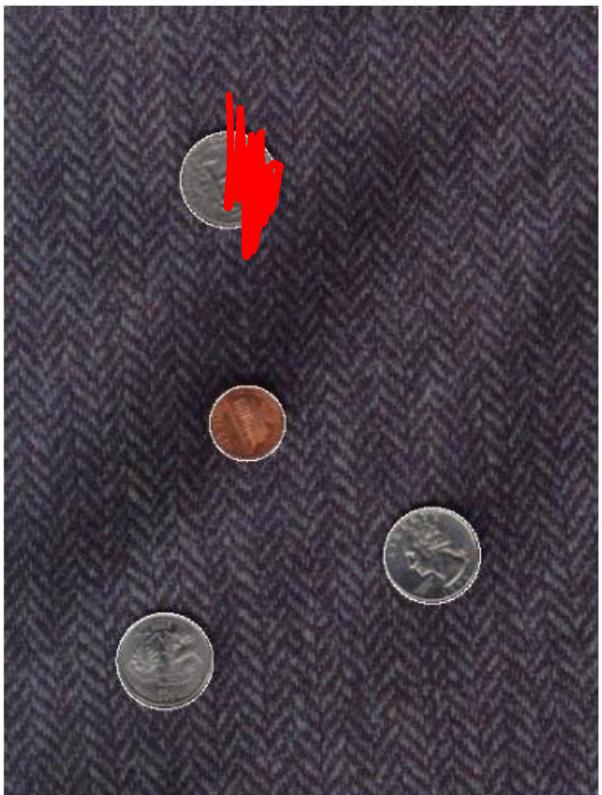
Accumulator for radius equal to radius of a penny



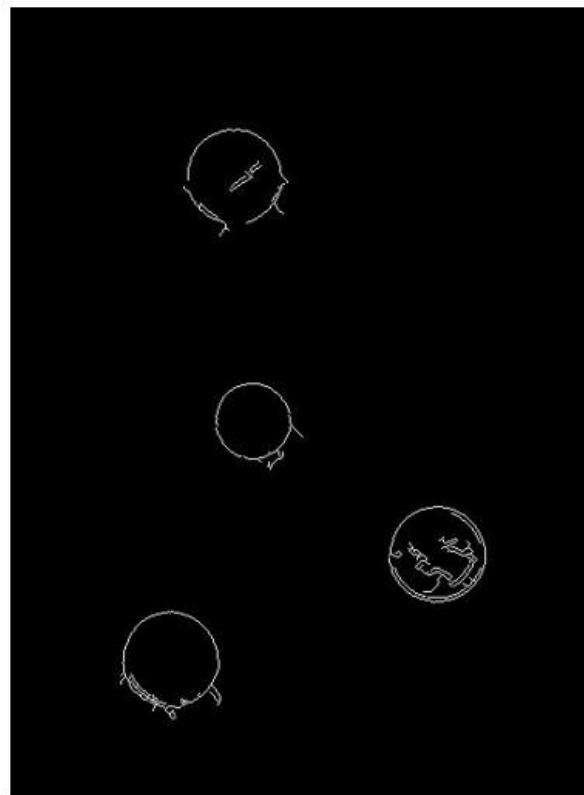
# An example

Accumulator for radius equal to radius of a quarter

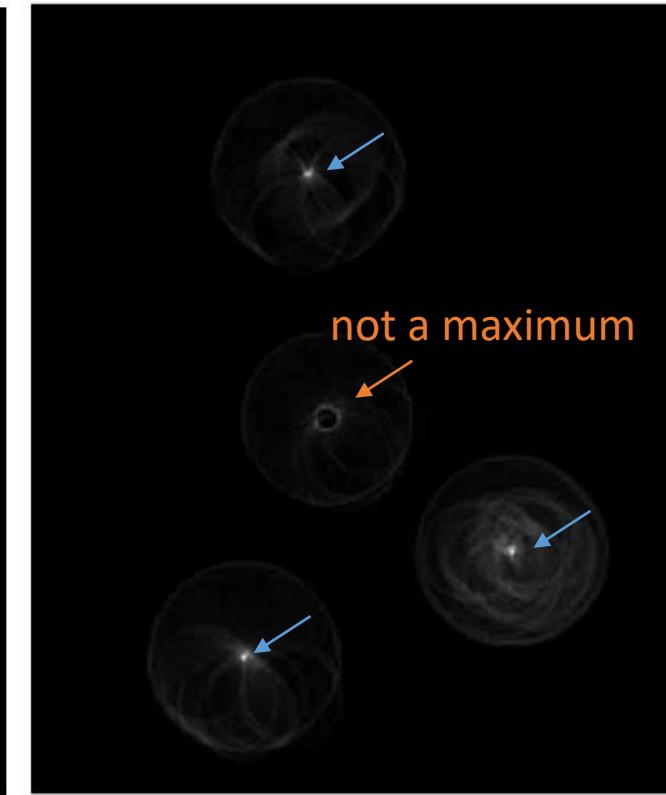
Image



Edges



Accumulator for radius=quarter



# Conclusions

## Advantages

- All points are processed independently, so **the algorithm can cope with occlusions and gaps**
- Voting algorithms are **robust to clutter**, because points not corresponding to any model are unlikely to contribute consistently to any single bin
- Can detect **multiple instances of a model** in a single pass

## Disadvantages

- Only suitable for models with **few parameters**
- Must filter out spurious peaks in hough accumulator
- Quantization of hough space is tricky

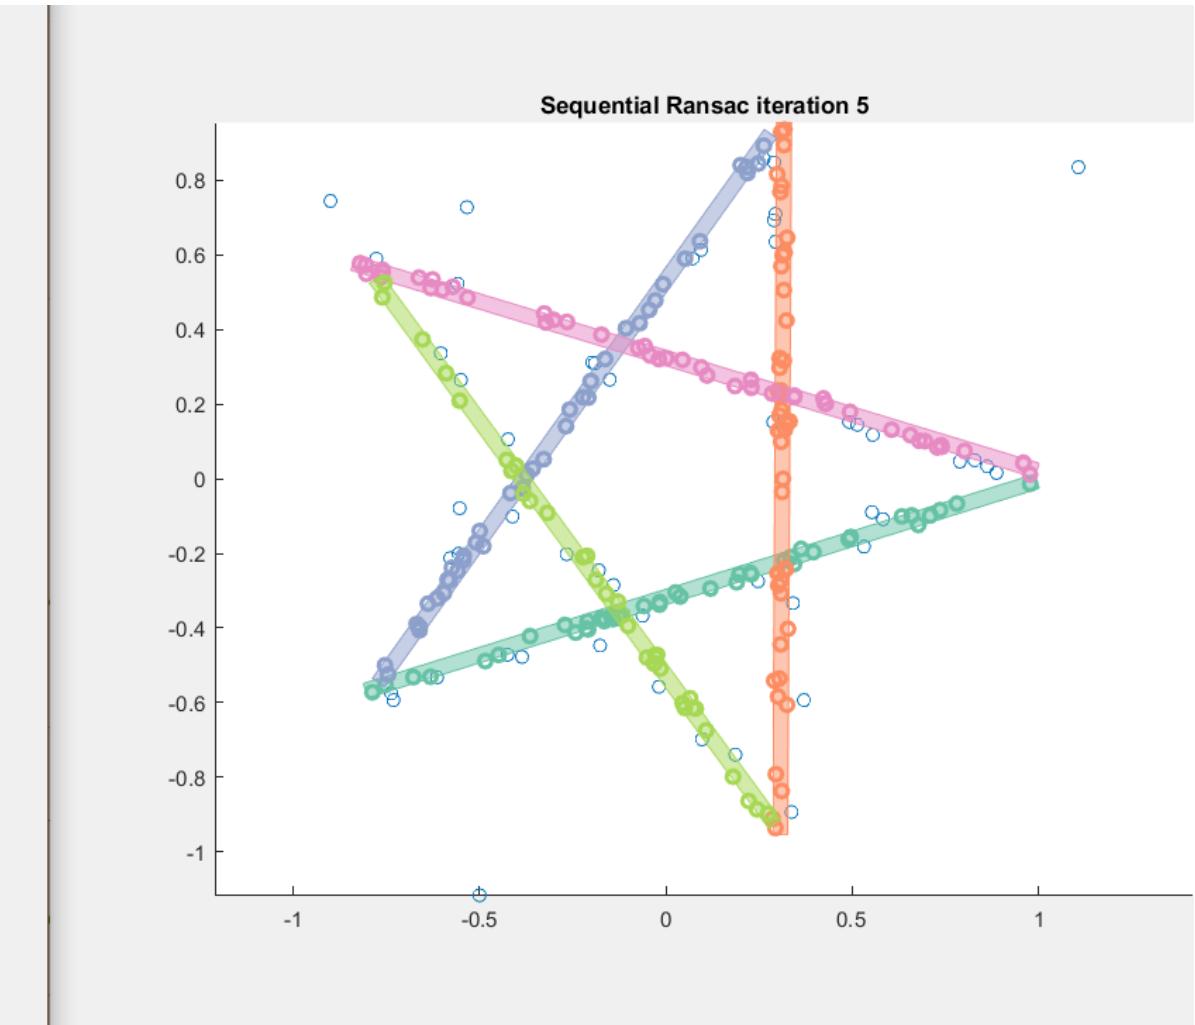
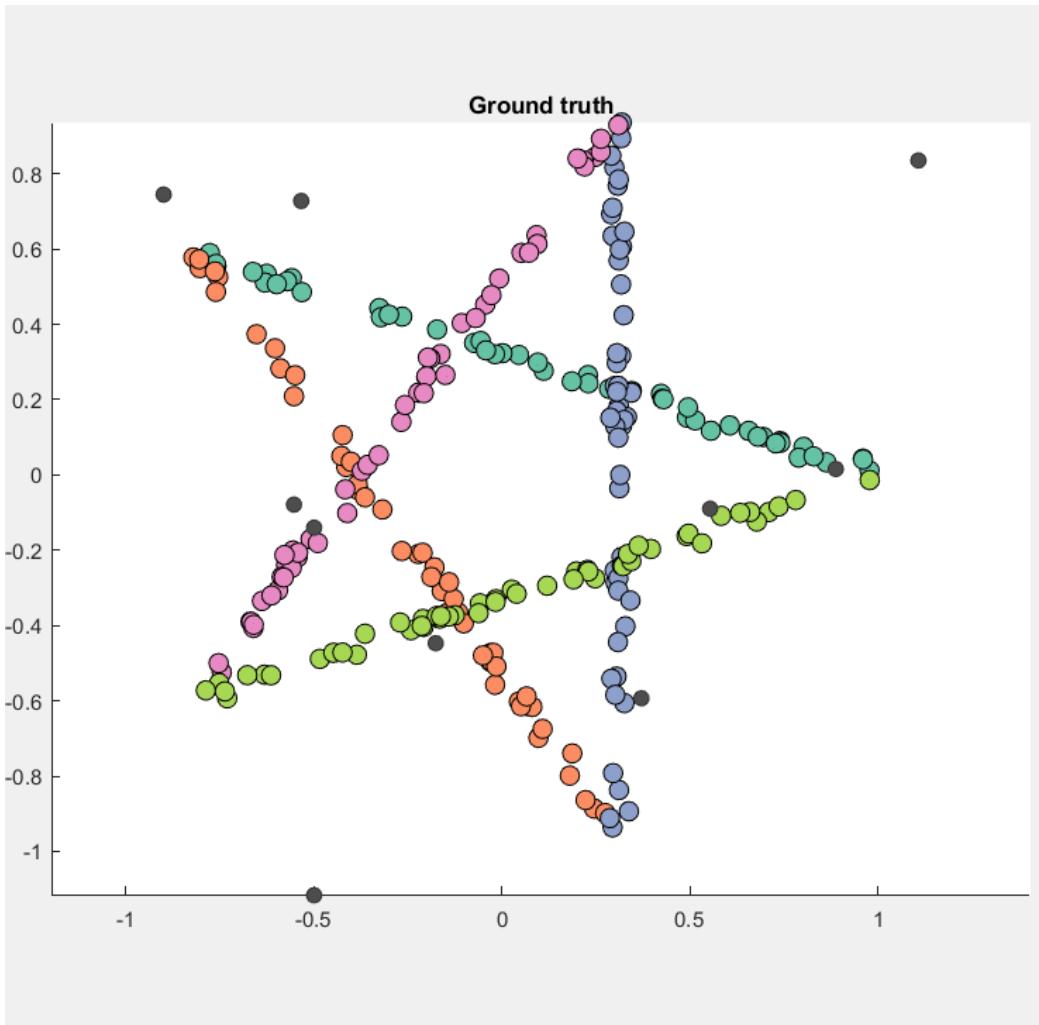
# Assigments

# **demo\_robustmmf\_TODO**

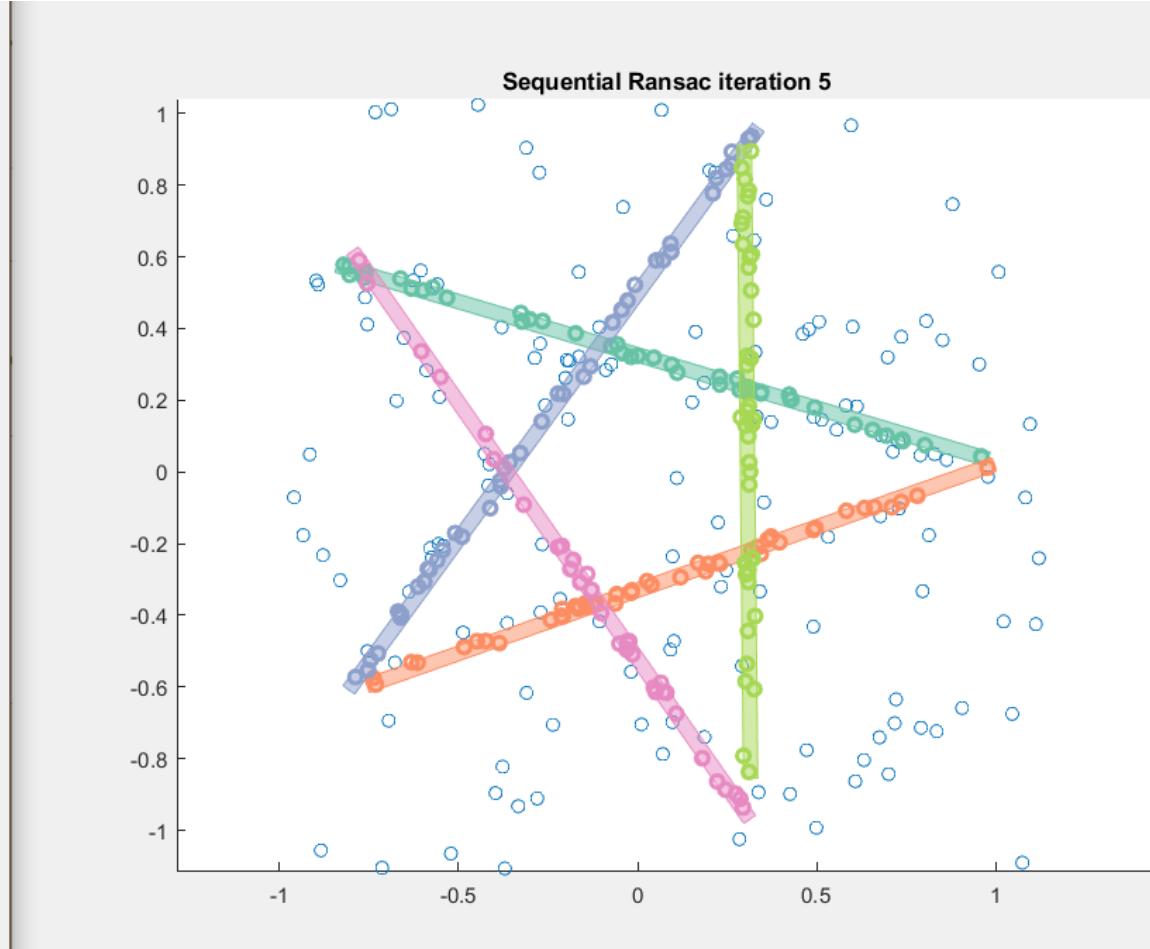
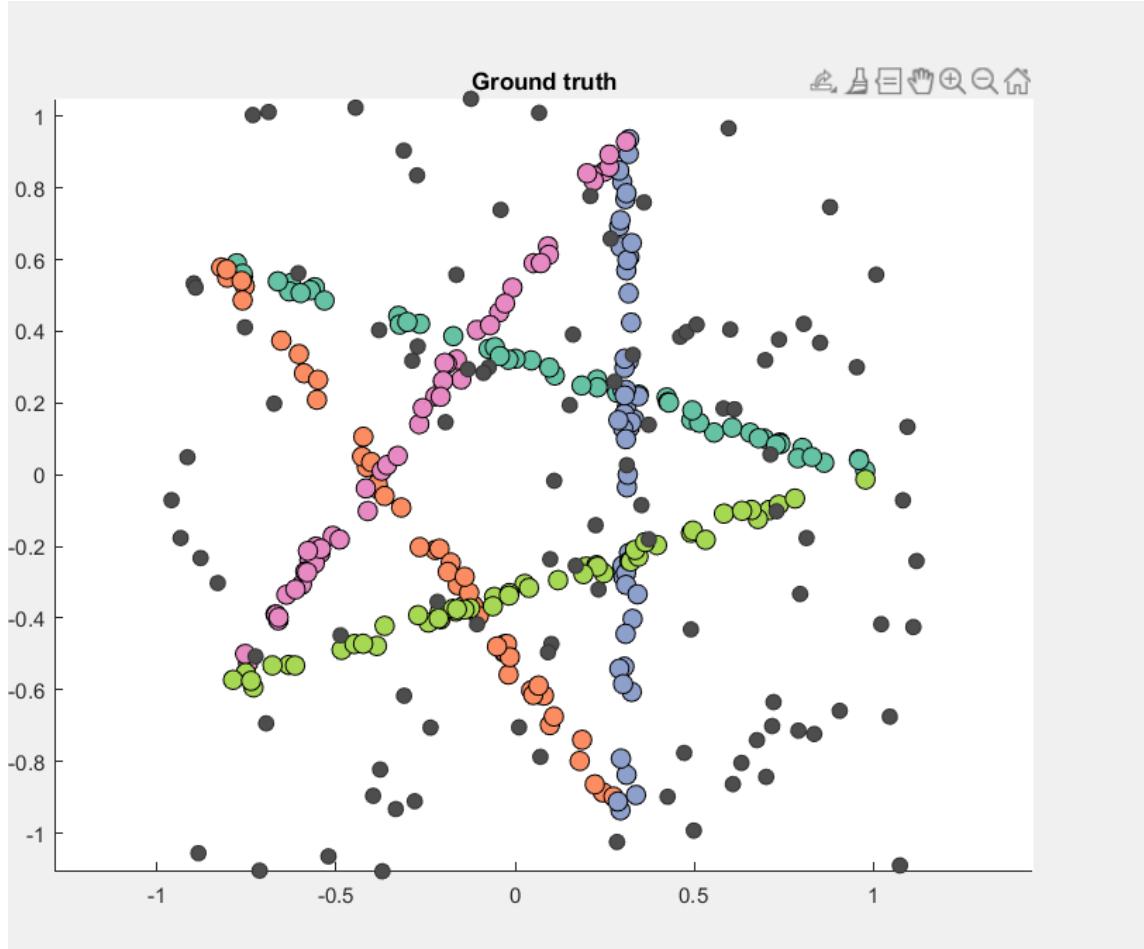
1. Implement Sequential Ransac for line fitting over
  1. the star5 data,
  2. The stair4 data

introducing different amount of outliers
2. Check the limitations of sequential ransac and test different stopping criteria (number of models retrieved, minimum consensus of the last model found)
3. Implement Ransac (thus run sequential Ransac) to fit circles

# Only 10 outliers



# 100 outliers



# The stairs case.. Sequential ransac terribly fails

